

Machine Learning 1.01: Introduction

Tom S. F. Haines
T.S.F.Haines@bath.ac.uk



Before we get started I...

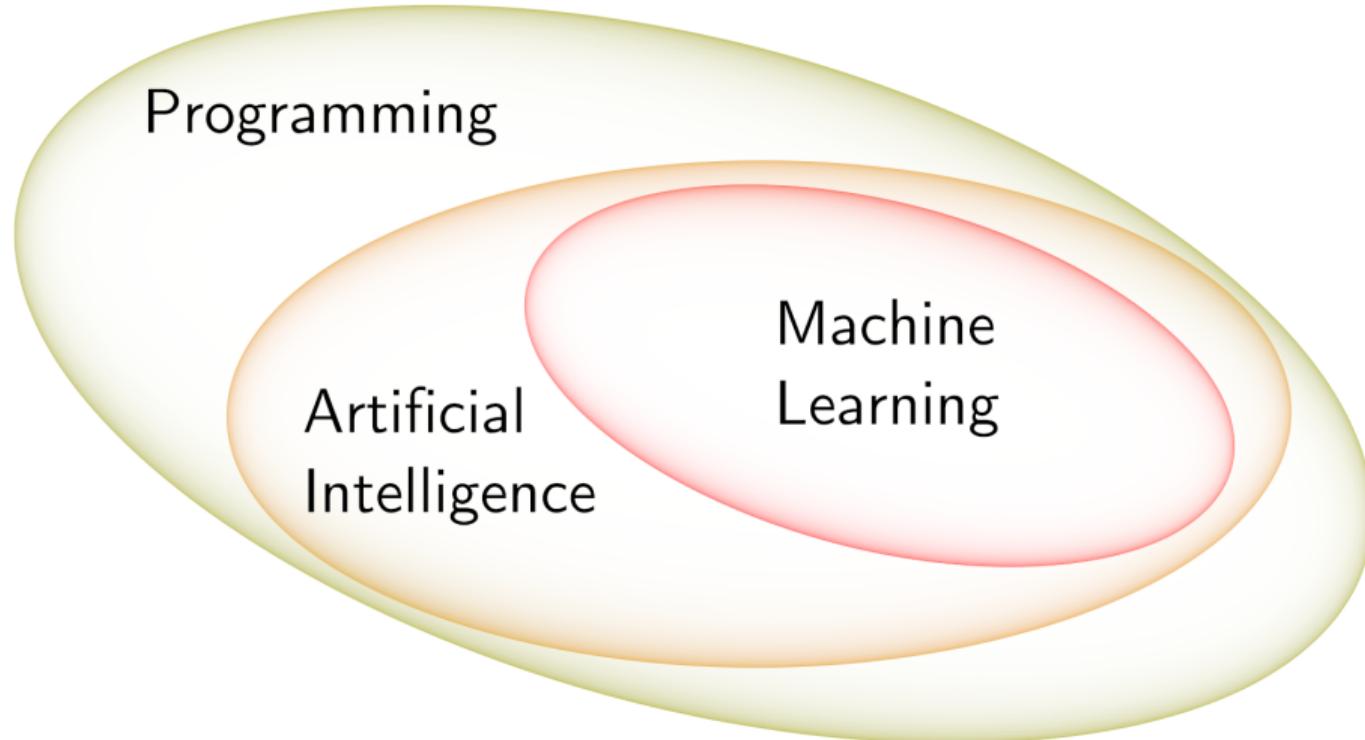
- Per week:
 - 2 hour lab session.
 - 2 lectures.
- Lectures:
 - 16 normal.
 - 1 to introduce optional modules.
 - 1 for recap.

Before we get started II. . .

- Marks:
 - 60% from 4 lab writeups (15% each).
 - 40% from final report (maximum 3000 words).
 - You will be submitting Jupyter workbooks!
- Using Jupyter:
 - Workbook via a web interface.
 - Python 3 + machine learning libraries.
 - Jupyter hub running at <http://dia.bath.ac.uk> (not yet!).
 - Of questionable value – best to use local install (Anaconda 3).
 - Do not run Anaconda from start menu – use shortcut on Moodle.

What is ML?

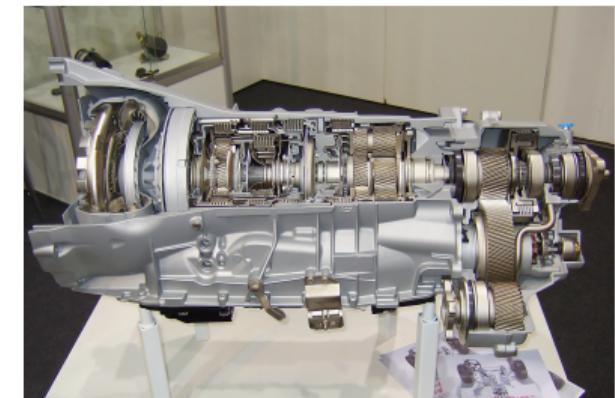
What is ML?



Imagine a car...

- Programming:
**Computer as an idiot –
it only does exactly what you tell it to.**
- e.g. automatic gear box:

```
while True:  
    if engine.revs > 5000 and  
        transmission.gear < 5:  
        clutch.disengage()  
        transmisison.gear += 1  
        clutch.engage()  
    elif engine.revs < 1000 and  
        transmission.gear > 1:  
        clutch.disengage()  
        transmisison.gear -= 1  
        clutch.engage()
```



Imagine a car...

- Artificial Intelligence:
Computer uses optimisation to find the best solution to a problem, but only knows exactly what you tell it.

- e.g. gps navigation:

```
graph.load_map('uk.h5')
graph.set_start('bath')
graph.set_end('bletchley')
route = graph.shortest_route()
```



Imagine a car...

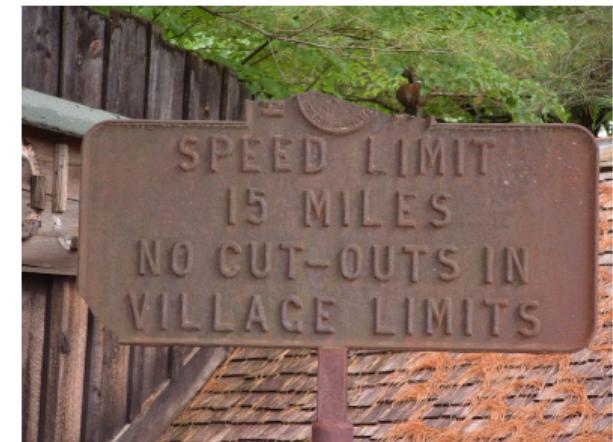
- Machine Learning:

Computer learns from examples and generalises to all cases.

- e.g. recognising road signs:

```
model = Recogniser('15mph_sign.h5')
while True:
    if model.search(camera.image()):
        engine.target = 6.7 # m/s
```

- How can this go wrong?



What is ML?

- Learning from data.

What is ML?

- **Learning from data.**
- Built on (you also need to understand):
 - Maths, especially probability.
 - Optimisation.
 - Programming.

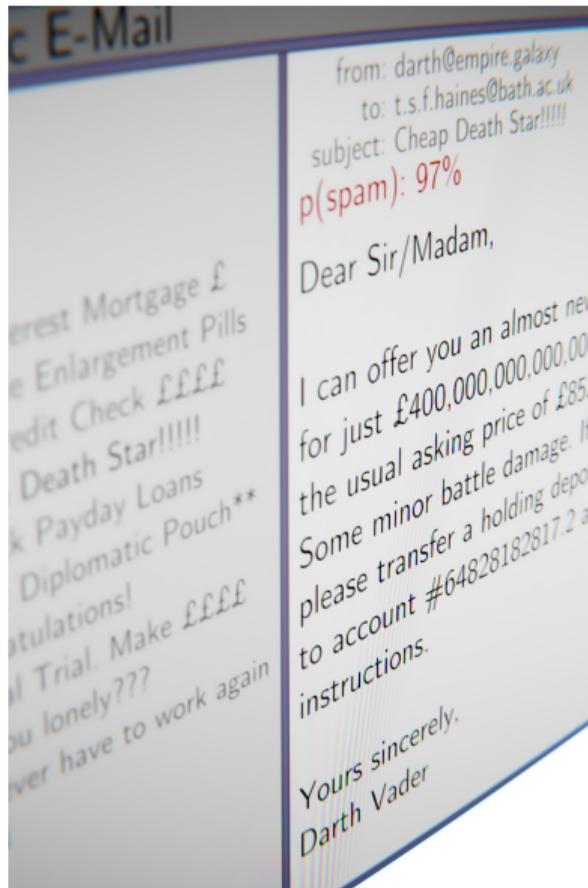
What is ML?

- **Learning from data.**
- Built on (you also need to understand):
 - Maths, especially probability.
 - Optimisation.
 - Programming.
- Warning 1: Not everyone would agree with this definition.
 - overlaps with statistical models, data mining

- **Learning from data.**
- Built on (you also need to understand):
 - Maths, especially probability.
 - Optimisation.
 - Programming.
- Warning 1: Not everyone would agree with this definition.
 - overlaps with statistical models, data mining
- Warning 2: ML is often used instead of AI because it's “fashionable”.

What can you do with it?

Motivation: Text



- Spam filtering.
 - Organising content (e.g. classifying news by area).
 - Search engines (e.g. classifying if a web page is accurate).
 - Automatic marking.

Motivation: Medical



- Identifying a disease.
- Identifying an abnormality.
- Drug development (e.g. predicting which chemical to test).

Motivation: Human-computer interaction



- Handwriting or speech recognition.
- Personal assistants (e.g. Siri).
- Recommendation systems (e.g. product suggestions on Amazon).
- Automated design.

Motivation: People



- Advertising, polling.
- Text analysis for contracts and law enforcement.
- Detecting terrorism.

Motivation: Financial



- Detecting fraud.
- Mortgage applications.
- Algorithmic trading.

Motivation: Robots



- Walking robots.
- Autonomous vehicles.
- Playing games (e.g. Deep Mind's Alpha Go).

The process

1. Choose a **problem**.
2. Obtain required **data**.
3. Choose or design a **model**.
4. Fit model to data using **optimisation**.
5. **Measure** performance.

(there are variants...)

1. Choose a problem

e.g. this toy problem:

Given something in the ocean identify if it is a **fish** or an **invertebrate**.

1. Choose a problem

e.g. this toy problem:

Given something in the ocean identify if it is a **fish** or an **invertebrate**.

Input: Yes/no answers to questions such as:

Does it have teeth?

Output: Fish or invertebrate.

2. Obtain required data

Animal name	bass	clam	carp	crab	catfish	crayfish	chub	lobster
Does it have teeth?	1	0	1	0	1	0	1	0
Does it breathe?	0	0	0	0	0	0	0	0
Does it have a backbone?	1	0	1	0	1	0	1	0
Is it aquatic?	1	0	1	1	1	1	1	1
Does it have a tail?	1	0	1	0	1	0	1	0
Is it a predator?	1	1	0	1	1	1	1	1
Is it an invertebrate?	0	1	0	1	0	1	0	1

- Use of 1 for “yes” and 0 for “no” is typical.
- This module will be ignoring real data collection.

3. Choose or design a model

- This is a **classification** problem – the output is a discrete label.
- It is one of the two main problem types in machine learning.

3. Choose or design a model

- This is a **classification** problem – the output is a discrete label.
- It is one of the two main problem types in machine learning.
- There are hundreds of models for solving it.
- Lets use another “model”: A rule (algorithm) created by a human!

4. Fit model to data using optimisation

You have three minutes to come up with an algorithm:

Animal name	bass	clam	carp	crab	catfish	crayfish	chub	lobster
Does it have teeth?	1	0	1	0	1	0	1	0
Does it breathe?	0	0	0	0	0	0	0	0
Does it have a backbone?	1	0	1	0	1	0	1	0
Is it aquatic?	1	0	1	1	1	1	1	1
Does it have a tail?	1	0	1	0	1	0	1	0
Is it a predator?	1	1	0	1	1	1	1	1
Is it an invertebrate?	0	1	0	1	0	1	0	1

Write your algorithm down!

5. Measure performance I

- Previous slide was a **training set**.
- Below is a **testing set**:

Animal letter Animal name	A	B	C	D	E	F	G
Does it have teeth?	1	0	0	1	1	0	1
Does it breathe?	0	0	1	0	0	0	1
Does it have a backbone?	1	0	0	1	1	0	1
Is it aquatic?	1	1	0	1	1	1	0
Does it have a tail?	1	0	1	1	1	0	1
Is it a predator?	1	1	1	0	1	1	1

- Apply algorithm and record results.

5. Measure performance II

Animal letter	A	B	C	D	E	F	G
Animal name	dogfish	octopus	scorpion	haddock	pike	seawasp	bear
Does it have teeth?	1	0	0	1	1	0	1
Does it breathe?	0	0	1	0	0	0	1
Does it have a backbone?	1	0	0	1	1	0	1
Is it aquatic?	1	1	0	1	1	1	0
Does it have a tail?	1	0	1	1	1	0	1
Is it a predator?	1	1	1	0	1	1	1
Is it an invertebrate?	0	1	1	0	0	1	mammal

- How well did your algorithm do? (ignore the bear!)

5. Measure performance II

Animal letter	A	B	C	D	E	F	G
Animal name	dogfish	octopus	scorpion	haddock	pike	seawasp	bear
Does it have teeth?	1	0	0	1	1	0	1
Does it breathe?	0	0	1	0	0	0	1
Does it have a backbone?	1	0	0	1	1	0	1
Is it aquatic?	1	1	0	1	1	1	0
Does it have a tail?	1	0	1	1	1	0	1
Is it a predator?	1	1	1	0	1	1	1
Is it an invertebrate?	0	1	1	0	0	1	mammal

- How well did your algorithm do? (ignore the bear!)
- Is the bear unreasonable?
- Does the algorithm really ask “*Is it an invertebrate?*”?

What was your algorithm?

What was your algorithm?

1. Has tail \implies fish.

- This is true for the training set, but violated by scorpions.



What was your algorithm?

1. Has tail \implies fish.

- This is true for the training set, but violated by scorpions.

2. Has backbone \implies fish.

- Official biological definition.
- Not always obvious! e.g. caterpillars.



What was your algorithm?

1. Has tail \implies fish.

- This is true for the training set, but violated by scorpions.

2. Has backbone \implies fish.

- Official biological definition.
- Not always obvious! e.g. caterpillars.

3. Has teeth \implies fish.

- Defined to be true, and much more visible.
- Invertebrates can have teeth-equivalent structures, e.g. a snail.



What was your algorithm?

1. Has tail \implies fish.

- This is true for the training set, but violated by scorpions.

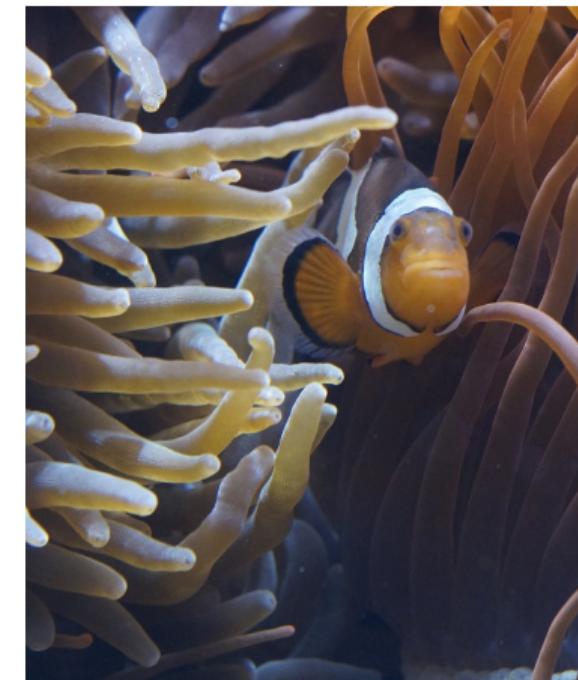
2. Has backbone \implies fish.

- Official biological definition.
- Not always obvious! e.g. caterpillars.

3. Has teeth \implies fish.

- Defined to be true, and much more visible.
- Invertebrates can have teeth-equivalent structures, e.g. a snail.

4. Any others?



What was your algorithm?

1. Has tail \implies fish.

- This is true for the training set, but violated by scorpions.

2. Has backbone \implies fish.

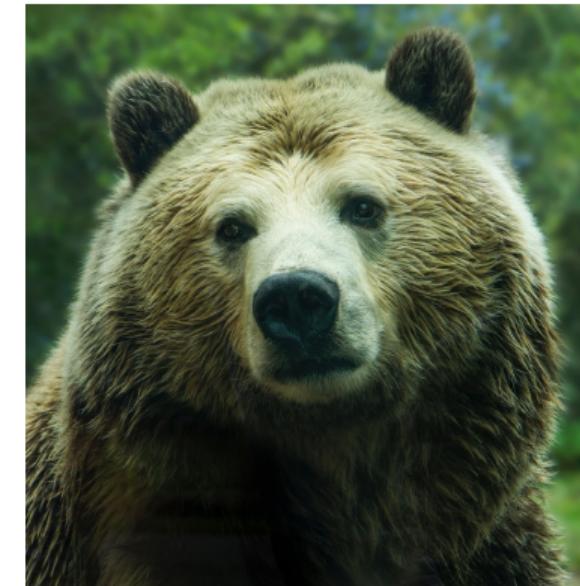
- Official biological definition.
- Not always obvious! e.g. caterpillars.

3. Has teeth \implies fish.

- Defined to be true, and much more visible.
- Invertebrates can have teeth-equivalent structures, e.g. a snail.

4. Any others?

If you get the right answer, does how really matter?



What happened?

1. You found a rule that solved the problem for training data.
2. You applied the rule to (testing) data.

What happened?

1. You found a rule that solved the problem for training data.
 2. You applied the rule to (testing) data.
-
- You could program step 2, e.g.

```
def invertebrate(fv):  
    return fv[ 'teeth' ] == False
```

What happened?

1. You found a rule that solved the problem for training data.
 2. You applied the rule to (testing) data.
-
- You could program step 2, e.g.

```
def invertebrate(fv):  
    return fv[ 'teeth' ] == False
```

- But what was your algorithm for step 1?

Supplementary definition

- Machine Learning is discovering the rule (step 1).
- Using the rule is just programming (step 2).

Supplementary definition

- Machine Learning is discovering the rule (step 1).
- Using the rule is just programming (step 2).
- Supplementary definition: A Machine Learning algorithm outputs code!

Supplementary definition

- Machine Learning is discovering the rule (step 1).
- Using the rule is just programming (step 2).
- Supplementary definition: A Machine Learning algorithm outputs code!
- But parameters are more practical than code, e.g.

```
# Learn these:  
feature = 'teeth'  
match = False  
  
# Code for model:  
def evaluate(fv):  
    return fv[feature] == match
```

Summary

- What is ML?
- Use cases.
- The typical process.
- Walk through with yourselves instead of computers.

Further reading & sources

- “Information Theory, Inference and Learning Algorithms” by **David J. C. MacKay**.
- “Pattern Recognition and Machine Learning” by **Christopher M. Bishop**.
- “Computer Vision: Models, Learning, and Inference” by **Simon J. D. Prince**.
- Zoo animal classification: <https://archive.ics.uci.edu/ml/datasets/Zoo>.

Image of automatic transmission: CC BY-ND 2.0 Kecko

<https://www.flickr.com/photos/70981241@N00/1876479840>

Image of speed sign: CC BY-SA 3.0 Jayron32 https://en.wikipedia.org/wiki/File:Antique_New_Hampshire_speed_limit_sign.jpg

HAL: Used without permission under fair use. From film 2001: A Space Odyssey.

Google car: Stolen from internet, used without permission under fair use.