

# Machine Learning 1.10: Curse of Dimensionality

Tom S. F. Haines  
T.S.F.Haines@bath.ac.uk



## Curse of Dimensionality

- Given a name by Richard Bellman  
(“Adaptive Control Processes: A Guided Tour”, 1961)  
(Inventor of dynamic programming – next lecture)
- Short version: More dimensions  $\implies$  need more data  
(dimension = feature)

# Curse of Dimensionality

- Given a name by Richard Bellman  
(“Adaptive Control Processes: A Guided Tour”, 1961)  
(Inventor of dynamic programming – next lecture)
- Short version: More dimensions  $\implies$  need more data  
(dimension = feature)
- Why?
- How much more?

# Combinatorics

- Imagine a problem:
  - Feature vector of length  $n$
  - Binary – answers to *yes/no* questions
  - Any output


Typical of surveys and psychological test.

## Combinatorics

- Imagine a problem:
    - Feature vector of length  $n$
    - Binary – answers to *yes/no* questions
    - Any output
- Typical of surveys and psychological test.

- There are  $2^n$  possible inputs
- If you have 1000 unique exemplars:

n	coverage
10	100.0%
50	40.0%
100	10.0%
500	0.4%
1000	0.1%



- Machine learning is built on similarity,  
the assumption that a similar feature vector implies a similar answer
- Low coverage  $\implies$  nothing is similar

- Machine learning is built on similarity,  
the assumption that a similar feature vector implies a similar answer
- Low coverage  $\implies$  nothing is similar
- There is still something “most similar” and “least similar” however...


## Distance in $100D$ Space

- 1000 points in  $nD$  space, uniform distribution,  $x \in [0, 1]^n$
- How far away is everything? (Euclidean)




## Distance in $100D$ Space

- 1000 points in  $nD$  space, uniform distribution,  $x \in [0, 1]^n$
- How far away is everything? (Euclidean)
- Right: Typical values  
Red = maximum  
Blue = mean  
Green = minimum



## Distance in $100D$ Space

- 1000 points in  $nD$  space, uniform distribution,  $x \in [0, 1]^n$
- How far away is everything? (Euclidean)
- Right: Typical values  
Red = maximum  
Blue = mean  
Green = minimum
- Distance increases, range does not.
- $\frac{\text{maximum}}{\text{minimum}}$  decreases.
- Everything starts to look the same!



## Never Enough

- Poor coverage = failing similarity
- However much data you have, the curse always gets you


## Never Enough

- Poor coverage = failing similarity
- However much data you have, the curse always gets you
- Consider nearest neighbours (set point to same class as nearest)
- First dimension,  $x[0]$ :
  - $x[0] < \frac{1/3}{\implies}$  Class 1,
  - $x[0] > \frac{2/3}{\implies}$  Class 2
  - (gap between empty)
- All further dimensions are noise,  
 $\sim \text{Normal}(0, 1)$

## Never Enough

- Poor coverage = failing similarity
- However much data you have, the curse always gets you
- Consider nearest neighbours (set point to same class as nearest)
- First dimension,  $x[0]$ :  
$$x[0] < \frac{1}{3} \implies \text{Class 1,}$$
$$x[0] > \frac{2}{3} \implies \text{Class 2}$$

(gap between empty)
- All further dimensions are noise,  
 $\sim \text{Normal}(0, 1)$
- This should be easy, but...




## Not the Distributions

- It also affects probability distributions  
e.g. the standard multivariate Gaussian


## Not the Distributions

- It also affects probability distributions e.g. the standard multivariate Gaussian
- Consider draws from the distribution
- Calculate the distribution of their distances from the mean (called a  $\chi$  (chi) distribution)



## Not the Distributions

- It also affects probability distributions e.g. the standard multivariate Gaussian
- Consider draws from the distribution
- Calculate the distribution of their distances from the mean (called a  $\chi$  (chi) distribution)
- In high dimensions most draws are in a thin shell, nowhere near the mean!





## Solutions

- Feature selection
- Feature engineering
- Problem doesn't exist for real data! (manifolds)

## Feature Selection

- Select a small number of features to use, ignore rest
- Three approaches:
  - **Model Selection:** Train model with many subsets of the features and select best
  - **Filtering:** Use an estimate of feature usefulness and filter accordingly
  - **Embedded:** An algorithm with feature selection built in
- Poor relative of structure learning – later lecture

## Feature Selection: Model Selection

- Clearly the best approach, but slow
- Impractical to search exhaustively

## Feature Selection: Model Selection

- Clearly the best approach, but slow
- Impractical to search exhaustively
- Typical approach:
  1. Start with empty set
  2. Consider  $n$  modifications: Adding (not in set) or removing (in set) each variable
  3. Select the best; usually with an extra cost on number of features used
  4. If best changed return to step 2, otherwise exit

## Feature Selection: Model Selection

- Clearly the best approach, but slow
- Impractical to search exhaustively
- Typical approach:
  1. Start with empty set
  2. Consider  $n$  modifications: Adding (not in set) or removing (in set) each variable
  3. Select the best; usually with an extra cost on number of features used
  4. If best changed return to step 2, otherwise exit
- **Stepwise regression:** The above with a regression model, e.g. Logistic regression
- Possibility that considering larger “moves” would find a better solution

## Feature Selection: Filtering

- Questionable approach, but fast:
  1. Calculate how useful every feature is
  2. Select set based on a threshold
  3. Train model

## Feature Selection: Filtering

- Questionable approach, but fast:
  1. Calculate how useful every feature is
  2. Select set based on a threshold
  3. Train model
- Example: Threshold absolute correlation

## Feature Selection: Filtering

- Questionable approach, but fast:
  1. Calculate how useful every feature is
  2. Select set based on a threshold
  3. Train model
- Example: Threshold absolute correlation
- “Usefulness” is dependent on other variables:
  - Two variables might contain the same information  $\implies$  including both is pointless
  - A variable might be useless on it's own, but really helpful with others

There are variants that consider such relationships



## Feature Selection: Embedded

- Good trade off of performance/speed
- Example: Random forests – Feature/split selection is feature filtering

## Feature Selection: Embedded

- Good trade off of performance/speed
- Example: Random forests – Feature/split selection is feature filtering
- Random forest can have an infinite number of features!
- Example:
  - Kinect person tracker labels pixels – left forearm, head, background. . .
  - Generates pairs of random offsets from current pixel
  - Splits based on relative depth between pixel at offsets
  - (perspective is factored in – offsets are in meters)

# Feature Engineering


- Use domain knowledge to design new features
- Small number, based on what you think will work  
(use data exploration to help)

# Feature Engineering

- Use domain knowledge to design new features
- Small number, based on what you think will work  
(use data exploration to help)
- Example: Radius trick to fit a circle with Logistic regression
- Invariance and equivariance

## Data in the Real World


Consider a video of a cube falling:



- How many dimensions does each frame have?

## Data in the Real World

Consider a video of a cube falling:






- How many dimensions does each frame have?
- Formally: 3145728:  $1024 \times 1024$  images, 3 channels per pixel.
- Actually: 6: 3 for position of cube, 3 for orientation of cube.  
(everything else is a (complex) function of these 6 parameters)
- **Real data is mostly low dimensional!**

# Manifolds

- A manifold is the low dimensional surface the data is actually on.

# Manifolds


- A manifold is the low dimensional surface the data is actually on.
- $2D$  manifolds embedded in  $3D$  space:





## Dimensionality Reduction I


- Assign a coordinate system to the manifold, use that instead:



- Often used for visualisation (you used PCA previously)

## Dimensionality Reduction II



- Doesn't work if you have loops however!



# Principal Component Analysis

- Invented in 1901 by Karl Pearson!  
“On Lines and Planes of Closest Fit to Systems of Points in Space”
- Really simple
- Great for visualisation and dimensionality reduction
- Only handles linear manifolds (e.g. a  $4D$  hyper-cube embedded in  $12D$  space)
- But still useful, even for non-linear manifolds

## Example



## One Feature

- Imagine we have to represent a  $nD$  data set using one feature!
- How do we preserve the most information?

## One Feature

- Imagine we have to represent a  $nD$  data set using one feature!
- How do we preserve the most information?
- Don't have to select one feature – can take a linear combination  
e.g.  $\text{keep} = 0.2x_0 - 0.3x_1 + 0.1x_2$

## One Feature

- Imagine we have to represent a  $nD$  data set using one feature!
- How do we preserve the most information?
- Don't have to select one feature – can take a linear combination  
e.g.  $\text{keep} = 0.2x_0 - 0.3x_1 + 0.1x_2$
- Scale of linear combination does not matter, so assume  
 $\text{keep} = \hat{d} \cdot x, |\hat{d}| = 1$

## One Feature

- Imagine we have to represent a  $nD$  data set using one feature!
- How do we preserve the most information?
- Don't have to select one feature – can take a linear combination  
e.g.  $\text{keep} = 0.2x_0 - 0.3x_1 + 0.1x_2$
- Scale of linear combination does not matter, so assume  
 $\text{keep} = \hat{d} \cdot x, |\hat{d}| = 1$
- The most information is preserved if we select the direction of maximum variance.



## Maximum Variance

- Whiteboard!
- Note: Always passes through mean
- Obtains first principal component

## More Principal Components

- Repeat process, for remaining information
- That is, each new principal component must be orthogonal to all previous
- Orthogonality means no shared information
- In practise, there is an analytic solution

- Subtract mean
- Calculate eigenvectors and eigenvalues of covariance matrix
- Energy is the sum of eigenvalues
- Keep the eigenvectors associated with the largest eigenvalues, to obtain 99.9% of energy (or another percentage, or  $2/3$  if for visualisation)
- Transform data with matrix of kept eigenvectors
- Can return to original space using transpose of kept eigenvectors

```
data -= data.mean(axis=1)[:,None]
covar = numpy.cov(data.T) # Symmetric

evals , evcs = numpy.linalg.eigh(covar)
# evals is ordered lowest to highest , with evcs[:,i] matching evals[i]

energy = numpy.cumsum(evals)
start = 0
while energy[start] < 0.001 * energy[-1]: # Keep 99.9% of energy
    start += 1

project = evcs[:,start:][:,:-1].T
projected = project.dot(data.T).T
```

## Further Manifold Approaches

- Iso-map – non-linear, much better than PCA for visualisation
- “Learning a Manifold as an Atlas” by N. Pitelis, C. Russell and L. Agapito – non-linear by segmenting manifold into sections where PCA (linear) can be applied

## Future Curse-avoiding Approaches

Approaches that will be covered latter:

- Learning features with convolutional NN – next term
- Using structure (conditional independence) – next lecture

## Summary

- The curse of dimensionality  
(beware your intuitions in high dimensions)
- Feature selection
- Feature engineering
- Manifolds & dimensionality reduction

## Further Reading

- “A Few Useful Things to Know about Machine Learning”, by P. Domingos,  
<https://homes.cs.washington.edu/~pedrod/papers/cacm12.pdf>