



# Lecture 16-17

## CM50264: Machine Learning 1

### Bayesian Linear Regression and Gaussian Process Regression

Probabilistic  
interpretation

Bayesian Linear  
regression

Bayesian non-linear  
regression

Gaussian process  
regression

Kwang In Kim

## Previously in machine learning...

### Regularized linear least-squares regression

Data:

$$D = \{(\mathbf{x}^1, y^1), \dots, (\mathbf{x}^N, y^N)\} \subset \mathbb{R}^n \times \mathbb{R}.$$

Linear regression function:

$$f(\mathbf{x}) = \mathbf{x}^\top \mathbf{w}.$$

- (Plain) linear least-squares regression minimizes

$$\mathcal{O}(\mathbf{w}) = \sum_{i=1}^N (\mathbf{w}^\top \mathbf{x}^i - y^i)^2$$

- Regularized linear least-squares regression minimizes

$$\mathcal{O}(\mathbf{w}) = \sum_{i=1}^N (\mathbf{w}^\top \mathbf{x}^i - y^i)^2 + \lambda \|\mathbf{w}\|^2,$$

for the regularization (hyper-)parameter  $\lambda \geq 0$ .

Probabilistic  
interpretation

Bayesian Linear  
regression

Bayesian non-linear  
regression

Gaussian process  
regression

# Regularized linear least-squares regression

Data matrix:  $\mathbf{X} = [\mathbf{x}^1, \dots, \mathbf{x}^N]$ ; Label vector:  $\mathbf{y} = [y^1, \dots, y^N]^\top$ .

- (Plain) linear least-squares regression minimizes

$$\mathcal{O}(\mathbf{w}) = \sum_{i=1}^N (\mathbf{w}^\top \mathbf{x}^i - y^i)^2 = \|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2.$$

The minimizer  $\mathbf{w}^*$  is obtained by solving a linear system

$$\mathbf{X}\mathbf{X}^\top \mathbf{w} = \mathbf{X}\mathbf{y}.$$

For high-dimensional problems ( $N < n$ ),  $\mathbf{X}\mathbf{X}^\top$  is rank deficient: Infinitely many solutions exist.

- Regularized linear least-squares regression minimizes for  $\lambda \geq 0$

$$\mathcal{O}(\mathbf{w}) = \sum_{i=1}^N (\mathbf{w}^\top \mathbf{x}^i - y^i)^2 + \lambda \|\mathbf{w}\|^2 = \|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2 + \lambda \|\mathbf{w}\|^2.$$

The minimizer  $\mathbf{w}^*$  is obtained by solving a linear system

$$(\mathbf{X}\mathbf{X}^\top + \lambda \mathbf{I})\mathbf{w} = \mathbf{X}\mathbf{y}.$$

For  $\lambda > 0$ ,  $\mathbf{X}\mathbf{X}^\top + \lambda \mathbf{I}$  is always full rank: A unique solution exists.

Probabilistic  
interpretation

Bayesian Linear  
regression

Bayesian non-linear  
regression

Gaussian process  
regression

## (Deterministic) linear regression summary



- Input: Data  $\{(\mathbf{x}^1, y^1), \dots, (\mathbf{x}^N, y^N)\} \subset \mathbb{R}^n \times \mathbb{R}$ ; regularization parameter  $\lambda \geq 0$ .
- Training:
  - Build the data matrix:  $\mathbf{X} = [\mathbf{x}^1, \dots, \mathbf{x}^N]$  and label vector  $\mathbf{y} = [y^1, \dots, y^N]$ ;
  - Solve a linear system to obtain  $\mathbf{w}^*$ :  $(\mathbf{X}\mathbf{X}^\top + \lambda\mathbf{I})\mathbf{w} = \mathbf{X}\mathbf{y}$ ;
- Testing:  $f(\mathbf{x}') = (\mathbf{w}^*)^\top \mathbf{x}'$ .

We will discuss probabilistic interpretations of (plain) linear regression and regularized linear regression algorithms (see 'L09 Regularisation & Model Types' slides).

- Input and output variables  $\mathbf{x}$  and  $y$  are **random variables**.
- Training data  $\mathcal{D} = \{(\mathbf{x}^N, y^N), \dots, (\mathbf{x}^N, y^N)\}$  is sampled from an **unknown** probability distribution  $p(\mathbf{x}, y)$ .
- There is an underlying ground-truth function  $f^*(\mathbf{x}) = y$ , but our observations (training data) are noisy:

$$y^i = f^*(\mathbf{x}^i) + \text{noise}^i, 1 \leq i \leq N.$$

- $\text{noise}^i$  is a random variable.



## Independent and identically distributed (i.i.d.) Gaussian noise model

Probabilistic  
interpretation

Bayesian Linear  
regression

Bayesian non-linear  
regression

Gaussian process  
regression

$$y^i = f(\mathbf{x}^i) + \text{noise}^i$$

- $\text{noise}^i$  represents the deviation between the observed label  $y^i$  and the prediction  $f(\mathbf{x}^i)$ : training error.
- $\text{noise}^i$  is **independent** of  $\text{noise}^j$  ( $i \neq j$ ).
- Distribution of  $\text{noise}^i$  is Gaussian  $\mathcal{N}(\mu, \sigma^2)$  with
  - mean  $\mu$  zero
  - variance  $\sigma^2$  **identical** across  $i$
- In linear regression:

$$\begin{aligned} y^i &= f(\mathbf{x}^i) + \text{noise} \\ &= \mathbf{w}^\top \mathbf{x}^i + \text{noise}. \end{aligned}$$

Why i.i.d. Gaussian noise model?

## Maximum likelihood (ML) estimation

Training data:  $\mathcal{D} = \{(\mathbf{x}^N, y^N), \dots, (\mathbf{x}^N, y^N)\}$ .

Data matrix:  $\mathbf{X} = [\mathbf{x}^1, \dots, \mathbf{x}^N]$ .

Label vector:  $\mathbf{y} = [y^1, \dots, y^N]^\top$ .

Our model:

$$y^i = \mathbf{w}^\top \mathbf{x}^i + \text{noise}$$

$$\text{noise} \sim \mathcal{N}(0, \sigma^2).$$

The maximum likelihood (ML) estimation chooses  $\mathbf{w}^*$  that maximizes the likelihood of  $\mathbf{w}$  given  $\mathcal{D}$ , the possibility of observing training data points  $\mathcal{D}$  given the hypothesized solution  $\mathbf{w}$ :

$$\begin{aligned} p(\mathbf{y}|\mathbf{X}, \mathbf{w}) &= \prod_{i=1}^N p(y^i|\mathbf{x}^i, \mathbf{w}) \text{ (i.i.d. noise)} \\ &= \prod_{i=1}^N \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(\mathbf{w}^\top \mathbf{x}^i - y^i)^2}{2\sigma^2}\right) \text{ (Gaussian noise)} \\ &= \frac{1}{\sqrt{(2\pi\sigma^2)^N}} \exp\left(-\frac{\|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2}{2\sigma^2}\right). \end{aligned}$$

Probabilistic  
interpretation

Bayesian Linear  
regression

Bayesian non-linear  
regression

Gaussian process  
regression

## Maximum likelihood (ML) estimation

Our likelihood model:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \frac{1}{\sqrt{(2\pi\sigma^2)^N}} \exp\left(-\frac{\|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2}{2\sigma^2}\right).$$

ML estimation maximizes  $p(\mathbf{y}|\mathbf{X}, \mathbf{w})$ :

$$\begin{aligned} \mathbf{w}^* &= \arg \max_{\mathbf{w} \in \mathbb{R}^n} \frac{1}{\sqrt{(2\pi\sigma^2)^N}} \exp\left(-\frac{\|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2}{2\sigma^2}\right) \\ &= \arg \max_{\mathbf{w} \in \mathbb{R}^n} \exp\left(-\frac{\|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2}{2\sigma^2}\right) \\ &= \arg \max_{\mathbf{w} \in \mathbb{R}^n} \left(-\frac{\|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2}{2\sigma^2}\right) \\ &= \arg \min_{\mathbf{w} \in \mathbb{R}^n} \frac{\|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2}{2\sigma^2} \\ &= \arg \min_{\mathbf{w} \in \mathbb{R}^n} \|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2 \end{aligned}$$

$$\Leftrightarrow \mathbf{X}\mathbf{X}^\top \mathbf{w}^* = \mathbf{X}\mathbf{y}.$$

ML under i.i.d. Gaussian noise is the same as least-squares regression.

Probabilistic  
interpretation

Bayesian Linear  
regression

Bayesian non-linear  
regression

Gaussian process  
regression



# Maximum a posteriori (MAP) estimation

In ML, we maximize the likelihood:

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} p(\mathbf{y}|\mathbf{X}, \mathbf{w}).$$

In MAP, we maximize the **posterior**:

$$\mathbf{w}^* = \arg \max_{\mathbf{w}} p(\mathbf{w}|\mathbf{X}, \mathbf{y}).$$



## Maximum a posteriori (MAP) estimation

Applying Bayes' rule,

$$\begin{aligned}\arg \max_{\mathbf{w}} p(\mathbf{w}|\mathbf{y}, \mathbf{X}) &= \arg \max_{\mathbf{w}} \frac{p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w}|\mathbf{X})}{p(\mathbf{y}|\mathbf{X})} \\ &= \arg \max_{\mathbf{w}} \frac{p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w})}{p(\mathbf{y}|\mathbf{X})} \\ &= \arg \max_{\mathbf{w}} p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w}).\end{aligned}$$

We know how to calculate  $p(\mathbf{y}|\mathbf{X}, \mathbf{w})$ . What about  $p(\mathbf{w})$ ?

This is where we apply our **a priori** knowledge of  $\mathbf{w}$ .

What priori knowledge?

Probabilistic  
interpretation

Bayesian Linear  
regression

Bayesian non-linear  
regression

Gaussian process  
regression

# Gaussian prior

Probabilistic  
interpretation

Bayesian Linear  
regression

Bayesian non-linear  
regression

Gaussian process  
regression

$$p(\mathbf{w}|\mathbf{X}, \mathbf{y}) \propto p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w})$$

If we assume that  $p(\mathbf{w})$  is a centered Gaussian  $\mathcal{N}(0, \mathbf{I})$ :

$$p(\mathbf{w}) = \frac{1}{\sqrt{(2\pi)^n}} \exp\left(-\frac{\|\mathbf{w}\|^2}{2}\right),$$

maximizing the posterior  $p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w})$  **biases the solution  $\mathbf{w}^*$  towards 0**:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w}) = \frac{1}{\sqrt{(2\pi\sigma^2)^N}} \exp\left(-\frac{\|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2}{2\sigma^2}\right) \frac{1}{\sqrt{(2\pi)^n}} \exp\left(-\frac{\|\mathbf{w}\|^2}{2}\right).$$

$$\arg \max_{\mathbf{w}} p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w}) = \arg \max_{\mathbf{w}} \exp\left(-\frac{\|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2}{2\sigma^2}\right) \exp\left(-\frac{\|\mathbf{w}\|^2}{2}\right)$$

$$= \arg \max_{\mathbf{w}} \exp\left(-\frac{\|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2}{2\sigma^2} - \frac{\|\mathbf{w}\|^2}{2}\right)$$

$$= \arg \max_{\mathbf{w}} \left(-\frac{\|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2}{2\sigma^2} - \frac{\|\mathbf{w}\|^2}{2}\right)$$

$$= \arg \min_{\mathbf{w}} \frac{\|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2}{2\sigma^2} + \frac{\|\mathbf{w}\|^2}{2}$$

$$= \arg \min_{\mathbf{w}} \|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2 + \sigma^2 \|\mathbf{w}\|^2.$$

$$p(\mathbf{w}|\mathbf{X}, \mathbf{y}) \propto p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w})$$

Assuming that  $p(\mathbf{w})$  is a Gaussian  $\mathcal{N}(0, \mathbf{I})$ ,  
maximizing  $p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w})$  biases the solution  $\mathbf{w}^*$  towards 0.

$$\begin{aligned}\mathbf{w}^* &= \arg \max_{\mathbf{w}} p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w}) \\ &= \arg \min_{\mathbf{w}} \|\mathbf{X}^\top \mathbf{w} - \mathbf{y}\|^2 + \sigma^2 \|\mathbf{w}\|^2 \\ &\vdots\end{aligned}$$

$$\Leftrightarrow (\mathbf{X}\mathbf{X}^\top + \sigma^2 \mathbf{I})\mathbf{w}^* = \mathbf{X}\mathbf{y}.$$

With a Gaussian prior and i.i.d. Gaussian noise model, MAP estimate becomes regularized least-squares solution with  $\sigma^2$  as the regularization hyper-parameter.

Why Gaussian prior?

- Input: Data  $\{(\mathbf{x}^1, y^1), \dots, (\mathbf{x}^N, y^N)\} \subset \mathbb{R}^n \times \mathbb{R}$ ; noise parameter  $\sigma^2 \geq 0$ .
- Training:
  - Build the data matrix:  $\mathbf{X} = [\mathbf{x}^1, \dots, \mathbf{x}^N]$  and label vector  $\mathbf{y} = [y^1, \dots, y^N]$ ;
  - Solve a linear system to obtain  $\mathbf{w}^*$ :  $(\mathbf{X}\mathbf{X}^\top + \sigma^2\mathbf{I})\mathbf{w} = \mathbf{X}\mathbf{y}$ ;
- Testing:  $f(\mathbf{x}') = (\mathbf{w}^*)^\top \mathbf{x}'$ .

We choose  $\mathbf{w}^*$  by maximizing  $p(\mathbf{w}|\mathbf{y}, \mathbf{X})$

and then,

apply the resulting linear regression function  $f$  to a new data point  $\mathbf{x}'$ :  $f(\mathbf{x}') = (\mathbf{w}^*)^\top \mathbf{x}'$ .

We don't really need to build  $f$  (equivalently  $\mathbf{w}^*$ ) explicitly as an **intermediate result**

if we just want to make a prediction  $y'$  for a given input  $\mathbf{x}'$  (or inputs):

We can maximize the **posterior** (or **predictive distribution**):<sup>1</sup>

$$p(y'|\mathbf{x}', \mathbf{y}, \mathbf{X}).$$

Rewrite this in plain text!

---

<sup>1</sup>Cf. the parameter posterior  $p(\mathbf{w}|\mathbf{y}, \mathbf{X})$ .

From a given joint distribution  $p(\mathbf{a}, \mathbf{b})$ ,  
one can calculate the **marginal** distribution  $p(\mathbf{a})$  by **integrating**  
**b out**:

$$p(\mathbf{a}) = \int p(\mathbf{a}, \mathbf{b}) d\mathbf{b}.$$

Similarly,

$$\begin{aligned} p(\mathbf{a}, \mathbf{c}) &= \int p(\mathbf{a}, \mathbf{b}, \mathbf{c}) d\mathbf{b} \\ &= \int p(\mathbf{a}|\mathbf{b}, \mathbf{c}) p(\mathbf{b}|\mathbf{c}) d\mathbf{b}. \end{aligned}$$



Applying the marginalization of  $\mathbf{w}$  to the predictive distribution

$$p(y'|\mathbf{x}', \mathbf{y}, \mathbf{X}) = \int p(y'|\mathbf{x}', \mathbf{w})p(\mathbf{w}|\mathbf{y}, \mathbf{X})d\mathbf{w}$$

and combining it with the parameter posterior and likelihood

$$p(\mathbf{w}|\mathbf{y}, \mathbf{X}) = \mathcal{N}((\mathbf{X}\mathbf{X}^\top + \sigma^2\mathbf{I})^{-1}\mathbf{X}\mathbf{y}, \sigma^2(\mathbf{X}\mathbf{X}^\top + \sigma^2\mathbf{I})^{-1}) \quad (1)$$

$$p(y'|\mathbf{x}', \mathbf{w}) = \mathcal{N}(\mathbf{w}^\top \mathbf{x}', \sigma^2),$$

we obtain

$$p(y'|\mathbf{x}', \mathbf{y}, \mathbf{X}) = \mathcal{N}(\mathbf{x}'^\top (\mathbf{X}\mathbf{X}^\top + \sigma^2\mathbf{I})^{-1}\mathbf{X}\mathbf{y}, \mathbf{x}'^\top \sigma^2(\mathbf{X}\mathbf{X}^\top + \sigma^2\mathbf{I})^{-1}\mathbf{x}'). \quad (2)$$

Prove Eqs. 1 and 2.

## Predictive distribution

$$p(y|\mathbf{x}', \mathbf{y}, \mathbf{X}) = \mathcal{N} \left( \mathbf{x}'^\top (\mathbf{X}\mathbf{X}^\top + \sigma^2 \mathbf{I})^{-1} \mathbf{X}\mathbf{y}, \mathbf{x}'^\top \sigma^2 (\mathbf{X}\mathbf{X}^\top + \sigma^2 \mathbf{I})^{-1} \mathbf{x}' \right) \quad (3)$$

$$p(y'|\mathbf{x}', \mathbf{y}, \mathbf{X}) = \int p(y'|\mathbf{x}', \mathbf{w}) p(\mathbf{w}|\mathbf{y}, \mathbf{X}) d\mathbf{w}. \quad (4)$$

- The output of Bayesian linear regression is a **probability distribution**: Gaussian for Gaussian prior + Gaussian noise.
- If we take the mean of this **predictive distribution**, the result is the same as the MAP solution:

$$\begin{aligned} \mathbf{x}'^\top (\mathbf{X}\mathbf{X}^\top + \sigma^2 \mathbf{I})^{-1} \mathbf{X}\mathbf{y} &= \mathbf{x}'^\top \mathbf{w}^*, \\ \mathbf{w}^* &= (\mathbf{X}\mathbf{X}^\top + \sigma^2 \mathbf{I})^{-1} \mathbf{X}\mathbf{y}. \end{aligned}$$

- Equation 3 represents the posterior  $p(y|\mathbf{x}', \mathbf{y}, \mathbf{X})$  without explicitly involving the parameter vector  $\mathbf{w}^*$ . This does not mean that our regression model

$$f(\mathbf{x}) = \mathbf{w}^\top \mathbf{x}$$

is removed: The posterior should be consistent with the marginalization rule (Eq. 4).

Mean = mode for Gaussian.

Probabilistic  
interpretation

Bayesian Linear  
regression

Bayesian non-linear  
regression

Gaussian process  
regression

$$p(y|\mathbf{x}', \mathbf{y}, \mathbf{X}) = \mathcal{N}(\mathbf{x}'^\top (\mathbf{X}\mathbf{X}^\top + \sigma^2 \mathbf{I})^{-1} \mathbf{X}\mathbf{y}, \mathbf{x}'^\top \sigma^2 (\mathbf{X}\mathbf{X}^\top + \sigma^2 \mathbf{I})^{-1} \mathbf{x}').$$

The prediction is a Gaussian distribution characterized by

predictive mean:  $\mathbf{x}'^\top (\mathbf{X}\mathbf{X}^\top + \sigma^2 \mathbf{I})^{-1} \mathbf{X}\mathbf{y}$

predictive variance:  $\mathbf{x}'^\top \sigma^2 (\mathbf{X}\mathbf{X}^\top + \sigma^2 \mathbf{I})^{-1} \mathbf{x}'$ .

The predictive variance represents how **confident** the prediction is

- Large variance  $\rightarrow$  low confidence.
- Under the i.i.d. Gaussian noise model, predictive variances are independent of training labels  $\mathbf{y}$  (and underlying  $f$ ). For other noise models, predictive variances might depend on  $\mathbf{y}$ .
- Similarly to the mean prediction, our predictive variance is limited by the model assumption:  
If  $\mathbf{x}' = 0$ , we have an absolutely confident prediction. Why is 0 special?

Probabilistic  
interpretation

Bayesian Linear  
regression

Bayesian non-linear  
regression

Gaussian process  
regression

## Demo

How can we choose the noise hyperparameter  $\sigma^2$  (or equivalently  $\lambda$ )?

When we were applying Bayes' rule

$$p(\mathbf{w}|\mathbf{y}, \mathbf{X}) = \frac{p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w})}{p(\mathbf{y}|\mathbf{X})},$$

we discarded  $p(\mathbf{y}|\mathbf{X})$  since it is independent of  $\mathbf{w}$ .

The **marginal likelihood**  $p(\mathbf{y}|\mathbf{X})$  is a function of  $\sigma^2$ :

$$p(\mathbf{y}|\mathbf{X}) = \int p(\mathbf{y}|\mathbf{X}, \mathbf{w})p(\mathbf{w})d\mathbf{w}.$$

This represents how well  $p(\mathbf{w})$  respects the observed data and it can be used as a criteria for optimizing  $\sigma^2$ .

- Input: Data  $\{(\mathbf{x}^1, y^1), \dots, (\mathbf{x}^N, y^N)\} \subset \mathbb{R}^n \times \mathbb{R}$ ;  
noise parameter  $\sigma^2 \geq 0$ .
- Construct the predictive distribution  $p(y'|\mathbf{x}', \mathbf{y}, \mathbf{X})$  for a  
given input  $\mathbf{x}'$ :

$$p(y|\mathbf{x}', \mathbf{y}, X) = \mathcal{N}(\mathbf{x}'^\top \mathbf{A} \mathbf{X} \mathbf{y}, \mathbf{x}'^\top \sigma^2 \mathbf{A} \mathbf{x}')$$
$$\mathbf{A} = (\mathbf{X} \mathbf{X}^\top + \sigma^2 \mathbf{I})^{-1}.$$

- No clear distinction of training and testing stages;  
 $\mathbf{A} \mathbf{X} \mathbf{y}$  could be pre-calculated.

Idea: map  $\mathbf{x}$  to a feature space  $\mathcal{F}$  using a **non-linear map**  $\phi$  and build a linear regressor in  $\mathcal{F}$ :

$$f(\mathbf{x}) = \mathbf{w}^\top \phi(\mathbf{x}).$$

$$\begin{aligned} p(y'|\mathbf{x}', \mathbf{y}, \mathbf{X}) &:= p(y'|\mathbf{x}', \mathbf{y}, \Phi) \\ &= \mathcal{N}\left(\phi(\mathbf{x}')^\top (\Phi\Phi^\top + \sigma^2\mathbf{I})^{-1} \Phi\mathbf{y}, \right. \\ &\quad \left. \phi(\mathbf{x}')^\top (\Phi\Phi^\top + \sigma^2\mathbf{I})^{-1} \sigma^2 \phi(\mathbf{x}')\right), \\ \Phi &= [\phi(\mathbf{x}^1), \dots, \phi(\mathbf{x}^N)]. \end{aligned}$$

$$\begin{aligned} p(y'|\mathbf{x}', \mathbf{y}, \Phi) \\ &= \mathcal{N}\left(\phi(\mathbf{x}')^\top (\Phi\Phi^\top + \sigma^2\mathbf{I})^{-1} \Phi\mathbf{y}, \phi(\mathbf{x}')^\top \sigma^2 (\Phi\Phi^\top + \mathbf{I})^{-1} \phi(\mathbf{x}')\right), \\ &= \mathcal{N}\left(\phi(\mathbf{x}')^\top \Phi (\Phi^\top \Phi + \sigma^2\mathbf{I})^{-1} \mathbf{y}, \right. \\ &\quad \left. \phi(\mathbf{x}')^\top \phi(\mathbf{x}') - \phi(\mathbf{x}')^\top \Phi (\Phi^\top \Phi + \sigma^2\mathbf{I})^{-1} \Phi^\top \phi(\mathbf{x}')\right), \end{aligned}$$

$$\Phi = [\phi(\mathbf{x}^1), \dots, \phi(\mathbf{x}^N)].$$

$\phi$  is always given in the inner-product form  $\phi(\mathbf{a})^\top \phi(\mathbf{b})$ .

see Sherman-Morrison-Woodbury formula (last slide) for the second equality.



Using a positive definite **kernel function**  $k(\mathbf{a}, \mathbf{b}) = \phi(\mathbf{a})^\top \phi(\mathbf{b})$ , we obtain

$$\begin{aligned} p(y' | \mathbf{x}', \mathbf{y}, \Phi) &= \mathcal{N} \left( \phi(\mathbf{x}')^\top \Phi (\Phi^\top \Phi + \sigma^2 \mathbf{I})^{-1} \mathbf{y}, \right. \\ &\quad \left. \phi(\mathbf{x}')^\top \phi(\mathbf{x}') - \phi(\mathbf{x}')^\top \Phi (\Phi^\top \Phi + \sigma^2 \mathbf{I})^{-1} \Phi^\top \phi(\mathbf{x}') \right), \\ &= \mathcal{N} \left( \mathbf{k}^\top (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{y}, k(\mathbf{x}', \mathbf{x}') - \mathbf{k}^\top (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{k} \right), \end{aligned}$$

$$\begin{aligned} \mathbf{k} &= [k(\mathbf{x}', \mathbf{x}^1), \dots, k(\mathbf{x}', \mathbf{x}^N)]^\top \\ [\mathbf{K}]_{i,j} &= k(\mathbf{x}^i, \mathbf{x}^j). \end{aligned}$$

## Two modes of Bayesian non-linear regression



Original: complexity  $\mathcal{O}(n^3)$  ( $n = \dim(\mathcal{F})$ )

$$p(y|\mathbf{x}', \mathbf{y}, \mathbf{X}) := p(y|\mathbf{x}', \mathbf{y}, \Phi) \\ = \mathcal{N}\left(\sigma^{-2}\phi(\mathbf{x}')^\top (\Phi\Phi^\top + \sigma^2\mathbf{I})^{-1}\Phi\mathbf{y}, \phi(\mathbf{x}')^\top \sigma^2(\Phi\Phi^\top + \sigma^2\mathbf{I})^{-1}\phi(\mathbf{x})\right).$$

Kernelized version: complexity  $\mathcal{O}(N^3)$  ( $N = \#$  data points)

$$p(y|\mathbf{x}', \mathbf{y}, \mathbf{X}) = \mathcal{N}\left(\mathbf{k}^\top (K + \sigma^2\mathbf{I})^{-1}\mathbf{y}, k(\mathbf{x}', \mathbf{x}') - \mathbf{k}^\top (K + \sigma^2\mathbf{I})^{-1}\mathbf{k}\right).$$

When  $n > N$ , kernelized version is preferable.

$$k(\mathbf{a}, \mathbf{b}) = \exp \left( -\frac{\|\mathbf{a} - \mathbf{b}\|^2}{\sigma_k^2} \right)$$

- Simple linear regression is powerful in high-dimensional spaces (see ‘L15 Regularized Regression and Support Vector Machines’ slides).
- For any positive definite kernel  $k$ , there is a (non-unique) feature map  $\phi : \mathbb{R}^n \rightarrow \mathcal{H}$  such that  $k(\mathbf{a}, \mathbf{b}) = \phi(\mathbf{a})^\top \phi(\mathbf{b})$ .
- For the Gaussian kernel, the dimensionality of the feature space  $\mathcal{H}$  is infinite.



## Bayesian nonlinear regression summary

- Input: Data  $\{(\mathbf{x}^1, y^1), \dots, (\mathbf{x}^N, y^N)\} \subset \mathbb{R}^n \times \mathbb{R}$ ; noise parameter  $\sigma^2 \geq 0$ .
- Construct the predictive distribution  $p(y'|\mathbf{x}', \mathbf{y}, \mathbf{X})$  for a given input  $\mathbf{x}'$ :

$$p(y'|\mathbf{x}', \mathbf{y}, \mathbf{X}) = \mathcal{N}\left(\mathbf{k}^\top (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{y}, \right. \\ \left. k(\mathbf{x}', \mathbf{x}') - \mathbf{k}^\top (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{k}\right).$$

- No clear distinction of training and testing stages;  $(\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{y}$  could be pre-calculated.

The deterministic part (i.e. predictive mean) of kernelized Bayesian nonlinear regression is also called **kernel ridge regression** and **regularization networks**.

# Demo

- Effect of varying hyper-parameters, noise variance  $\sigma^2$  and kernel parameter  $\sigma_k^2$ ?
- How do we select hyper-parameters?

The **marginal likelihood**  $p(\mathbf{y}|\mathbf{X})$  is a function of  $\sigma^2$  and  $\sigma_k^2$ :

$$p(\mathbf{y}|\mathbf{X}) := p(\mathbf{y}|\Phi) = \int p(\mathbf{y}|\Phi, \mathbf{w})p(\mathbf{w})d\mathbf{w}.$$

This represents how well  $p(\mathbf{w})$  respects the observed data and it can be used as a criteria for optimizing  $\sigma^2$  and  $\sigma_k^2$  (see ‘Calculating the marginal likelihood’ in the last slide).

This is not a truly Bayesian approach. How do we **choose** the hyper-parameters in a fully Bayesian way?

For Bayesian linear regression, we adopted the model assumption

$$f(\mathbf{x}) = (\mathbf{w}^*)^\top \mathbf{x}$$

and made predictions using the marginalization

$$p(y'|\mathbf{x}', \mathbf{y}, \mathbf{X}) = \int p(y'|\mathbf{x}', \mathbf{w})p(\mathbf{w}|\mathbf{y}, \mathbf{X})d\mathbf{w}.$$

The parameter posterior  $p(\mathbf{w}|\mathbf{y}, \mathbf{X})$  can be evaluated by combining the likelihood  $p(\mathbf{y}|\mathbf{X}, \mathbf{w})$  and the prior  $p(\mathbf{w}) = p(\mathbf{w}|\mathbf{X})$ .

We used a centered Gaussian  $\mathcal{N}(0, \mathbf{I})$  prior:

$$p(\mathbf{w}) = \frac{1}{\sqrt{(2\pi)^n}} \exp\left(-\frac{\|\mathbf{w}\|^2}{2}\right).$$

If we remove the model assumption  $f(\mathbf{x}) = (\mathbf{w}^*)^\top \mathbf{x}$  and use  $f$  as a variable,  
our prediction rule will look like

$$p(y'|\mathbf{x}', \mathbf{y}, \mathbf{X}) = \int p(y'|\mathbf{x}', f)p(f|\mathbf{y}, \mathbf{X})df.$$

The **function posterior**  $p(f|\mathbf{y}, \mathbf{X})$  depends on the likelihood  $p(\mathbf{y}|f, \mathbf{X})$  and the prior  $p(f|\mathbf{X})$ .

Now we need a Gaussian distribution  $p(f|\mathbf{X})$  on the space of functions.

## Gaussian random vectors

- A Gaussian random variable  $w$  follows a Gaussian distribution:

$$w \sim \mathcal{N}(\mu, \sigma^2)$$

$$p(w) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(w - \mu)^2}{2\sigma^2}\right)$$

with mean  $\mu$  and variance  $\sigma^2$ .

- A Gaussian random vector  $\mathbf{w} \in \mathbb{R}^n$  is a collection of random variables  $\{\mathbf{w}_j\}_{j=1}^n$  that has a joint Gaussian distribution

$$\mathbf{w} \sim \mathcal{N}(\boldsymbol{\mu}, \Sigma)$$

$$p(\mathbf{w}) = \frac{1}{\sqrt{(2\pi\sigma^2)^n |\Sigma|}} \exp\left(-\frac{1}{2}(\mathbf{w} - \boldsymbol{\mu})\Sigma^{-1}(\mathbf{w} - \boldsymbol{\mu})\right)$$

with mean vector  $\boldsymbol{\mu}$  and covariance matrix  $\Sigma$ .

Elements  $\{\mathbf{w}_j\}_{j=1}^n$  of  $\mathbf{w}$  is indexed by an integer  $j \in 1, \dots, n$ .

Probabilistic  
interpretation

Bayesian Linear  
regression

Bayesian non-linear  
regression

Gaussian process  
regression



A **Gaussian process** (GP)  $f$  is a collection of random variables, any finite subset of which has a joint Gaussian distribution.

- A GP is specified by a mean function  $m(\mathbf{x})$  and a covariance function  $k(\mathbf{x}, \mathbf{x}')$ :

$$m(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})]$$

$$k(\mathbf{x}, \mathbf{x}') = \mathbb{E}[(f(\mathbf{x}) - m(\mathbf{x}))(f(\mathbf{x}') - m(\mathbf{x}'))].$$

- The elements of a GP  $f$  is indexed by a continuous variable  $\mathbf{x} \in \mathbb{R}^n$ .
- For any set  $\mathbf{X} = \{\mathbf{x}^1, \dots, \mathbf{x}^N\}$ ,  
 $\mathbf{f}_{\mathbf{X}} =: \{f(\mathbf{x}^1), \dots, f(\mathbf{x}^N)\}$  is a Gaussian random vector characterized by  
mean vector  $\boldsymbol{\mu}_{\mathbf{X}} = [m(\mathbf{x}^1), \dots, m(\mathbf{x}^N)]^\top$   
covariance matrix  $\mathbf{K}_{\mathbf{X}} : [\mathbf{K}_{\mathbf{X}}]_{i,j} = k(\mathbf{x}^i, \mathbf{x}^j)$ .

- A GP is specified by a mean function  $m(\mathbf{x})$  and a covariance function  $k(\mathbf{x}, \mathbf{x}')$ :

$$m(\mathbf{x}) = \mathbb{E}[f(\mathbf{x})]$$

$$k(\mathbf{x}, \mathbf{x}') = \mathbb{E}[(f(\mathbf{x}) - m(\mathbf{x}))(f(\mathbf{x}') - m(\mathbf{x}'))].$$

- For any set  $\mathbf{X} = \{\mathbf{x}^1, \dots, \mathbf{x}^N\}$ ,  
 $\mathbf{f}_{\mathbf{X}} =: \{f(\mathbf{x}^1), \dots, f(\mathbf{x}^N)\}$  is a Gaussian random vector characterized by  
mean vector  $\boldsymbol{\mu}_{\mathbf{X}} = [m(\mathbf{x}^1), \dots, m(\mathbf{x}^N)]^\top$   
covariance matrix  $\mathbf{K}_{\mathbf{X}} : [\mathbf{K}_{\mathbf{X}}]_{i,j} = k(\mathbf{x}^i, \mathbf{x}^j)$ .
- A GP is the generalization of a Gaussian random vector to infinite-dimensional objects, e.g. functions:

$$f \sim \mathcal{GP}(m, k).$$

For  $\bar{\mathbf{X}} := \mathbb{R}^n \setminus \{\mathbf{X}, \mathbf{x}'\}$  (all inputs other than training and test inputs),

$$\begin{aligned} p(y' | \mathbf{x}', \mathbf{y}, \mathbf{X}) &= \int p(y' | \mathbf{x}', f) p(f | \mathbf{y}, \mathbf{X}) df \\ &= \int p(y' | \mathbf{x}', \mathbf{f}_{\mathbf{X}}, \mathbf{f}_{\bar{\mathbf{X}}}) p(\mathbf{f}_{\mathbf{X}}, \mathbf{f}_{\bar{\mathbf{X}}} | \mathbf{y}, \mathbf{X}) d\mathbf{f}_{\mathbf{X}} d\mathbf{f}_{\bar{\mathbf{X}}} \\ &= \int p(y' | \mathbf{x}', \mathbf{f}_{\mathbf{X}}) p(\mathbf{f}_{\mathbf{X}} | \mathbf{y}, \mathbf{X}) d\mathbf{f}_{\mathbf{X}}. \end{aligned}$$

The third equality is called the **marginalization property** of GPs, generalizing

$$p(\mathbf{a}) = \int p(\mathbf{a}, \mathbf{b}) d\mathbf{b}.$$

$$p(y'|\mathbf{x}', \mathbf{y}, \mathbf{X}) = \int p(y'|\mathbf{x}', \mathbf{f}_{\mathbf{x}})p(\mathbf{f}_{\mathbf{x}}|\mathbf{y}, \mathbf{X})d\mathbf{f}_{\mathbf{x}}$$
$$p(\mathbf{f}_{\mathbf{x}}|\mathbf{y}, \mathbf{X}) = \frac{p(\mathbf{y}|\mathbf{X}, \mathbf{f}_{\mathbf{x}})p(\mathbf{f}_{\mathbf{x}}|\mathbf{X})}{p(\mathbf{y}|\mathbf{X})}$$

- We place a zero mean GP prior on  $f$ :  
For any set  $\mathbf{X} = \{\mathbf{x}^1, \dots, \mathbf{x}^N\}$ ,  
 $\mathbf{f}_{\mathbf{X}}|\mathbf{X}$  is a Gaussian random vector:

$$p(\mathbf{f}_{\mathbf{x}}|\mathbf{X}) = \mathcal{N}(\mathbf{0}, \mathbf{K}_{\mathbf{X}}).$$

- We place a zero-mean i.i.d. Gaussian noise model on  $\mathbf{y}$ :

$$p(\mathbf{y}|\mathbf{X}, \mathbf{f}_{\mathbf{X}}) = \mathcal{N}(\mathbf{f}_{\mathbf{X}}, \sigma^2\mathbf{I}).$$

With the zero mean GP prior:

$$p(\mathbf{f}_\mathbf{x}|\mathbf{X}) = \mathcal{N}(\mathbf{0}, \mathbf{K}_\mathbf{X}).$$

and i.i.d. Gaussian noise model:

$$p(\mathbf{y}|\mathbf{X}, \mathbf{f}_\mathbf{X}) = \mathcal{N}(\mathbf{f}_\mathbf{X}, \sigma^2 \mathbf{I}),$$

the joint distribution of  $\mathbf{y}$  and  $f(\mathbf{x}')$  is obtained as

$$\begin{pmatrix} \mathbf{y} \\ f(\mathbf{x}') \end{pmatrix} \sim \mathcal{N} \left( \mathbf{0}, \begin{pmatrix} (\mathbf{K}_\mathbf{X} + \sigma^2 \mathbf{I}) & \mathbf{k} \\ \mathbf{k}^\top & k(\mathbf{x}', \mathbf{x}') \end{pmatrix} \right) \quad (5)$$

with  $\mathbf{k} = [k(\mathbf{x}', \mathbf{x}^1), \dots, k(\mathbf{x}', \mathbf{x}^N)]^\top$ .

Using the Gaussian conditioning formula (last slide), we obtain the posterior

$$p(y'|\mathbf{x}', \mathbf{y}, \mathbf{X}) = \mathcal{N} \left( \mathbf{k}^\top (\mathbf{K}_\mathbf{X} + \sigma^2 \mathbf{I})^{-1} \mathbf{y}, k(\mathbf{x}', \mathbf{x}') - \mathbf{k}^\top (\mathbf{K}_\mathbf{X} + \sigma^2 \mathbf{I})^{-1} \mathbf{k} \right).$$

How do we derive the joint distribution (Eq. 5)?

# Equivalence of GP regression and kernelized Bayesian nonlinear regression



Lo  ve's theorem [Ber]:

$k(\cdot, \cdot)$  is a covariance function of a GP  
 $\Leftrightarrow k(\cdot, \cdot)$  is a symmetric positive definite function (kernel).

Nonlinear feature map + linear Bayesian regression is the same as GP regression.

- Sherman-Morrison-Woodbury formula:

For matrices  $\mathbf{A} \in \mathbb{R}^{m \times m}$ ,  $\mathbf{U} \in \mathbb{R}^{m \times n}$ ,  $\mathbf{C} \in \mathbb{R}^{n \times n}$ ,  $\mathbf{V} \in \mathbb{R}^{n \times m}$ ,

$$(\mathbf{A} + \mathbf{UCV})^{-1} = \mathbf{A}^{-1} - \mathbf{A}^{-1}\mathbf{U}(\mathbf{C}^{-1} + \mathbf{VA}^{-1}\mathbf{U})^{-1}\mathbf{VA}^{-1}$$

$$(\mathbf{A} + \mathbf{UCV})^{-1}\mathbf{UC} = \mathbf{A}^{-1}\mathbf{U}(\mathbf{C}^{-1} + \mathbf{VA}^{-1}\mathbf{U})^{-1}.$$

When  $n < m$ ,  $(\mathbf{C}^{-1} + \mathbf{VA}^{-1}\mathbf{U})^{-1}$  is less costly to calculate than  $(\mathbf{A} + \mathbf{UCV})^{-1}$ .

- Conditioning of a joint Gaussian is a Gaussian:

$$p\left(\begin{bmatrix} \mathbf{a} \\ \mathbf{b} \end{bmatrix}\right) = \mathcal{N}\left(\mathbf{0}, \begin{bmatrix} \mathbf{A} & \mathbf{C}^\top \\ \mathbf{C} & \mathbf{B} \end{bmatrix}\right)$$
$$\Rightarrow p(\mathbf{a}|\mathbf{b}) = \mathcal{N}(\mathbf{C}^\top \mathbf{B}^{-1} \mathbf{b}, \mathbf{A} - \mathbf{C}^\top \mathbf{B}^{-1} \mathbf{C}).$$

- Calculating the marginal likelihood:

$$-2 \log p(\mathbf{y}|\mathbf{X}) = \mathbf{y}^\top (\mathbf{K} + \sigma^2 \mathbf{I})^{-1} \mathbf{y} + \log |\mathbf{K} + \sigma^2 \mathbf{I}| + N \log 2\pi.$$

- Ber** Berlinet and Thomas-Agnan, *Reproducing Kernel Hilbert Spaces in Probability And Statistics*, Kluwer Academic, 2004.
- Pet** Petersen and Pedersen, *The Matrix Cookbook*
- Teu** Teukolsky, Vetterling, Flannery, *Numerical Recipes: The Art of Scientific Computing*, Cambridge University Press (any edition)  
<http://www.nr.com/>