

— CM50248 — 2017/2018 —


Visual Understanding 1

Dr Christian Richardt

Recap


Additive and subtractive primary colours (recap)

RGB – additive



Mixing lights
e.g. LCD monitors

CMY – subtractive



Filters and paints
e.g. printers

Digital image formation (recap)

- **Rasterisation:** convert a continuous or vector image representation to a rectangular sampled grid of pixels
 - e.g. Bresenham line drawing
- Rasterising vector graphics
 - e.g. Bresenham line drawing
- Sampling an analogue signal
 - e.g. in a digital camera
 - Digital cameras have a CCD or CMOS sensor with light sensitive elements, typically arranged in a “Bayer” pattern






Image storage (greyscale, recap)

- Digital images are arranged in memory in scanline order:
 - pixels from left to right within a row
 - rows from top to bottom





Image storage (colour)


- Colour images are arranged in memory with rows of RGB pixels:




- In general, access pixel (x, y) at memory address $3 \times (y \times \text{width}[\text{in pixels}] + x)$
- Other possibilities:
 - indexed images, 4 values per pixel (RGBA: RGB and “alpha” for opacity)

Histograms

- Summarise distribution of intensities in an image



RGB histogram
(Photoshop)




intensity
histogram

© 2016 Cambridge in Colour

Linear filtering (convolution, recap)

$$(K * I)(x, y) = \sum_i \sum_j K(i, j) I(x - i, y - j)$$


- Slide a filter over the image to compute a new value for each pixel
 - Linearly combine pixel values within a window/neighbourhood
 - Weight of each neighbouring pixel defined by a **kernel**
- **Convolution:**
general filtering approach that supports arbitrary kernels
- (We'll look at speeding it up later)




Morphological filtering (recap)

- Filtering of images using set theory
- Using a **structuring element** (set of pixel offsets),
e.g. $B = \{(-2, -2), (-1, -1), (0, 0), (1, 1), (2, 2)\}$

Input image



Erosion



Dilation



OpenCV documentation

Opening



Closing



0	0	0	0	1
0	0	0	1	0
0	0	1	0	0
0	1	0	0	0
1	0	0	0	0

line

0	0	1	0	0
0	1	1	1	0
1	1	1	1	1
0	1	1	1	0
0	0	1	0	0

disc/diamond

1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

square

The Fourier Transform

Motivation

- Transformations are useful for analysing signals
 - Natural to analyse audio signals by decomposing into frequencies
 - Can also analyse images using frequencies in x - and y -directions
- **Fourier transform** represents signals in terms of their frequencies
- Applications:
 - Low- and high-pass filtering
 - Fast linear filtering using the convolution theorem
 - Removing structured noise
 - Image compression (JPEG)


Application: structured noise removal

- Discard regions in spectrum associated with patterned noise


Input image



Spectrum



Edited



Reconstructed image





Image difference

Images by Fred Weinhaus

Complex numbers

- Extending the 1D number line to 2D complex plane: $z = a + bi \in \mathbb{C}$
 - Real part $\operatorname{Re}(z) = a \in \mathbb{R}$, imaginary part $\operatorname{Im}(z) = b \in \mathbb{R}$
 - Imaginary unit: $i^2 = -1$
- Complex conjugate: $\bar{z} = \operatorname{Re}(z) - \operatorname{Im}(z) \cdot i = a - bi$
- Polar form: $z = r \cdot e^{i\varphi}$
 - Radius $r \in \mathbb{R}$, argument (angle) $\varphi \in \mathbb{R}$
 - Complex conjugate: $\bar{z} = r \cdot e^{-i\varphi}$
- Magnitude: $|z| = \sqrt{z\bar{z}} = \sqrt{(a + bi)(a - bi)}$
$$= \sqrt{a^2 - b^2 i^2} = \sqrt{a^2 + b^2}$$



Arithmetic with complex numbers

- Addition: $(a + bi) + (c + di) = (a + c) + (b + d) \cdot i$
- Subtraction: $(a + bi) - (c + di) = (a - c) + (b - d) \cdot i$
- Multiplication: $(a + bi) \cdot (c + di) = ac + adi + bci + bdi^2$
 $= (ac - bd) + (ad + bc) \cdot i$
 - In polar form: $(r_1 e^{i\varphi_1}) \cdot (r_2 e^{i\varphi_2}) = r_1 r_2 e^{i(\varphi_1 + \varphi_2)}$
- Division: $\frac{a + bi}{c + di} = \frac{(a + bi) \cdot (c - di)}{(c + di) \cdot (c - di)} = \frac{ac - adi + cbi - bdi^2}{c^2 - di^2}$
 - expand with with **complex conjugate** $= \frac{ac + bd}{c^2 + d^2} + \frac{bc - ad}{c^2 + d^2} \cdot i$

Euler's formula

$$e^{i\varphi} = \cos \varphi + i \sin \varphi$$

- Describes complex numbers on the unit circle:

$$\begin{aligned} |e^{i\varphi}| &= \sqrt{(\cos \varphi + i \sin \varphi)(\cos \varphi - i \sin \varphi)} \\ &= \sqrt{\cos^2 \varphi + \sin^2 \varphi} = 1 \end{aligned}$$

- Using $\sin(-\varphi) = -\sin \varphi$ and $\cos(-\varphi) = \cos \varphi$, we obtain:


$$\cos \varphi = \frac{e^{i\varphi} + e^{-i\varphi}}{2} \qquad \sin \varphi = \frac{e^{i\varphi} - e^{-i\varphi}}{2i}$$

Even and odd functions

- A function is
 - **even** if $f(-x) = f(x)$
 - **odd** if $f(-x) = -f(x)$

- Examples:

- even: $f(x) = \cos x$
- odd: $f(x) = \sin x$



- Properties:
 - If f, g are even, then fg is even
 - If f, g are odd, then fg is even
 - If f is even and g is odd, fg is odd
 - If g is odd, then $\int_{-a}^a g(x) dx = 0$ for $a > 0$
 - If g is even, then $\int_{-a}^a g(x) dx = 2 \int_0^a g(x) dx$

- For any function f :

$$f(x) = f_{\text{even}}(x) + f_{\text{odd}}(x)$$

$$f_{\text{even}}(x) = \frac{f(x) + f(-x)}{2}$$

$$f_{\text{odd}}(x) = \frac{f(x) - f(-x)}{2}$$

Fourier transform

- The **Fourier transform** of a function $f : \mathbb{R} \rightarrow \mathbb{R}$ is given by:

$$F(\omega) = \mathcal{F}[f](\omega) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} f(x) e^{-i\omega x} dx$$


- The **inverse Fourier transform** restores f from F :

$$f(x) = \mathcal{F}^{-1}[F](\omega) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} F(\omega) e^{i\omega x} d\omega$$

- Note the similarity between the Fourier transform and its inverse.
- We use $F(\omega)$ and $\mathcal{F}[f](\omega)$ interchangeably to refer to the representation of $f(x)$ in the frequency (or Fourier) domain.

Fourier transform example: box function

$$f(x) = \begin{cases} 1 & \text{if } -1 < x < 1 \\ 0 & \text{otherwise} \end{cases}$$




$$F(\omega) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^{\infty} f(x) e^{-i\omega x} dx$$

$$= \frac{1}{\sqrt{2\pi}} \int_{-1}^1 e^{-i\omega x} dx = \frac{1}{\sqrt{2\pi}} \left[\frac{e^{-i\omega x}}{-i\omega} \right]_{-1}^1$$

$$= \frac{e^{-i\omega} - e^{i\omega}}{-i\omega\sqrt{2\pi}} = \frac{-2i \sin(\omega)}{-i\omega\sqrt{2\pi}}$$

$$= \sqrt{\frac{2}{\pi}} \frac{\sin(\omega)}{\omega} = \sqrt{\frac{2}{\pi}} \text{sinc}(\omega)$$



Properties of the Fourier transform

- **Linearity:** for constants a and b , and functions f and g :

$$\mathcal{F}[af + bg](\omega) = a\mathcal{F}[f](\omega) + b\mathcal{F}[g](\omega)$$

- **Shifting:** for a constant a , if $g(x) = f(x - a)$ then:

$$\mathcal{F}[g](\omega) = e^{-ia\omega} \mathcal{F}[f](\omega)$$

- **Modulation:** for a constant a , if $g(x) = e^{iax} f(x - a)$ then:

$$\mathcal{F}[g](\omega) = \mathcal{F}[f](\omega - a)$$

- **Scaling:** for a constant a , if $g(x) = f(ax)$, then:

$$\mathcal{F}[g](\omega) = \frac{1}{|a|} \mathcal{F}[f]\left(\frac{\omega}{a}\right)$$


Properties of the Fourier transform

- If f is even and real-valued, then F is even and real-valued.
- If f is odd and real-valued, then F is odd and purely imaginary.
- Differentiation:
$$\mathcal{F}\left[\frac{df}{dx}\right](\omega) = i\omega\mathcal{F}[f](\omega)$$
$$\mathcal{F}\left[\frac{d^n f}{dx^n}\right](\omega) = (i\omega)^n\mathcal{F}[f](\omega)$$
- Taking the derivative in the spatial domain multiplies the Fourier transform with the frequency:
 - higher frequencies are amplified, such as noise
 - Lower frequencies are attenuated

Convolution visualised

- The convolution $f * g$ of two functions f and g is given by:

$$(f * g)(x) = \int_{-\infty}^{\infty} f(x - y)g(y) dy$$



Cmglee / CC BY-SA 3.0

Convolution theorem

- The convolution $f * g$ of two functions f and g is given by:

$$(f * g)(x) = \int_{-\infty}^{\infty} f(x - y)g(y) dy$$

- **Convolution theorem:** $\mathcal{F}[f * g](\omega) = \sqrt{2\pi}\mathcal{F}[f](\omega) \cdot \mathcal{F}[g](\omega)$
 - Convolution becomes multiplication in the Fourier domain
 - Afterwards, transform results back to spatial domain
- As we will see, this can save a lot of time for convolutions with large kernels
- Reciprocal convolution theorem: $\mathcal{F}[f \cdot g](\omega) = \sqrt{2\pi}\mathcal{F}[f](\omega) * \mathcal{F}[g](\omega)$

Discrete Fourier transform (DFT)

- Discrete analogue to the continuous Fourier transform
 - deals with finite sampled signals, such as audio, images
 - N values decomposed into N frequency components
- **Discrete Fourier transform (DFT):**

$$F[k] = \sum_{n=0}^{N-1} f[n] e^{-2\pi i n k / N} \quad (k = 0, 1, \dots, N-1)$$

- **Inverse discrete Fourier transform:**

$$f[n] = \frac{1}{N} \sum_{k=0}^{N-1} F[k] e^{2\pi i n k / N} \quad (n = 0, 1, \dots, N-1)$$

Properties of the discrete Fourier transform

- Periodicity: $F[k]$ is periodic with period N since

$$F[k+N] = \sum_{n=0}^{N-1} f[n] e^{-2\pi i n (k+N)/N} = \sum_{n=0}^{N-1} f[n] e^{-2\pi i n k/N} = F[k]$$
$$\left(e^{-2\pi i n (k+N)/N} = e^{-2\pi i n k/N} e^{-2\pi i n} = e^{-2\pi i n k/N} \right)$$

- Cyclical convolution: $(f * g)[n] = \sum_{m=0}^{N-1} f[m] g[n-m]$
- **Convolution theorem:** the DFT of $f * g$ is $F[k] \cdot G[k]$

Fast Fourier transform (FFT)

- Naïve computation of DFT is expensive: $\mathcal{O}(N^2)$ complexity

- Compute N Fourier coefficients:

- N (complex-valued) multiplications

- $N-1$ (complex-valued) additions

$$\begin{aligned} F[k] &= \sum_{n=0}^{N-1} f[n] e^{-2\pi i n k / N} \\ &= f[0] + f[1] e^{-2\pi i k / N} + \dots + f[N-1] e^{-2\pi i k (N-1) / N} \end{aligned}$$

- For example, an image with 1 megapixel resolution ($N = 10^6$) would require $\mathcal{O}(N^2) = \mathcal{O}(10^{12})$ arithmetic operations

- Fast Fourier Transform (FFT) has $\mathcal{O}(N \log_2 N)$ complexity

[Gauss 1805, Cooley & Tukey 1965]

- much faster than $\mathcal{O}(N^2)$

- e.g. $\mathcal{O}(10^6 \log_2 10^6) \approx \mathcal{O}(2 \cdot 10^7)$ for $N = 10^6$

Fast Fourier transform (FFT)

- Key insight: divide and conquer
 - Split an input of size N into two inputs of size $N/2$ each
 - Repeat until inputs of size 1 is reached
 - This is has $\mathcal{O}(\log N)$ complexity: need this many levels of nesting
- Advantages:
 - Very efficient: $\mathcal{O}(N \log_2 N)$
 - Implemented in many libraries (e.g. FFTW used by MATLAB)
- Disadvantages:
 - Requires inputs of size $N = 2^k$, so input may need padding

Discrete Fourier transform (DFT) in 2D

- Fourier techniques can be generalised to n-D functions, e.g. for 2D:

- Discrete Fourier transform in 2D:

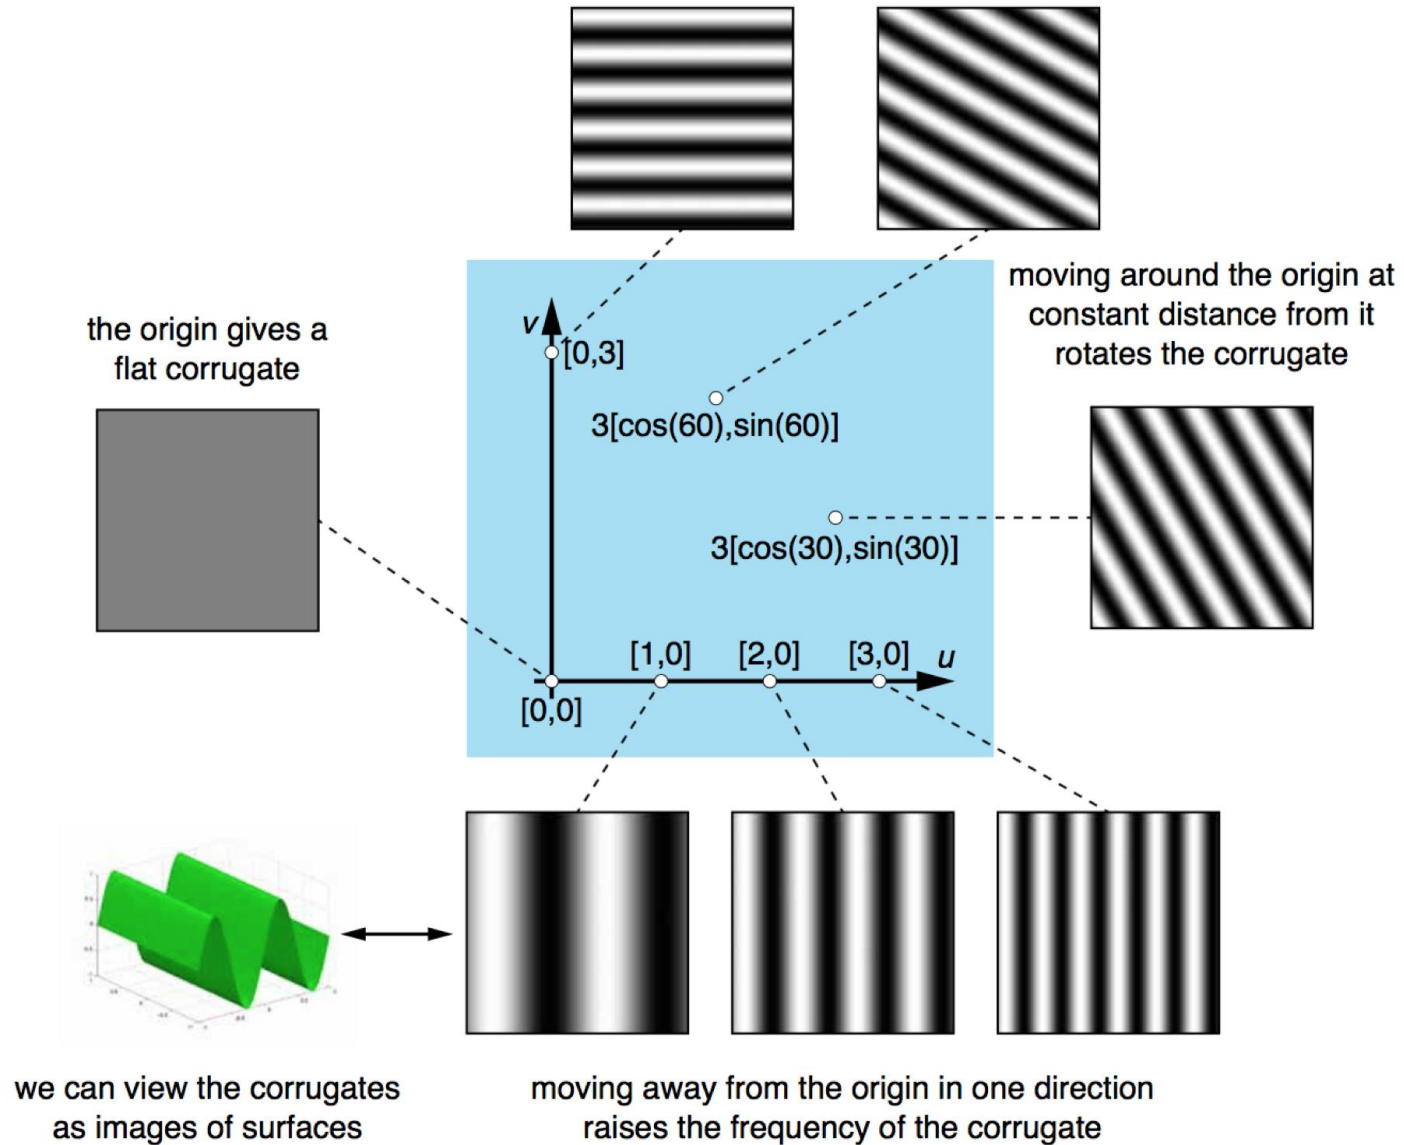
$$F[k, j] = \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} f[n, m] e^{-2\pi i(kn/N + jm/M)} \quad (j = 0, 1, \dots, M-1; k = 0, 1, \dots, N-1)$$

- Discrete inverse Fourier transform in 2D:

$$f[n, m] = \frac{1}{MN} \sum_{m=0}^{M-1} \sum_{n=0}^{N-1} F[k, j] e^{2\pi i(kn/N + jm/M)} \quad (n = 0, 1, \dots, N-1; m = 0, 1, \dots, M-1)$$


- Colour images: transform each channel independently

2D sinusoids



Low-pass filter



- passes signals with frequencies lower than a threshold
 - Removes fine details and hence blurs an image
- Example using hard cut-off in frequency space:
 - Ringing artefacts in image space
- Better low-pass filter: Gaussian blur
 - Removes ringing artefacts





High-pass filter

- passes signals with frequencies higher than a threshold
 - Removes low frequencies (~overall shape)
 - get edge image
- Example using hard cut-off in frequency space:
 - Ringing artefacts in image space
- Better high-pass filter: 1-Gaussian blur
 - Removes ringing artefacts

Input image



F



Images by Fred Weinhaus


Structured noise removal

- Discard regions in spectrum associated with patterned noise


Input image



Spectrum



Edited



Reconstructed image




Image difference



Images by Fred Weinhaus

Discrete cosine transform (DCT)

- Recall that the Fourier transform of an even, real-valued signal is even and real-valued (no imaginary part)
 - Audio signals and images are real-valued
 - Can make them even by reflection around $n=0$ or even $n=-0.5$
 - This cancels out the purely imaginary, sine-related terms
- DCTs are equivalent to DFTs of roughly twice the length
- input and/or output data shifted by half a sample in some variants
- Used for audio compression in MP3, image compression in JPEG

Application of the DCT in JPEG

- Discrete cosine transform is applied to 8×8 pixel blocks
- DCT coefficients are quantised (rounded) depending on frequency



Hanakis / Wikipedia

$$\text{FYI: } F[k_1, k_2] = \sum_{n_1=0}^{N_1-1} \sum_{n_2=0}^{N_2-1} x[n_1, n_2] \cos \left[\frac{\pi}{N_1} \left(n_1 + \frac{1}{2} \right) k_1 \right] \cos \left[\frac{\pi}{N_2} \left(n_2 + \frac{1}{2} \right) k_2 \right]$$

Questions?