

— CM50248 — 2017/2018 —

Visual Understanding 1

Dr Christian Richardt

Today's lecture

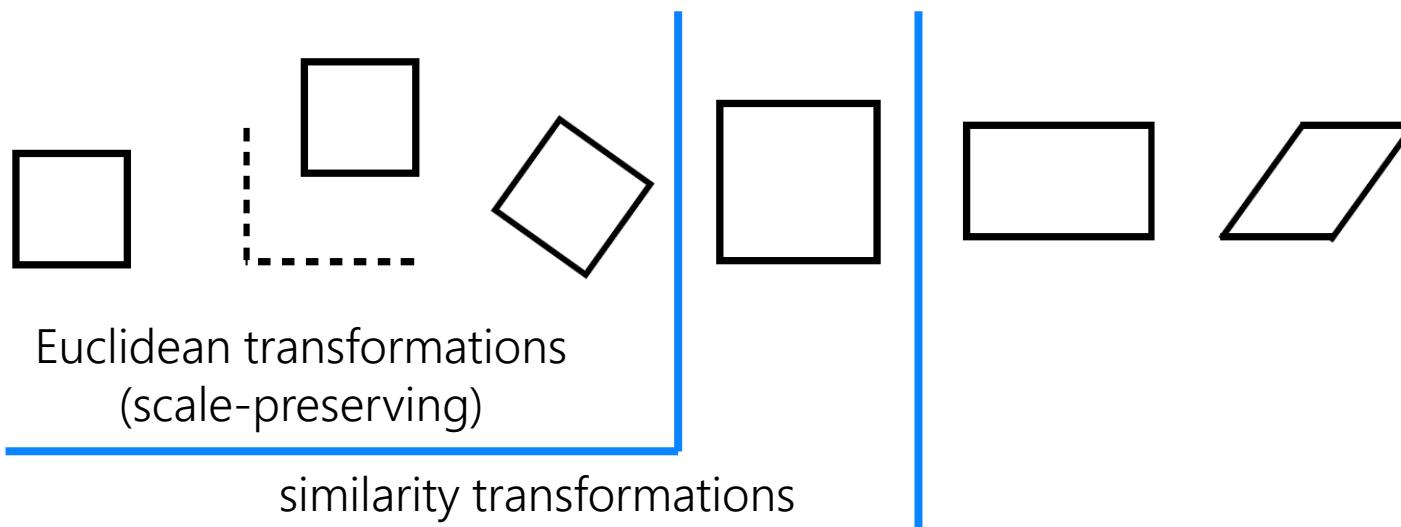
- Recap:
 - 2D transformations
 - RANSAC
 - Image warping, interpolation & stitching
- Camera geometry
- Camera calibration
- Epipolar geometry
- Fundamental matrix estimation
- Triangulation of 3D points
- Recommended reading:
 - CM50248 course notes:
 - 4.1 Projection
 - 4.2 Epipolar geometry
 - 4.3.2 Reconstruction
 - CM20219 course notes:
 - 4.3 Projection – 3D on a 2D display

Affine transformations

- What range of transformations can we represent using this?

$$\begin{bmatrix} p'_1 \\ p'_2 \\ 1 \end{bmatrix} = \begin{bmatrix} a_{11} & a_{12} & t_1 \\ a_{21} & a_{22} & t_2 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} p_1 \\ p_2 \\ 1 \end{bmatrix}$$

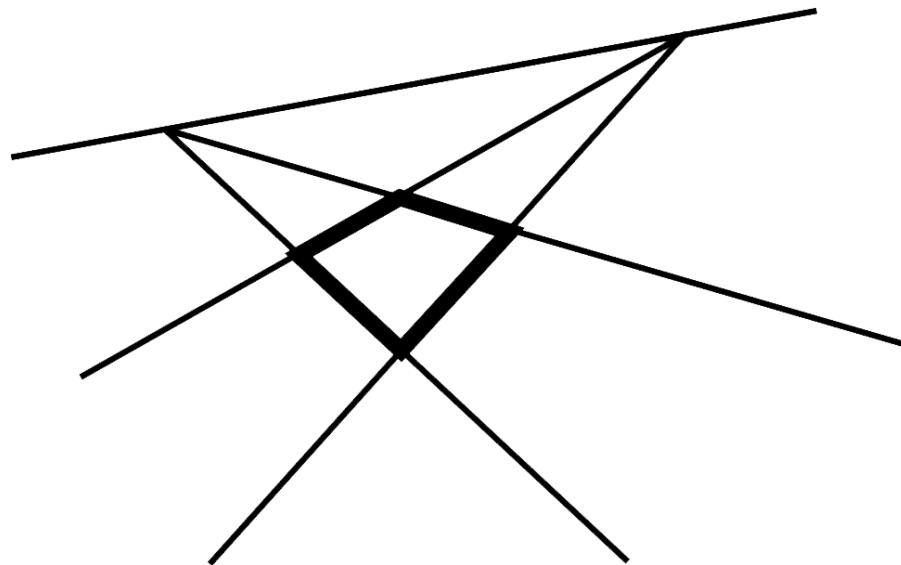
- 6 degrees of freedom: translation (2), rotation, scale, stretch, shear



Homographies

- Linear transform in homogeneous coordinates
- Represented by 3x3 matrix:

$$\mathbf{H} = \begin{bmatrix} h_{11} & h_{12} & h_{13} \\ h_{21} & h_{22} & h_{23} \\ h_{31} & h_{32} & h_{33} \end{bmatrix}$$



- 8 degrees of freedom (scale of \mathbf{H} arbitrary)
- 4 points can map to any 4 points
- Includes translation (2), rotation, scale, stretch, shear

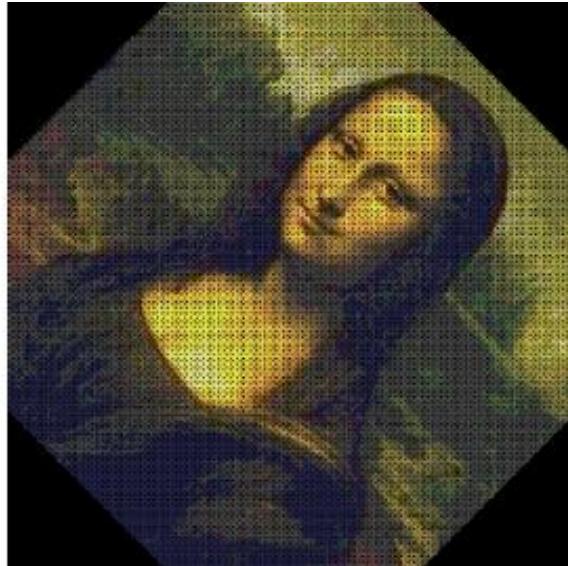
RANSAC for homography estimation

RANSAC loop:

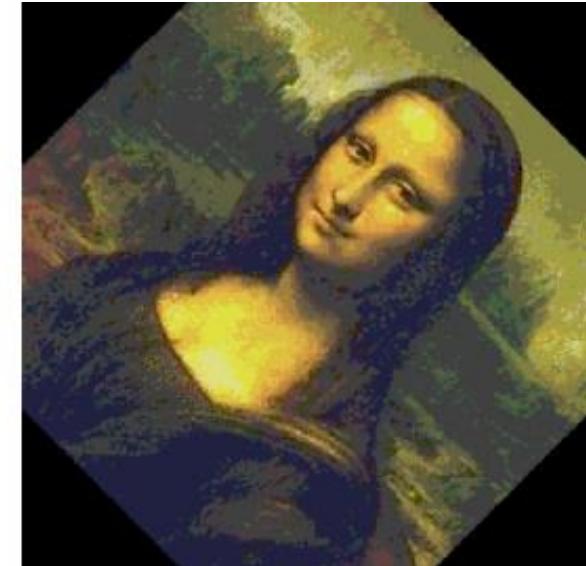
1. Select four feature pairs (at random)
2. Compute homography H (exact)
3. Compute *inliers* where $\|p' - Hp\| < \varepsilon$
4. Keep largest set of inliers
5. Re-compute least-squares H estimate
on all of the inliers

Image warping

- How do we warp images given a known affine transform?



Forward mapping



Backward mapping

$$\tilde{\mathbf{p}}' = \mathbf{T}\tilde{\mathbf{p}}$$

$$\tilde{\mathbf{p}} = \mathbf{T}^{-1}\tilde{\mathbf{p}}'$$

(if we round coordinates, image will contain gaps!)

[CM20219 Notes, Section 4.5]

Bundle adjustment

New images initialised with rotation,
focal length of best matching image



Bundle adjustment

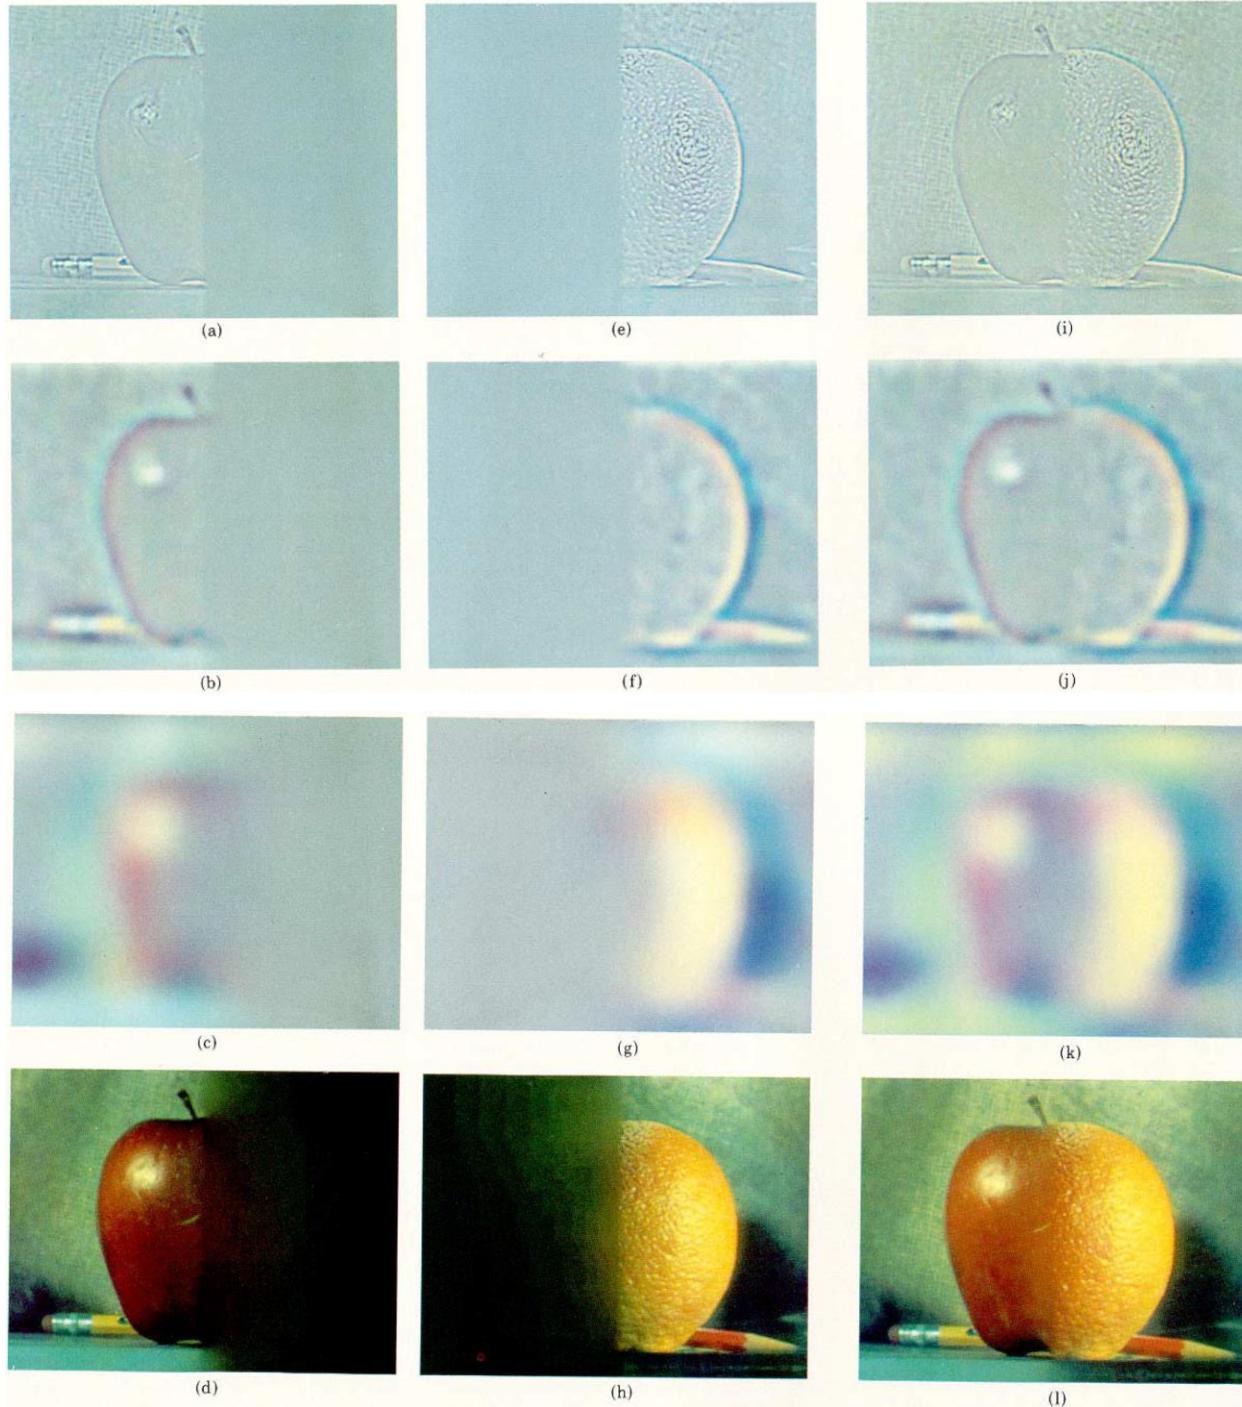
New images initialised with rotation,
focal length of best matching image



Panorama stitching

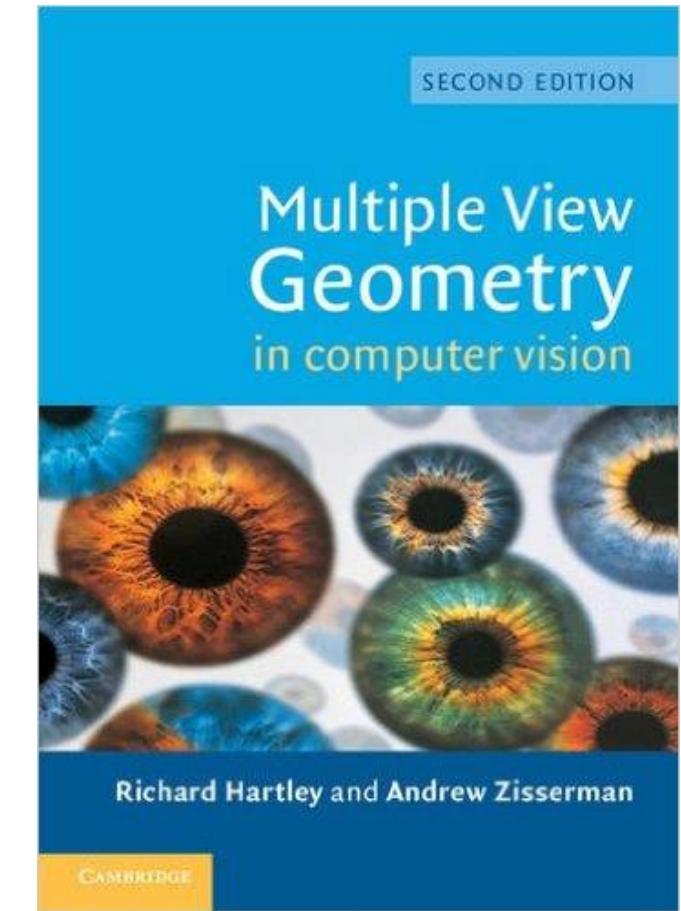
- Multi-band blending [Burt & Adelson 1983]
 - Blend frequency bands over different scales





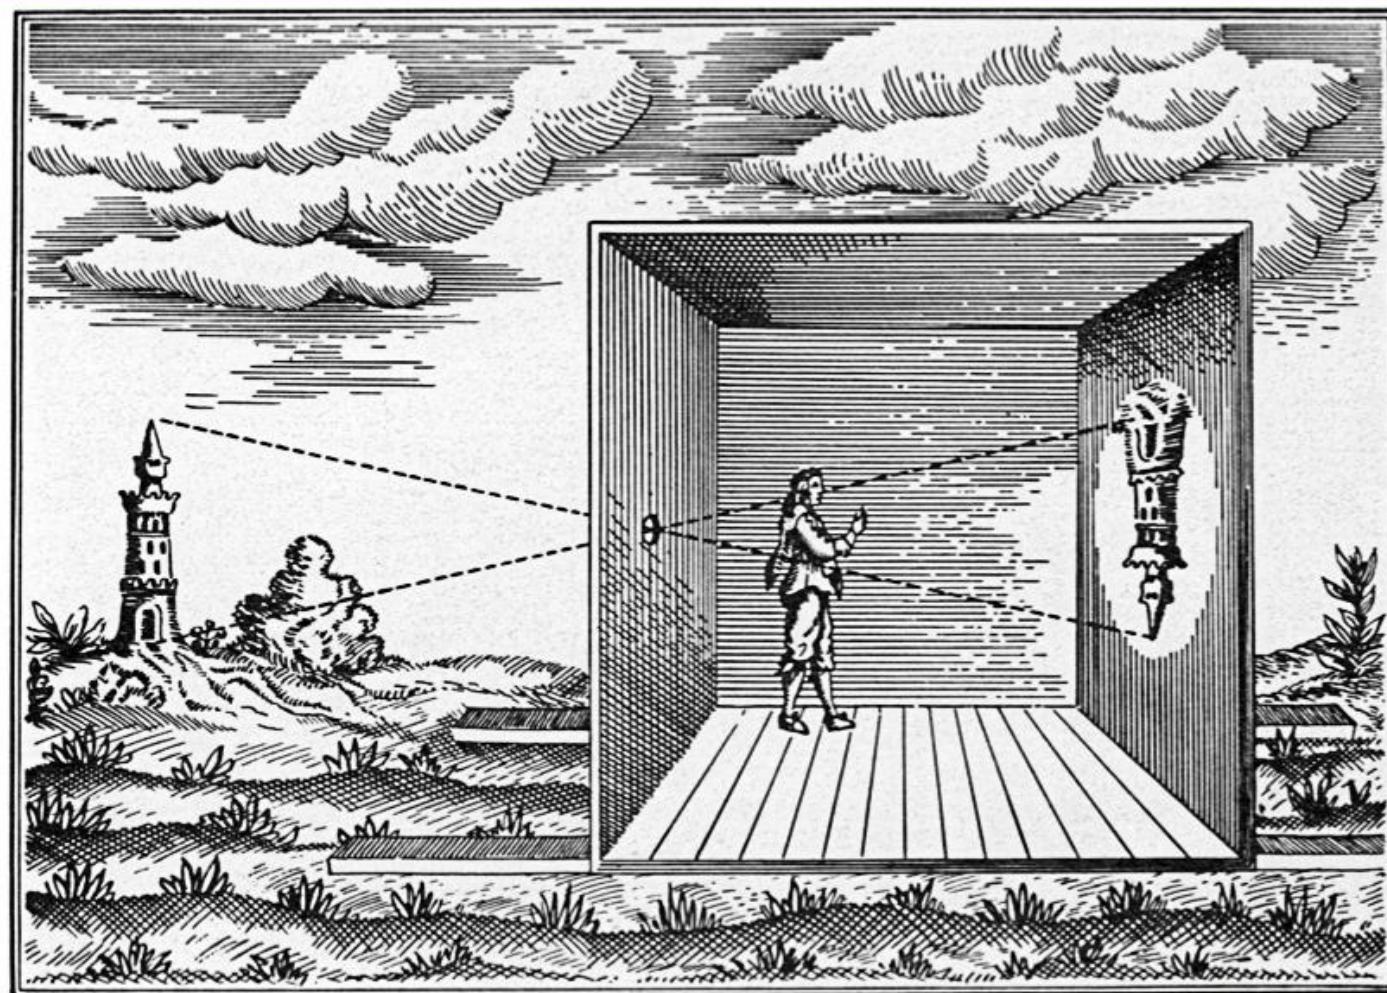
Burt & Adelson [ACM ToG 1983]

Camera geometry

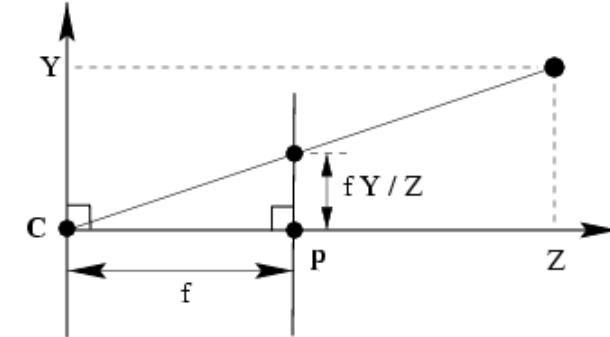
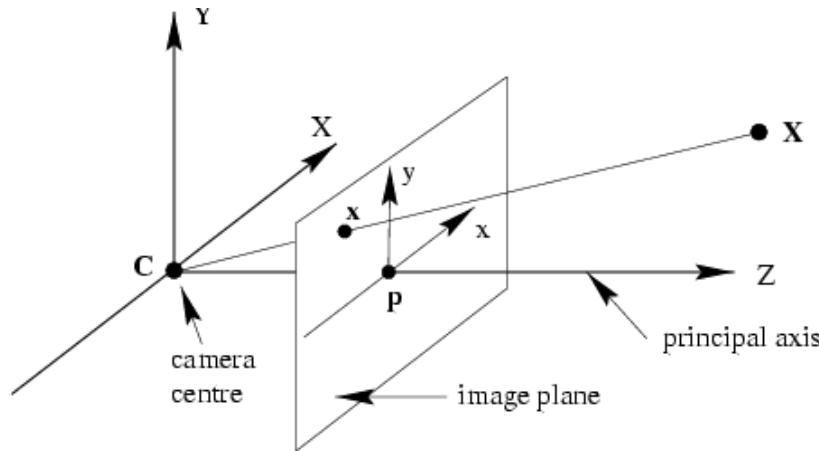


**Multiple View Geometry
in Computer Vision**
Hartley & Zisserman, 2004

Camera obscura



Pinhole camera model



$$(X, Y, Z)^T \mapsto (fX/Z, fY/Z)^T$$

$$\begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix} \mapsto \begin{pmatrix} fX \\ fY \end{pmatrix} = \begin{bmatrix} f & 0 & 0 & 1 \\ 0 & f & 0 & 0 \\ 0 & 0 & 1 & 0 \end{bmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix}$$

$$\begin{pmatrix} fX \\ fY \\ Z \end{pmatrix} = \begin{bmatrix} f & & & \\ & f & & \\ & & 1 & \end{bmatrix} \begin{pmatrix} 1 & & & \\ & 1 & & \\ & & 1 & \\ & & & 1 \end{pmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix}$$

$$P = \text{diag}(f, f, 1)[I | 0]$$

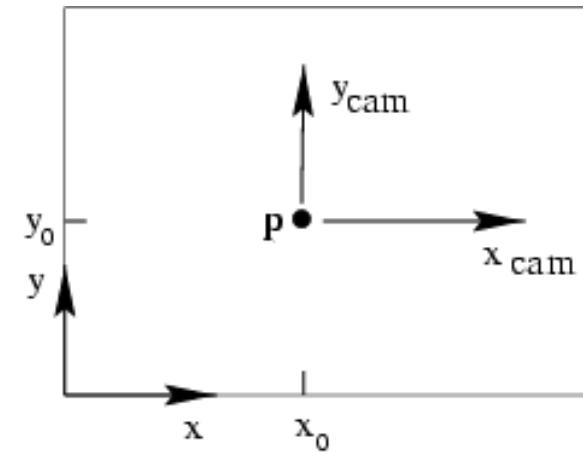
$$x = PX$$

Principal point offset

$$(X, Y, Z)^T \mapsto (fX/Z + p_x, fY/Z + p_y)^T$$

$(p_x, p_y)^T$ principal point

$$\begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix} \mapsto \begin{pmatrix} fX + Zp_x \\ fY + Zp_y \\ Z \end{pmatrix} = \begin{bmatrix} f & p_x & 0 \\ f & p_y & 0 \\ 1 & 0 & 0 \end{bmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix}$$

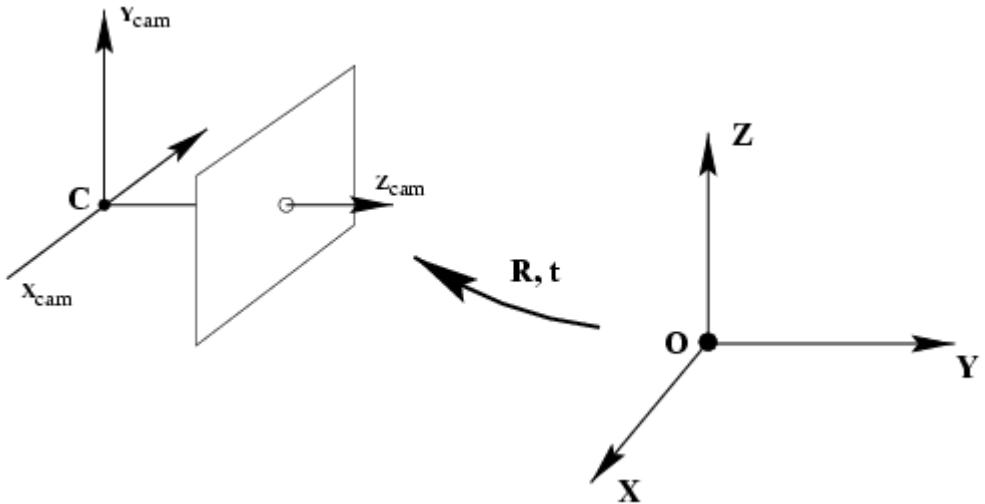


$$x = K[I | 0]X_{\text{cam}}$$

$$K = \begin{bmatrix} f & p_x \\ f & p_y \\ 1 & 0 \end{bmatrix}$$

calibration matrix

Camera rotation and translation



$$\tilde{X}_{\text{cam}} = R(\tilde{X} - \tilde{C})$$

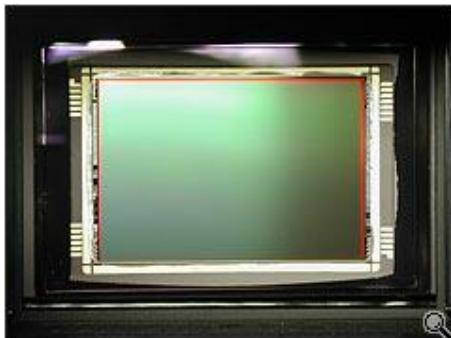
$$X_{\text{cam}} = \begin{bmatrix} R & -R\tilde{C} \\ 0 & 1 \end{bmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix} = \begin{bmatrix} R & -RC \\ 0 & 1 \end{bmatrix} X$$

$$x = K[I | 0]X_{\text{cam}}$$

$$x = KR[I | -\tilde{C}]X$$

$$x = PX \quad P = K[R | t] \quad t = -R\tilde{C}$$

Pixel coordinates



Pixel size: $\frac{1}{m_x} \times \frac{1}{m_y}$

- m_x pixels per meter in horizontal direction,
 m_y pixels per meter in vertical direction

$$K = \begin{bmatrix} m_x & & \\ & m_y & \\ & & 1 \end{bmatrix} \begin{bmatrix} f & p_x \\ & f & p_y \\ & & 1 \end{bmatrix} = \begin{bmatrix} \alpha_x & \beta_x \\ \alpha_y & \beta_y \\ & 1 \end{bmatrix}$$

pixels/m m pixels

Finite projective camera

$$K = \begin{bmatrix} \alpha_x & s & p_x \\ & \alpha_x & p_y \\ & & 1 \end{bmatrix}$$

$$P = \underbrace{KR}_{\text{non-singular}} [I \mid -\tilde{C}] \quad 11 \text{ dof (5+3+3)}$$

non-singular

decompose P in K,R,C?

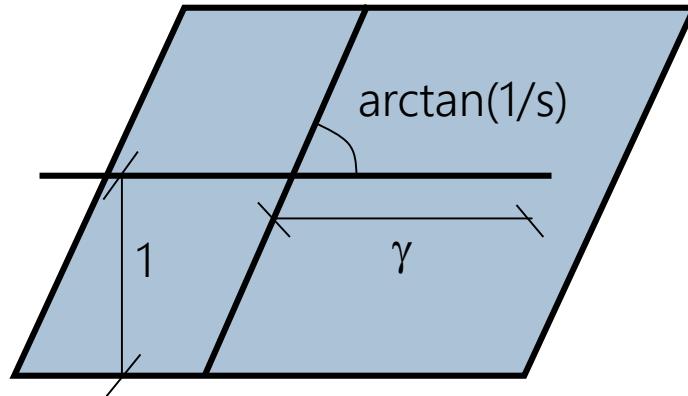
$$P = [M \mid p_4] \quad [K, R] = RQ(M) \quad \tilde{C} = -M^{-1}p_4$$

$$\{\text{finite cameras}\} = \{P_{4 \times 3} \mid \det M \neq 0\}$$

If rank P=3, but rank M<3, then cam at infinity

When is skew non-zero?

$$K = \begin{bmatrix} \alpha_x & s & p_x \\ & \alpha_x & p_y \\ & & 1 \end{bmatrix}$$



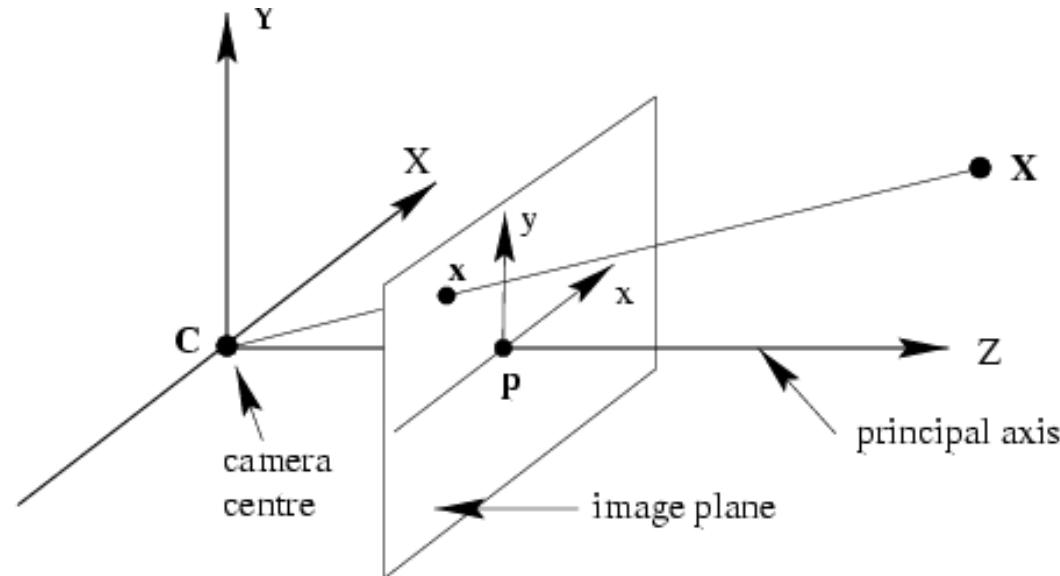
for CCD/CMOS, always $s=0$

Image from image, $s \neq 0$ possible
(non coinciding principal axis)

resulting camera: HP

Camera anatomy

- Camera center
- Column points
- Principal plane
- Axis plane
- Principal point
- Principal ray



Camera centre

null-space camera projection matrix

$$\mathbf{P}\mathbf{C} = \mathbf{0}$$

$$\mathbf{X} = \lambda \mathbf{A} + (1 - \lambda) \mathbf{C}$$

$$\mathbf{x} = \mathbf{P}\mathbf{X} = \lambda \mathbf{P}\mathbf{A} + (1 - \lambda) \mathbf{P}\mathbf{C}$$

For all \mathbf{A} , all points on \mathbf{AC} project on image of \mathbf{A} ,
therefore \mathbf{C} is camera centre

Image of camera center is $(0,0,0)^T$, i.e. undefined

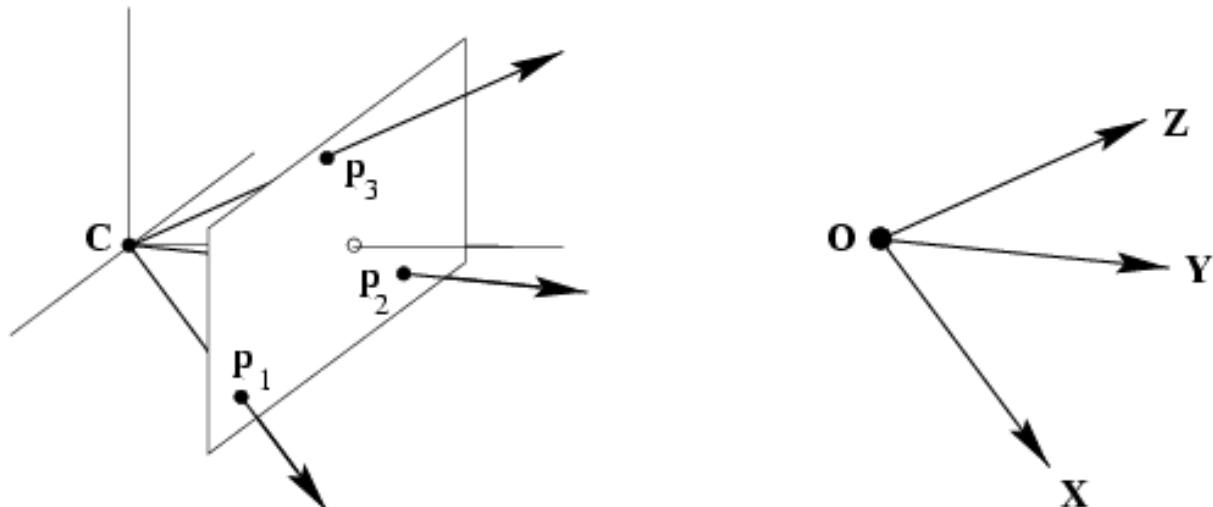
Finite cameras: $\mathbf{C} = \begin{pmatrix} -\mathbf{M}^{-1}\mathbf{p}_4 \\ 1 \end{pmatrix}$

Infinite cameras: $\mathbf{C} = \begin{pmatrix} \mathbf{d} \\ 0 \end{pmatrix}, \mathbf{M}\mathbf{d} = 0$

Column vectors of P

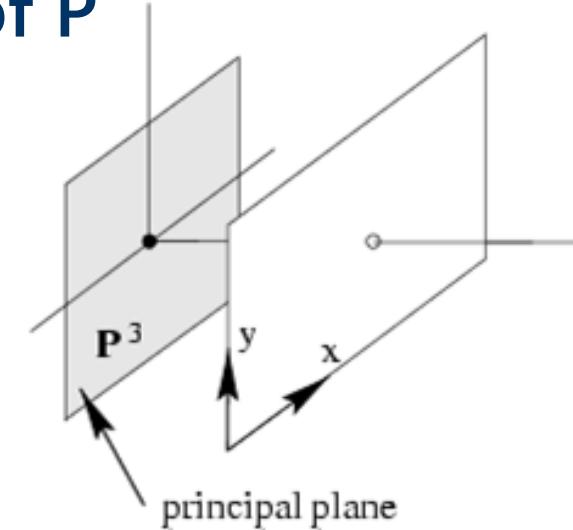
$$[p_2] = [p_1 p_2 p_3 p_4] \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

Image points
corresponding to x-, y-, z-
directions and origin

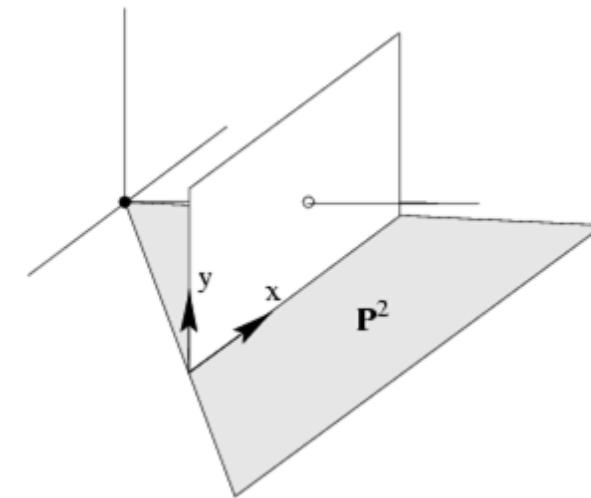


Row vectors of P

$$\begin{bmatrix} x \\ y \\ 0 \end{bmatrix} = \begin{bmatrix} p_1^T \\ p_2^T \\ p_3^T \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

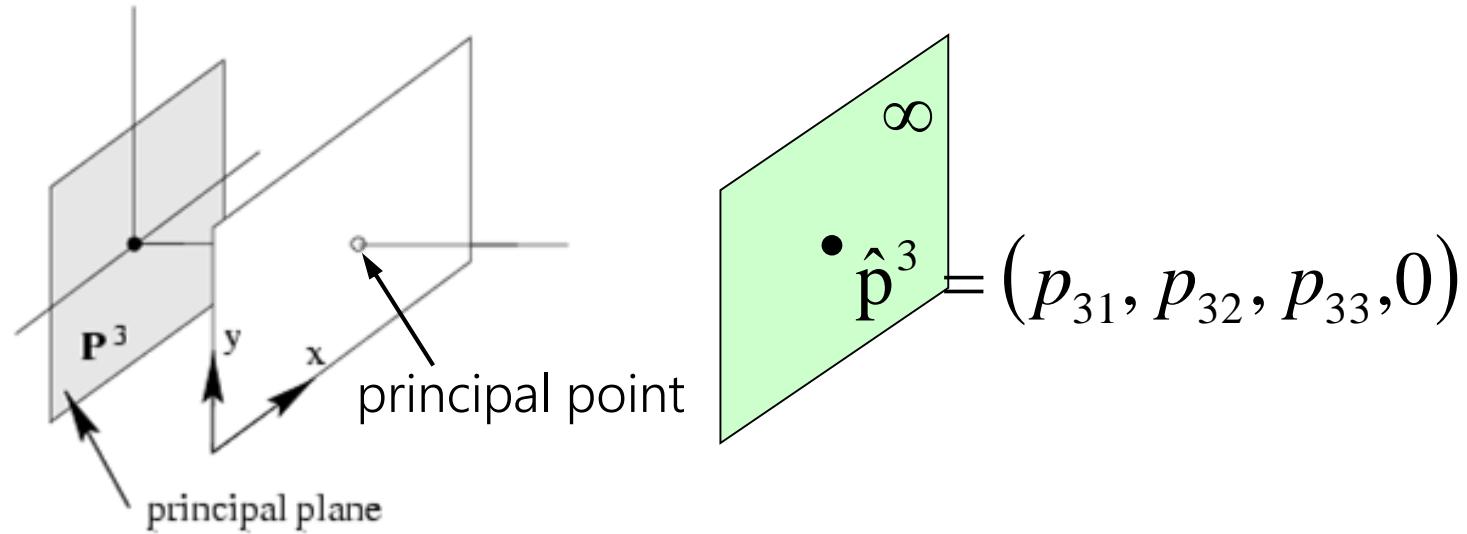


$$\begin{bmatrix} 0 \\ y \\ w \end{bmatrix} = \begin{bmatrix} p_1^T \\ p_2^T \\ p_3^T \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$



note: p^1, p^2 dependent on image parametrisation

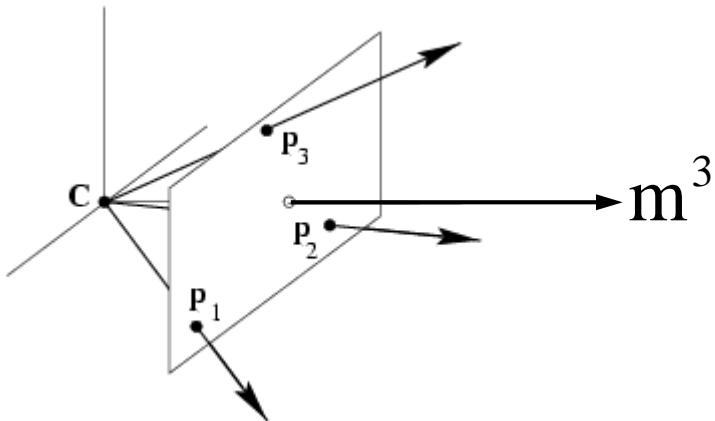
The principal point



$$\mathbf{x}_0 = \mathbf{P}\hat{\mathbf{p}}^3 = \mathbf{M}\mathbf{m}^3$$

The principal axis vector

vector defining front side of camera



$$\mathbf{x} = \mathbf{P}_{\text{cam}} \mathbf{X}_{\text{cam}} = \mathbf{K}[\mathbf{I} | \mathbf{0}] \mathbf{X}_{\text{cam}}$$

$$\mathbf{P}_{\text{cam}} \mapsto k \mathbf{P}_{\text{cam}}$$

$$\mathbf{P} = k \mathbf{K} \mathbf{R} [\mathbf{I} | -\tilde{\mathbf{C}}] = [\mathbf{M} | \mathbf{p}_4]$$

$$\mathbf{P}_{\text{cam}} \mapsto k \mathbf{P}_{\text{cam}}$$

$$\mathbf{v} = \det(\mathbf{M}) \mathbf{m}^3 = (0, 0, 1)^T$$

$$\mathbf{v} \mapsto k^4 \mathbf{v} \quad (\text{direction unaffected})$$

$$\mathbf{v} \mapsto \det(k \mathbf{M}) k \mathbf{m}^3 = k^4 \mathbf{v}$$

because $\det(\mathbf{R}) > 0$

Action of projective camera on a point

Forward projection

$$\mathbf{x} = \mathbf{P}\mathbf{X}$$

$$\mathbf{x} = \mathbf{P}\mathbf{D} = [\mathbf{M} \mid \mathbf{p}_4]\mathbf{D} = \mathbf{M}\mathbf{d}$$

Back-projection

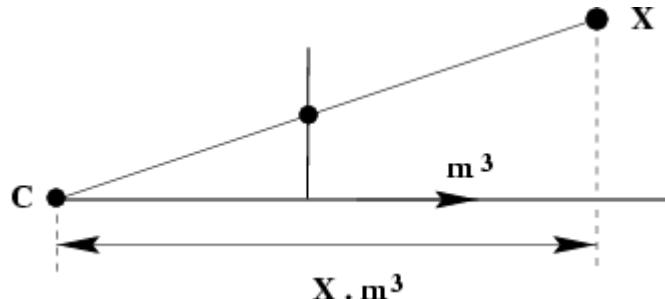
$$\mathbf{P}\mathbf{C} = 0$$

$$\mathbf{X} = \mathbf{P}^+ \mathbf{x} \quad \mathbf{P}^+ = \mathbf{P}^\top (\mathbf{P}\mathbf{P}^\top)^{-1} \quad \mathbf{P}\mathbf{P}^+ = \mathbf{I}$$

(pseudo-inverse)

$$\mathbf{X}(\lambda) = \mathbf{P}^+ \mathbf{x} + \lambda \mathbf{C}$$

Depth of points



$$w = P^3T X = P^3T(X - C) = m^3T(\tilde{X} - \tilde{C})$$

(PC=0) (dot product)

If $\det M > 0$; $\|m^3\| = 1$,
then m^3 unit vector in positive direction

$$\text{depth}(X; P) = \frac{\text{sign}(\det M) w}{T \|m^3\|}$$

$$X = (X, Y, Z, T)^T$$

Camera matrix decomposition

Finding the camera center

$$\mathbf{P}\mathbf{C} = \mathbf{0} \quad (\text{use SVD to find null-space})$$

$$X = \det([p_2, p_3, p_4]) \quad Y = -\det([p_1, p_3, p_4])$$

$$Z = \det([p_1, p_2, p_4]) \quad T = -\det([p_1, p_2, p_3])$$

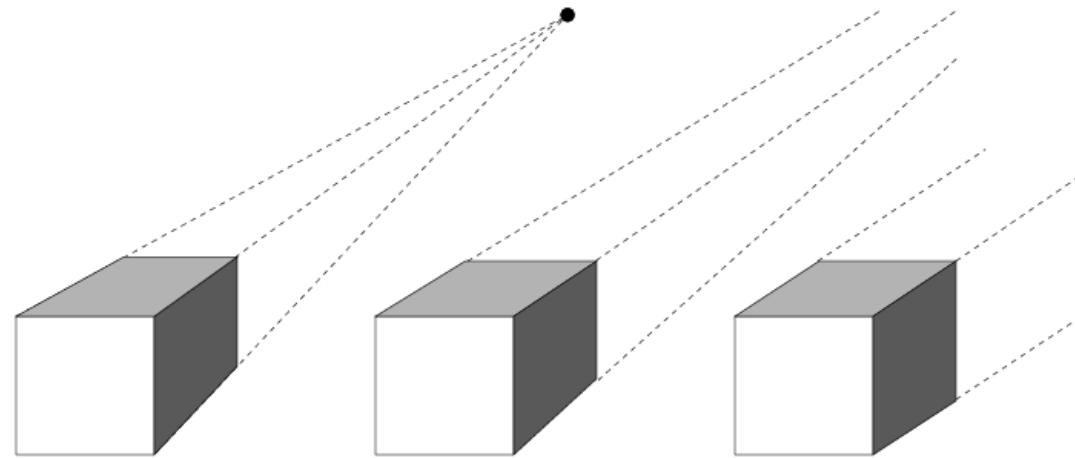
Finding the camera orientation and internal parameters

$$\mathbf{M} = \mathbf{K}\mathbf{R} \quad (\text{use RQ decomposition } \sim \text{QR})$$

(if only QR, invert)

$$\boxed{} = (\boxed{Q} \ \boxed{\triangledown R})^{-1} = \boxed{\triangledown R}^{-1} \ \boxed{Q}^{-1}$$

Affine cameras



perspective

weak perspective

increasing focal length →

increasing distance from camera →



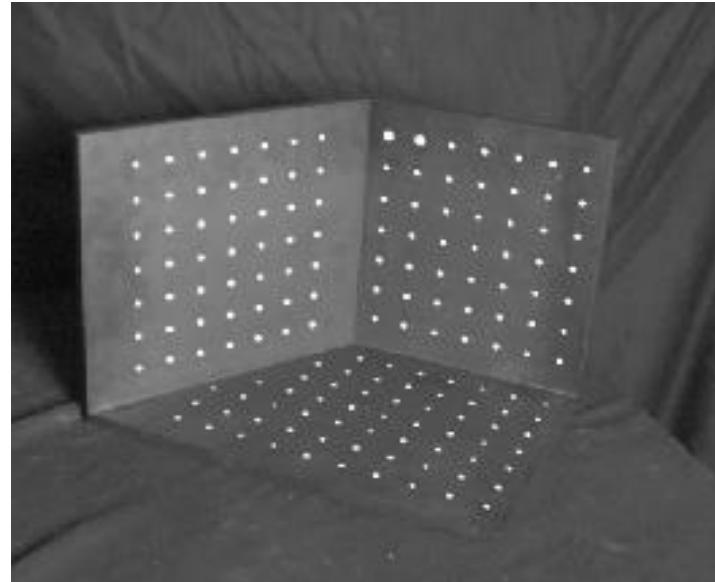
Camera calibration

Why calibrate?

- Want to solve for 3D geometry
- Could solve for 3D shape without known cameras:
 - Structure-from-motion (unknown extrinsics)
 - Self-calibration (unknown intrinsics & extrinsics)
- So why bother pre-calibrating the camera?
 - Simplifies the 3D reconstruction problem – fewer parameters to solve for
 - Improves accuracy
 - Not too hard to do
 - Eliminates certain ambiguities (scale of scene)

Camera calibration

- Determine camera parameters from known 3D points or calibration object(s)
 - internal or intrinsic parameters such as focal length, optical center, aspect ratio: what kind of camera?
 - external or extrinsic (pose) parameters: where is the camera?
- How can we do this?



Camera matrix

- Fold intrinsic calibration matrix \mathbf{K} and extrinsic pose parameters (\mathbf{R}, \mathbf{t}) together into a camera matrix
 - $\mathbf{M} = \mathbf{K} [\mathbf{R} \mid \mathbf{t}]$

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \sim \begin{bmatrix} m_{00} & m_{01} & m_{02} & m_{03} \\ m_{10} & m_{11} & m_{12} & m_{13} \\ m_{20} & m_{21} & m_{22} & 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \\ 1 \end{bmatrix}$$

- (put 1 in lower right-hand corner for 11 degrees of freedom)

Camera matrix calibration

- Directly estimate 11 unknowns in the \mathbf{M} matrix using known 3D points (X_i, Y_i, Z_i) and measured feature positions (u_i, v_i)

$$u_i = \frac{m_{00}X_i + m_{01}Y_i + m_{02}Z_i + m_{03}}{m_{20}X_i + m_{21}Y_i + m_{22}Z_i + 1}$$
$$v_i = \frac{m_{10}X_i + m_{11}Y_i + m_{12}Z_i + m_{13}}{m_{20}X_i + m_{21}Y_i + m_{22}Z_i + 1}$$

Camera matrix calibration

- Linear regression:
 - Bring denominator over, solve set of (over-determined) linear equations.

How?

$$\begin{aligned} u_i(m_{20}X_i + m_{21}Y_i + m_{22}Z_i + 1) &= \\ &m_{00}X_i + m_{01}Y_i + m_{02}Z_i + m_{03} \\ v_i(m_{20}X_i + m_{21}Y_i + m_{22}Z_i + 1) &= \\ &m_{10}X_i + m_{11}Y_i + m_{12}Z_i + m_{13} \end{aligned}$$

- Least squares (pseudo-inverse)
- Is this good enough?

Camera matrix calibration

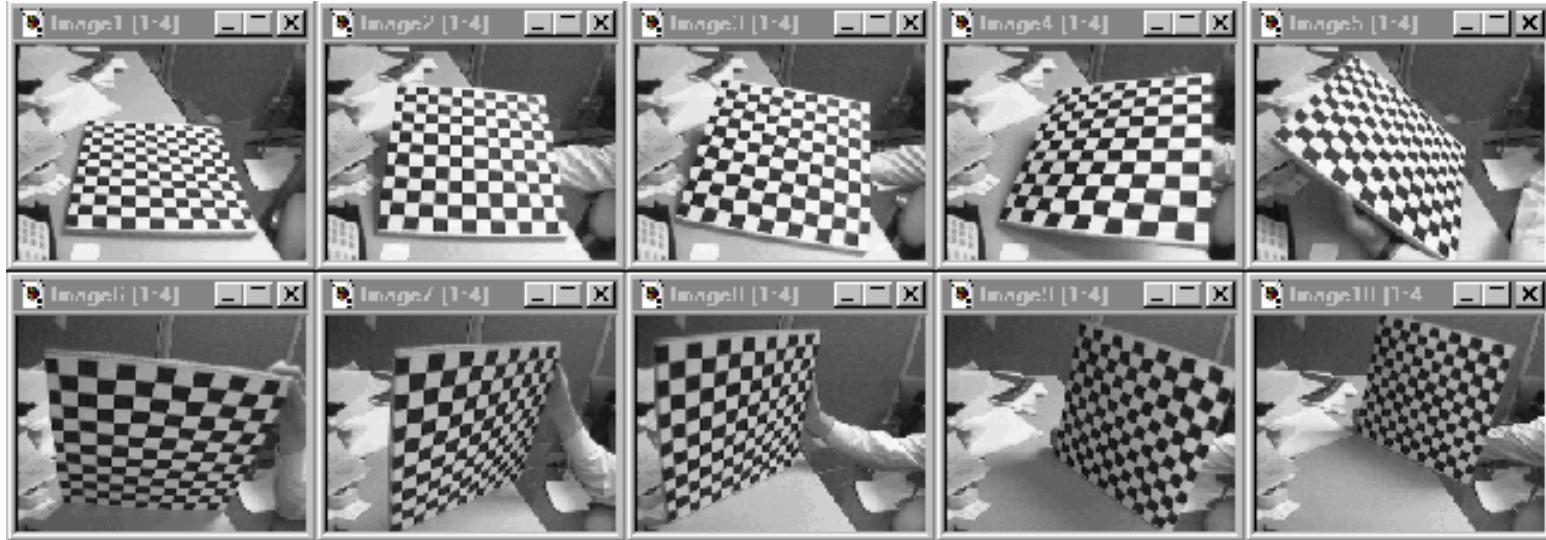
- Advantages:
 - very simple to formulate and solve
 - can recover \mathbf{K} $[\mathbf{R} \mid \mathbf{t}]$ from \mathbf{M} using QR decomposition
[Golub & VanLoan 96]
- Disadvantages:
 - doesn't compute internal parameters
 - more unknowns than true degrees of freedom
 - need a separate camera matrix for each new view

Separate intrinsic and extrinsic calibration

- Advantages:
 - can solve for more than one camera pose at a time
 - potentially fewer degrees of freedom

- Disadvantages:
 - more complex update rules
 - need a good initialization (recover $K [R | t]$ from M)

Multi-plane calibration



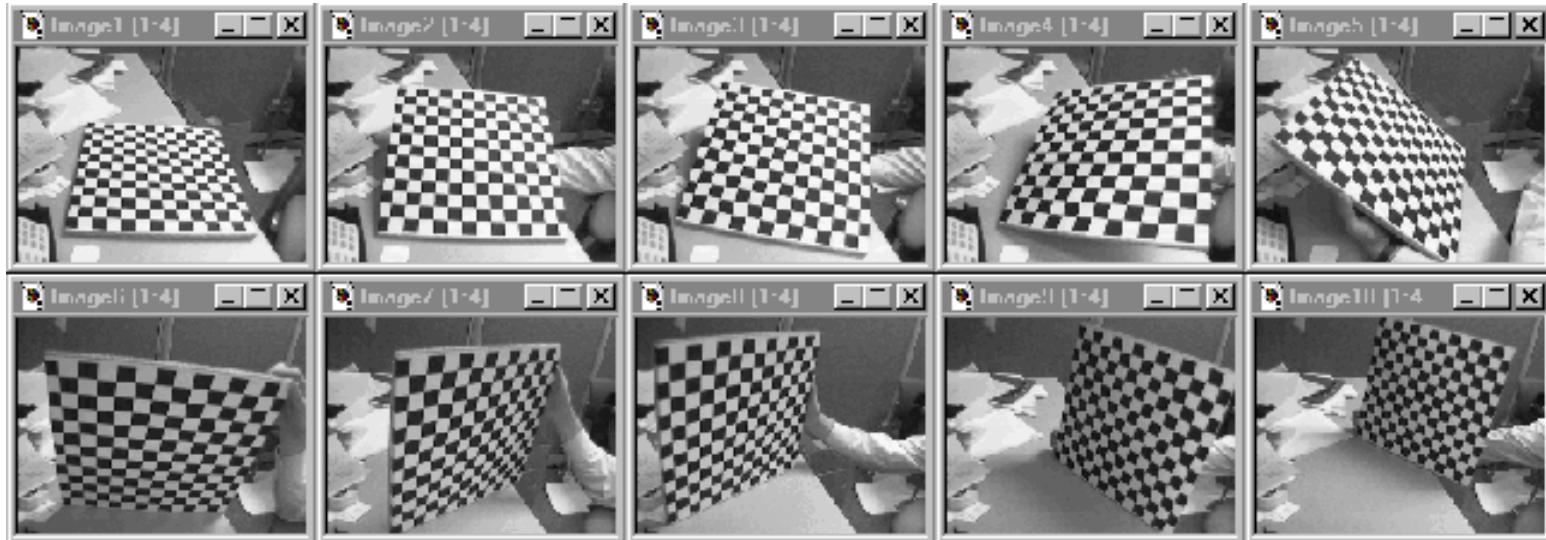
Images courtesy Jean-Yves Bouguet, Intel Corp.

Advantage

- Only requires a plane
- Don't have to know positions/orientations
- Good code available online!
 - OpenCV library: <http://www.opencv.org/>
 - MATLAB toolbox by Jean-Yves Bouguet:
http://www.vision.caltech.edu/bouguetj/calib_doc/index.html

Disadvantages?

Multi-plane calibration



Images courtesy Jean-Yves Bouguet, Intel Corp.

Need $3D \rightarrow 2D$ correspondence

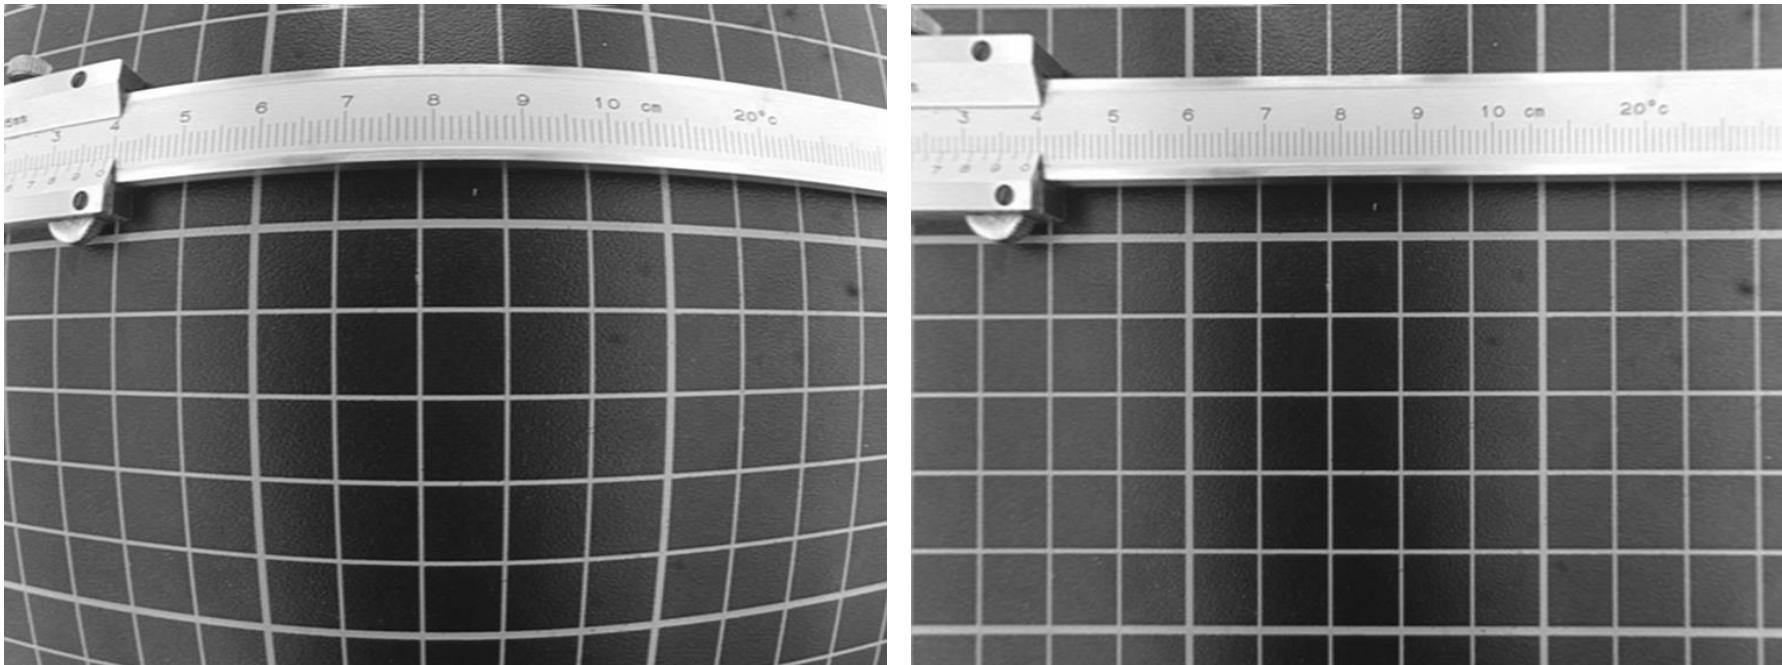
- User provided (lots of clicking)
- User seeded (some clicking)
- Fully automatic?

Lens distortion

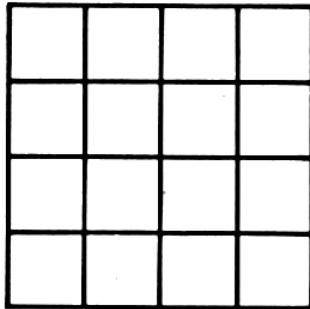
- radial & tangential distortion of the image
- caused by imperfect (read: cheap) lenses
- most noticeable near the edge of the lens



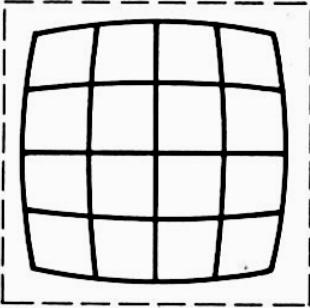
Making straight lines straight



Barrel distortion



No distortion

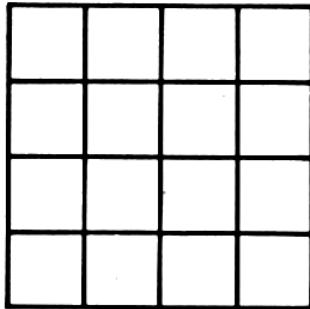


Barrel

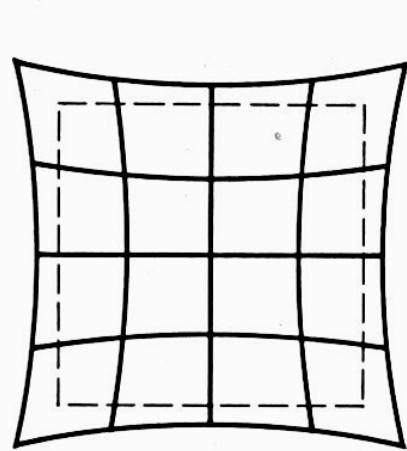


Wide Angle Lens

Pin-cushion distortion



No distortion



Pin cushion



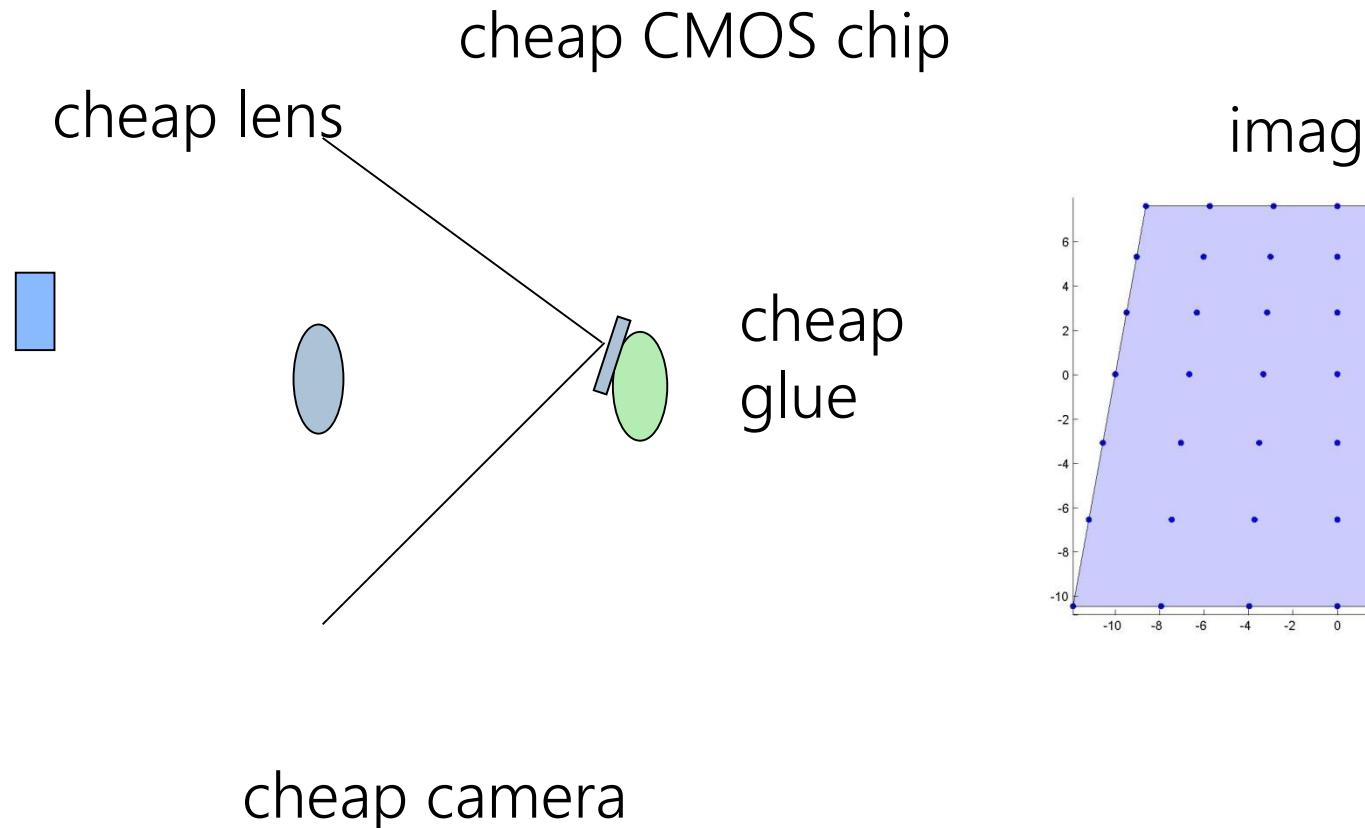
Telephoto lens

Radial distortion – k_1, k_2, k_5

$$\begin{pmatrix} x \\ y \end{pmatrix} = \begin{pmatrix} x_d \\ y_d \end{pmatrix} \cdot \left(1 + k_1 \cdot r^2 + k_2 \cdot r^4 + k_5 \cdot r^6 \right)$$

↑
distance from center

Tangential distortion – k_3, k_4



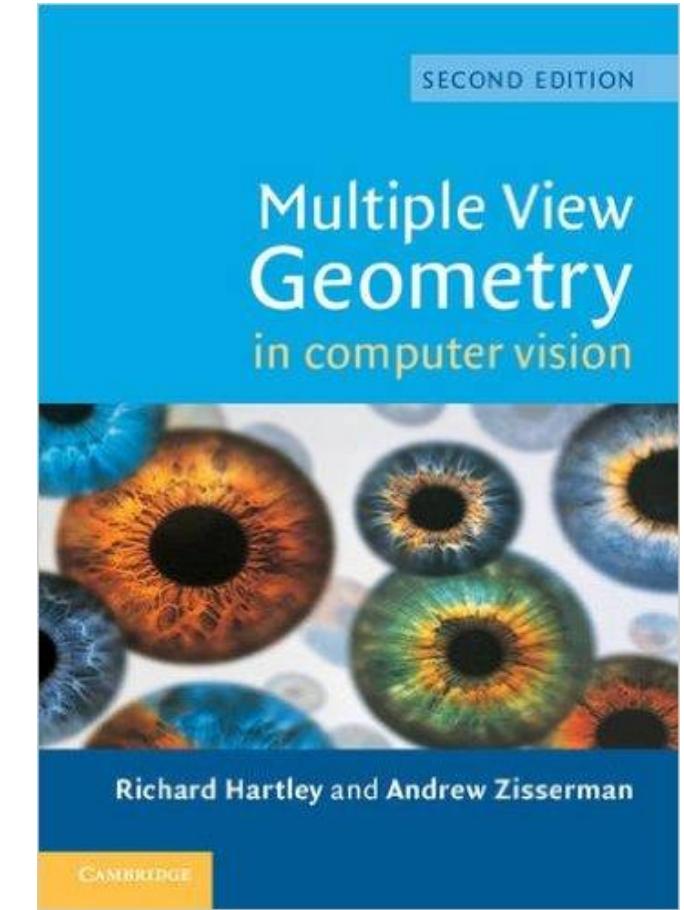
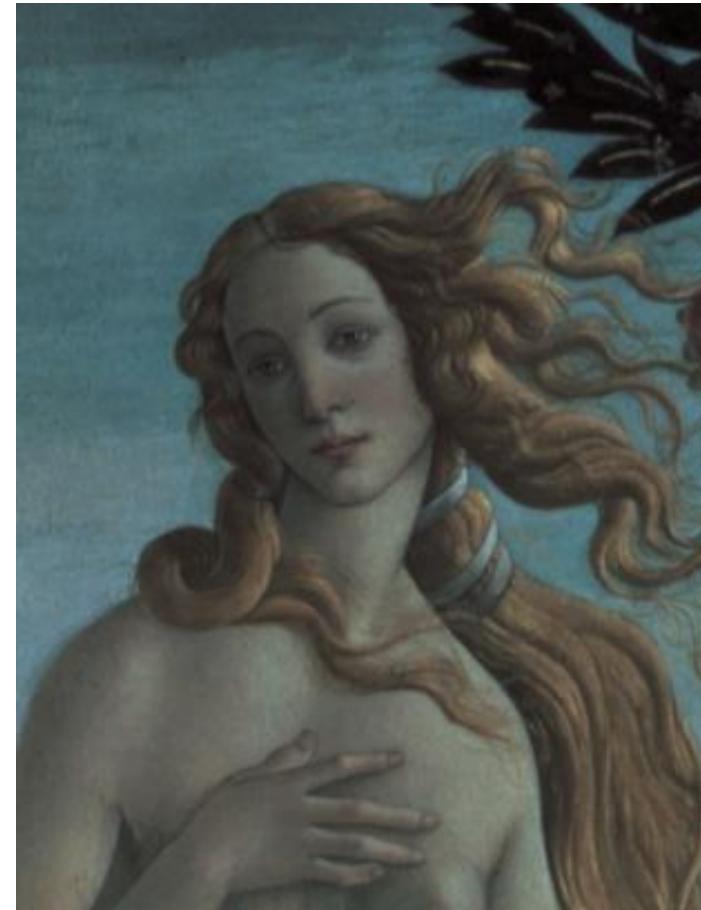
Recap of intrinsic calibration parameters

- Focal lengths f_x, f_y
- Principal point c_x, c_y
- Skew s

$$\mathbf{K} = \begin{bmatrix} f_x & s & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

- Lens distortion k_1, k_2, k_3, k_4, k_5

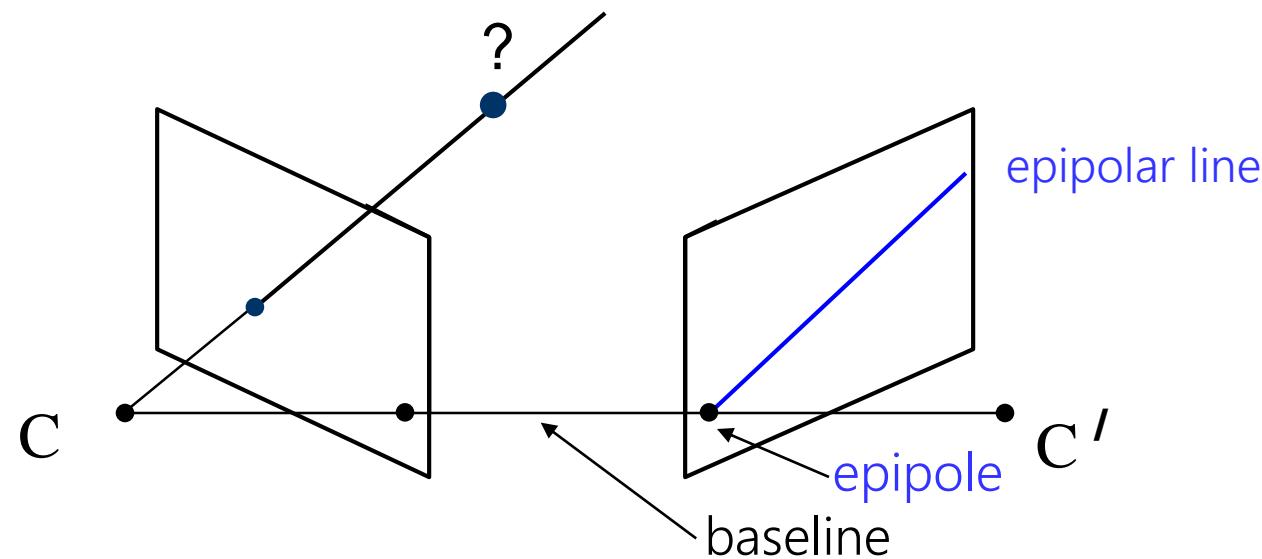
Epipolar geometry



**Multiple View Geometry
in Computer Vision**
Hartley & Zisserman, 2004

Epipolar geometry

- Given an image point in one view, where is the corresponding point in the other view?



- A point in one view “generates” an **epipolar line** in the other view
- The corresponding point lies on this line

Epipolar line

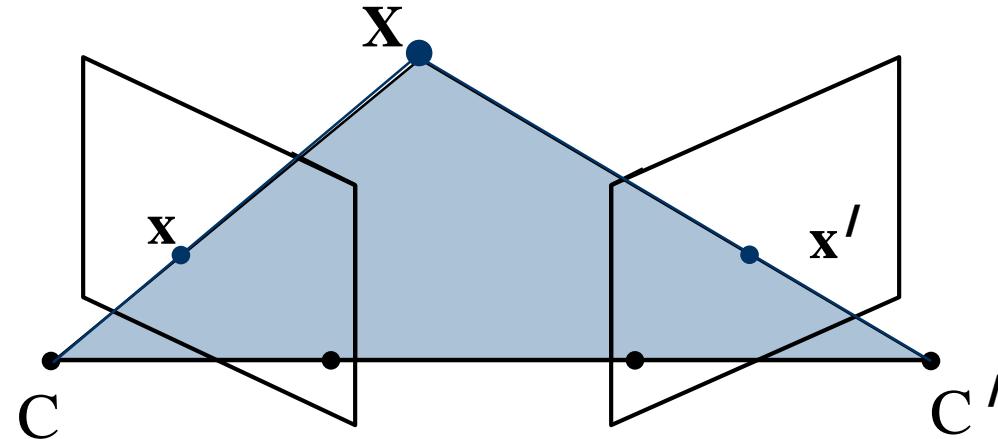


Epipolar constraint

- Reduces correspondence problem to 1D search along an epipolar line

Epipolar geometry continued

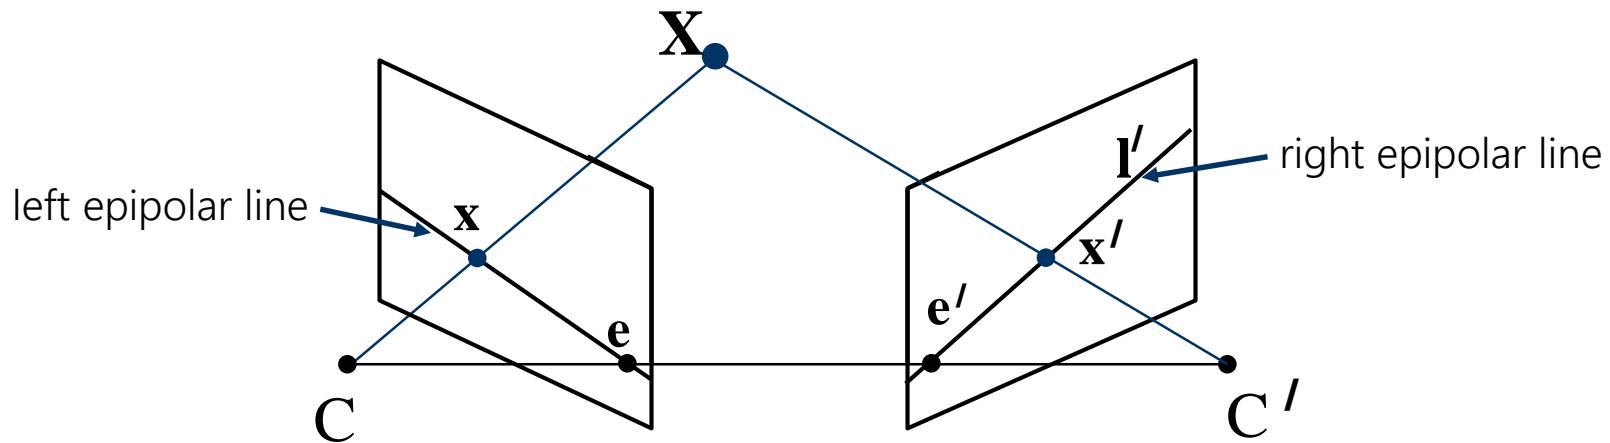
- Epipolar geometry is a consequence of the coplanarity of the camera centres and scene point



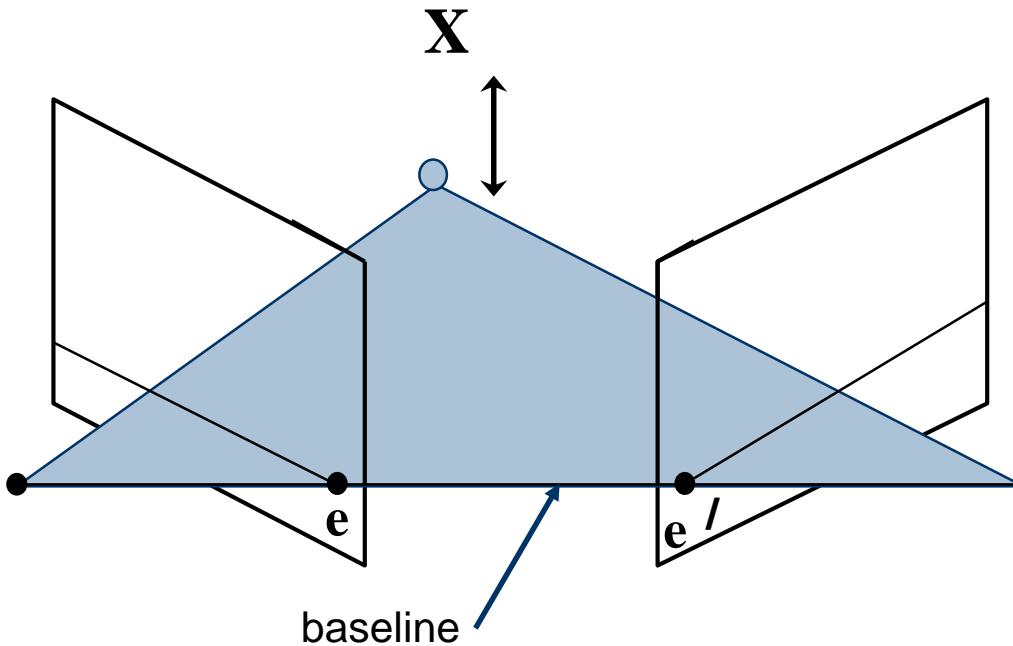
The camera centres, corresponding points and scene point lie in a single plane, known as the **epipolar plane**

Nomenclature

- The epipolar line l' is the image of the ray through x
- The epipole e is the point of intersection of the line joining the camera centres with the image plane
 - this line is the baseline for a stereo rig, and
 - the translation vector for a moving camera
- The epipole is the image of the other camera centre: $e=PC'$, $e'=P'C$



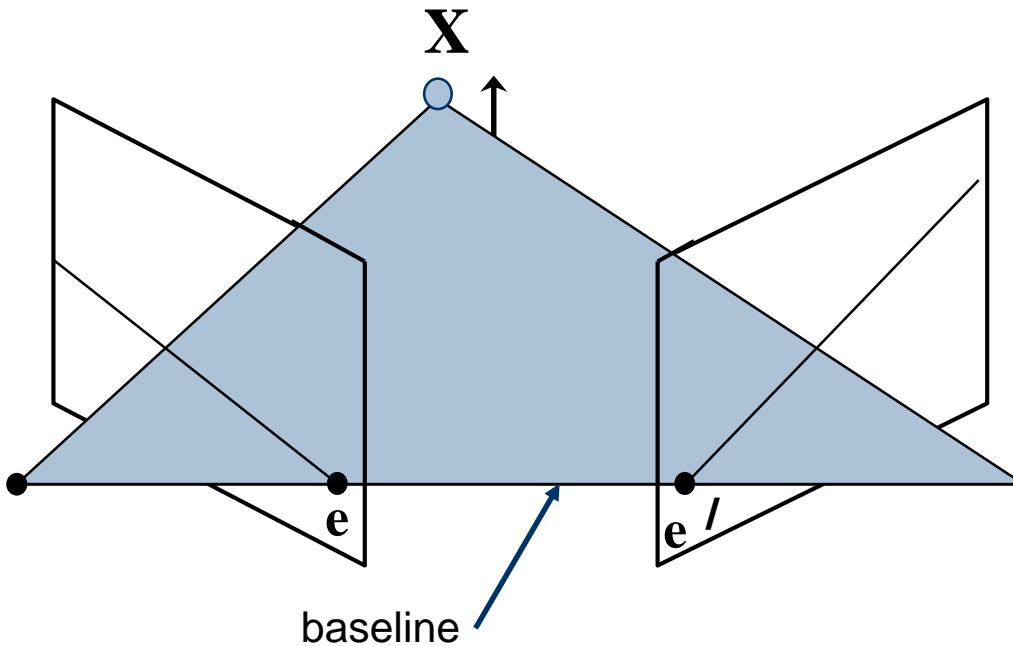
The epipolar pencil



As the position of the 3D point X varies, the epipolar planes “rotate” about the baseline. This family of planes is known as an **epipolar pencil**. All epipolar lines intersect at the epipole.

(a pencil is a one parameter family)

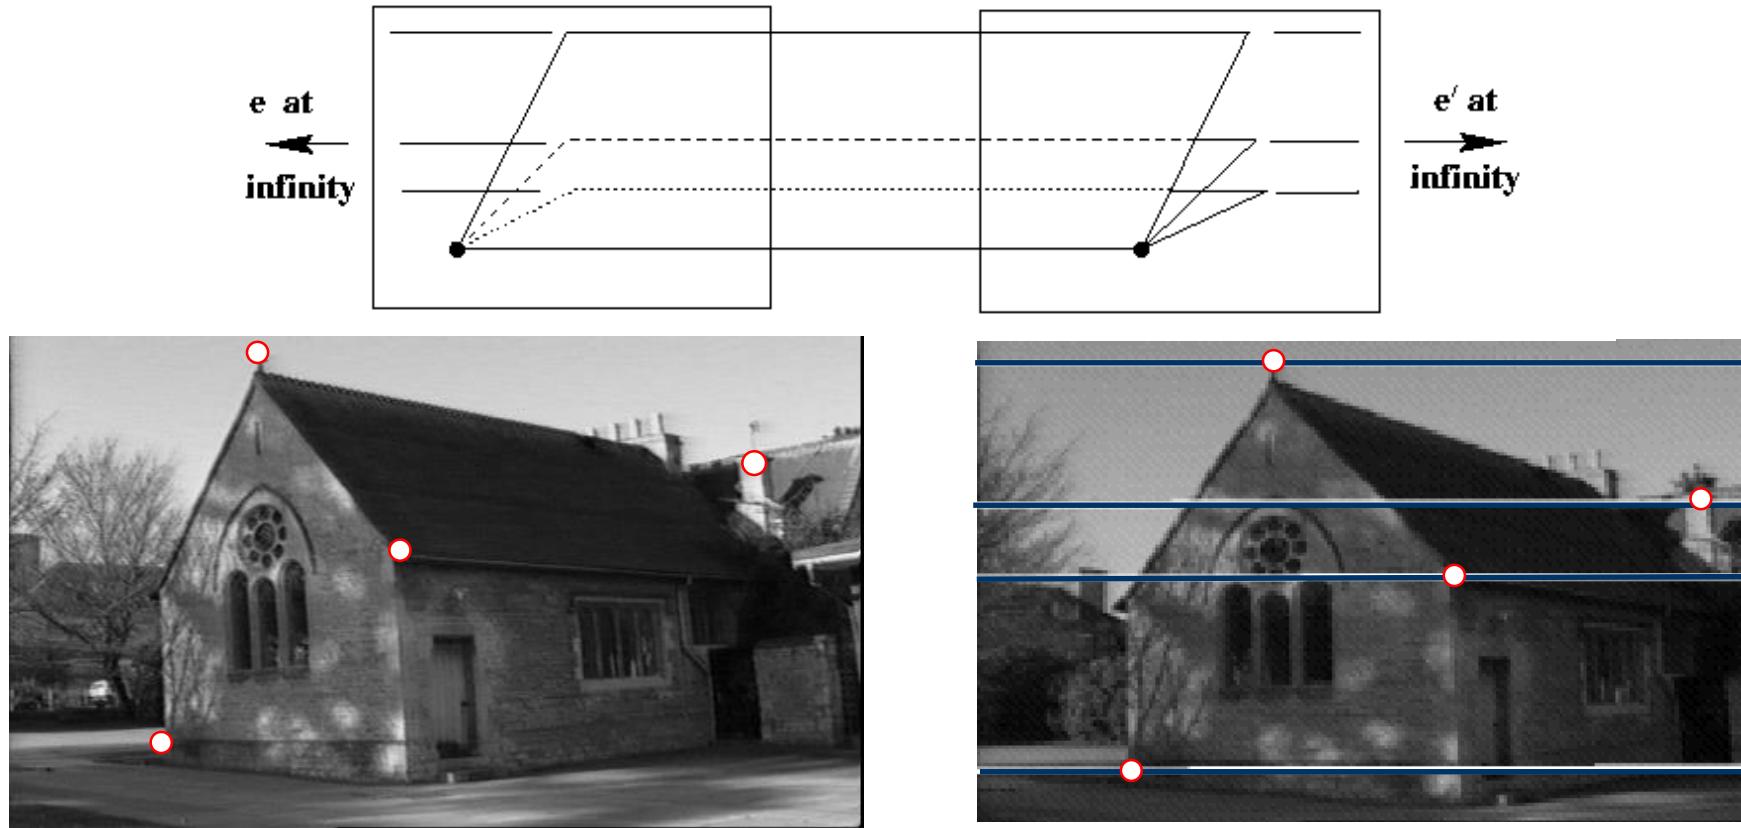
The epipolar pencil



As the position of the 3D point X varies, the epipolar planes “rotate” about the baseline. This family of planes is known as an **epipolar pencil**. All epipolar lines intersect at the epipole.

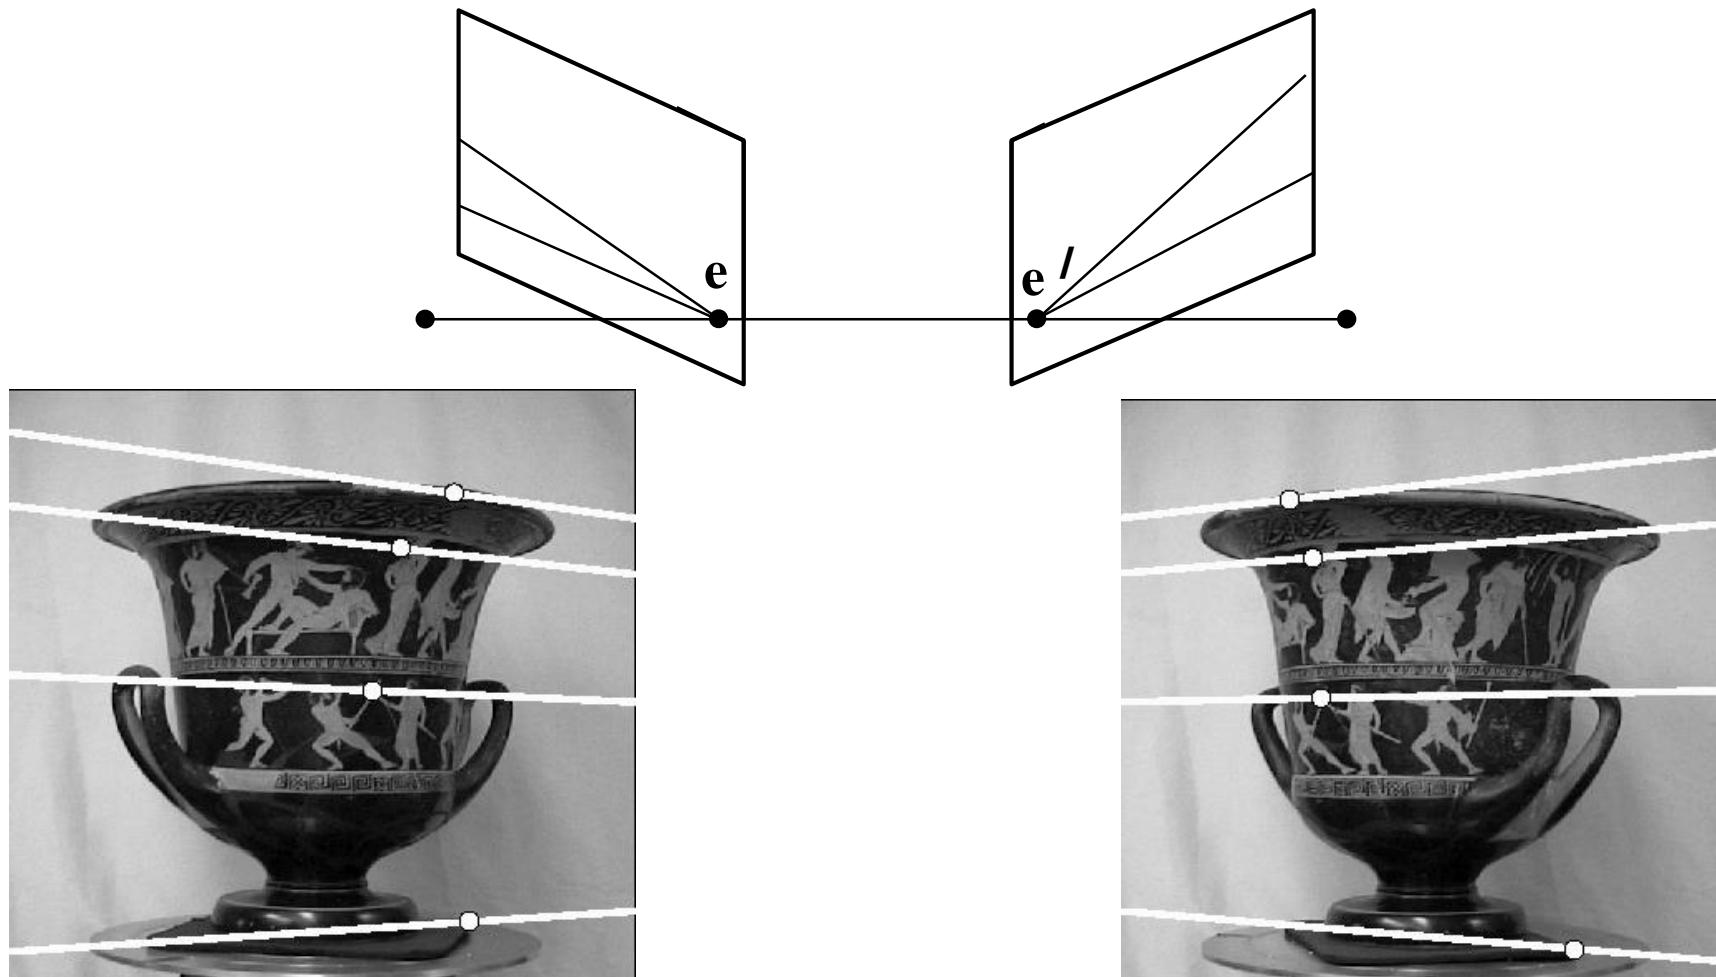
(a pencil is a one parameter family)

Epipolar geometry example I: parallel cameras



Epipolar geometry depends **only** on the relative pose (position and orientation) and internal parameters of the two cameras, i.e. the position of the camera centres and image planes. It does **not** depend on the scene structure (3D points external to the camera).

Epipolar geometry example II: converging cameras



Note: epipolar lines are in general **not** parallel

Homogeneous notation for lines

A **line** in 2D is represented by the homogeneous 3-vector

$$\mathbf{l} = \begin{pmatrix} l_1 \\ l_2 \\ l_3 \end{pmatrix}$$

which is the line $l_1x + l_2y + l_3 = 0$.

Example represent the line $y = 1$ as a homogeneous vector.

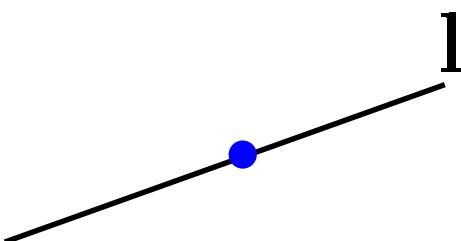
Write the line as $-y + 1 = 0$ then $l_1 = 0, l_2 = -1, l_3 = 1$, and $\mathbf{l} = (0, -1, 1)^\top$.

Note that $\mu(l_1x + l_2y + l_3) = 0$ represents the same line (only the ratio of the homogeneous line coordinates is significant).

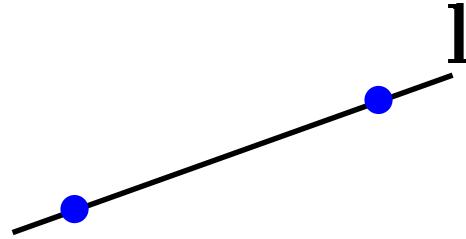
Writing both the point and line in homogeneous coordinates gives

$$l_1x_1 + l_2x_2 + l_3x_3 = 0$$

- **point on line** $\mathbf{l} \cdot \mathbf{x} = 0$ or $\mathbf{l}^\top \mathbf{x} = 0$ or $\mathbf{x}^\top \mathbf{l} = 0$



- The line l through the two points p and q is $l = p \times q$



Proof

$$l.p = (p \times q).p = 0 \quad l.q = (p \times q).q = 0$$

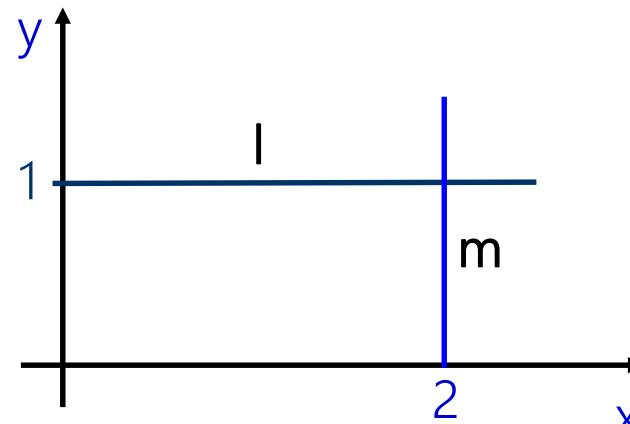
- The intersection of two lines l and m is the point $x = l \times m$

Example: compute the point of intersection of the two lines l and m in the figure below

$$l = \begin{pmatrix} 0 \\ -1 \\ 1 \end{pmatrix} \quad m = \begin{pmatrix} -1 \\ 0 \\ 2 \end{pmatrix}$$

$$x = l \times m = \begin{vmatrix} i & j & k \\ 0 & -1 & 1 \\ -1 & 0 & 2 \end{vmatrix} = \begin{pmatrix} -2 \\ -1 \\ -1 \end{pmatrix}$$

which is the point $(2,1)$



Matrix representation of the vector cross product

The vector product $\mathbf{v} \times \mathbf{x}$ can be represented as a matrix multiplication

$$\mathbf{v} \times \mathbf{x} = \begin{pmatrix} v_2x_3 - v_3x_2 \\ v_3x_1 - v_1x_3 \\ v_1x_2 - v_2x_1 \end{pmatrix} = [\mathbf{v}]_{\times}\mathbf{x}$$

where

$$[\mathbf{v}]_{\times} = \begin{bmatrix} 0 & -v_3 & v_2 \\ v_3 & 0 & -v_1 \\ -v_2 & v_1 & 0 \end{bmatrix}$$

- $[\mathbf{v}]_{\times}$ is a 3×3 skew-symmetric matrix of rank 2.
- \mathbf{v} is the null-vector of $[\mathbf{v}]_{\times}$, since $\mathbf{v} \times \mathbf{v} = [\mathbf{v}]_{\times}\mathbf{v} = \mathbf{0}$.

Example

Compute the cross product of \mathbf{l} and \mathbf{m}

$$\mathbf{l} = \begin{pmatrix} 0 \\ -1 \\ 1 \end{pmatrix}$$

$$\mathbf{m} = \begin{pmatrix} -1 \\ 0 \\ 2 \end{pmatrix}$$

$$[\mathbf{v}]_{\times} = \begin{bmatrix} 0 & -v_3 & v_2 \\ v_3 & 0 & -v_1 \\ -v_2 & v_1 & 0 \end{bmatrix}$$

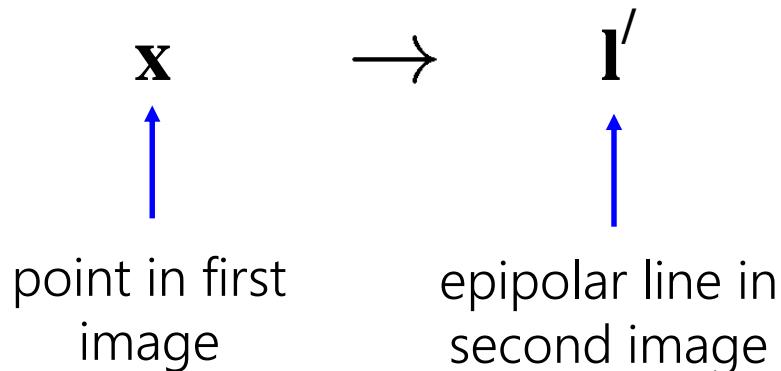
$$\mathbf{x} = \mathbf{l} \times \mathbf{m} = [\mathbf{l}]_{\times} \mathbf{m} = \begin{bmatrix} 0 & -1 & -1 \\ 1 & 0 & 0 \\ 1 & 0 & 0 \end{bmatrix} \begin{pmatrix} -1 \\ 0 \\ 2 \end{pmatrix} = \begin{pmatrix} -2 \\ -1 \\ -1 \end{pmatrix}$$

Note

$$\begin{bmatrix} 0 & -1 & -1 \\ 1 & 0 & 0 \\ 1 & 0 & 0 \end{bmatrix} \begin{pmatrix} 0 \\ -1 \\ 1 \end{pmatrix} = \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix}$$

Algebraic representation of epipolar geometry

We know that the epipolar geometry defines a mapping

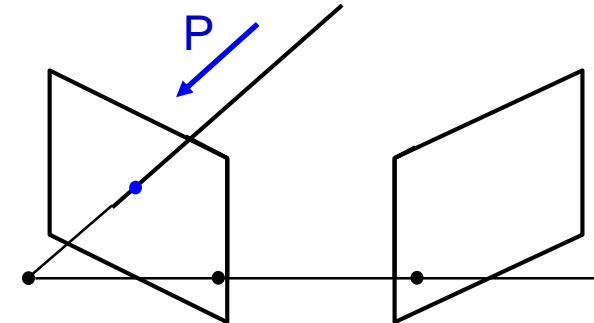


- the map only depends on the cameras P, P' (not on structure)
- it will be shown that the map is **linear** and can be written as $\mathbf{l}' = \mathbf{F}\mathbf{x}$, where \mathbf{F} is a 3×3 matrix called the **fundamental matrix**

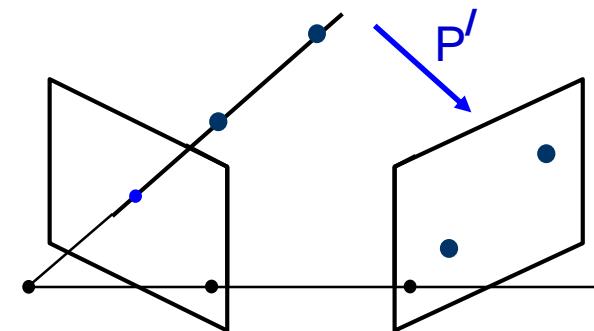
Derivation of the algebraic expression $\mathbf{l}' = \mathbf{F}\mathbf{x}$

Outline

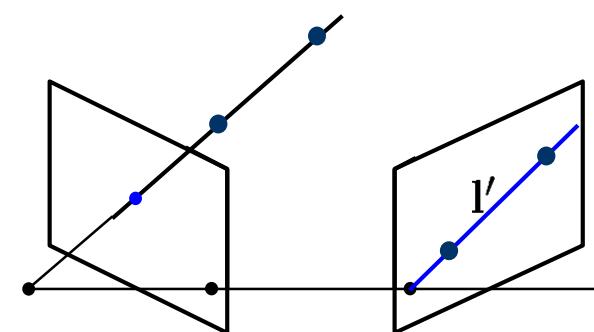
Step 1: for a point x in the first image back project a ray with camera P



Step 2: choose two points on the ray and project into the second image with camera P'



Step 3: compute the line through the two image points using the relation $\mathbf{l}' = \mathbf{p} \times \mathbf{q}$



- choose camera matrices

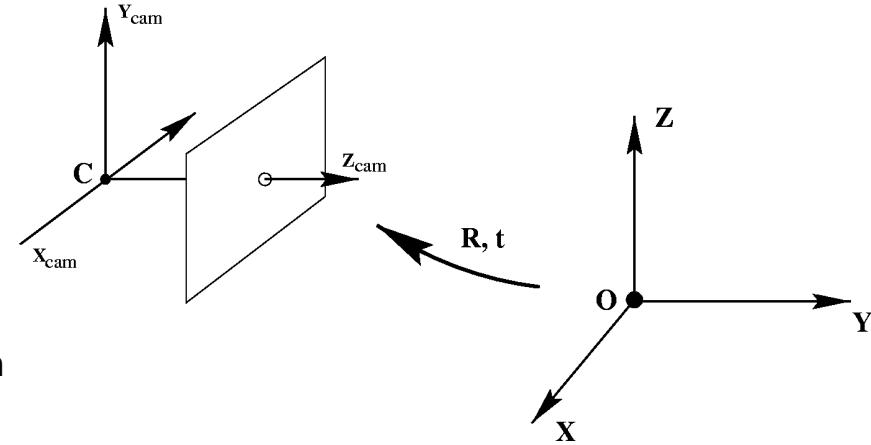
$$P = K [R \mid t]$$

internal
calibration

rotation

translation

from world to camera
coordinate frame



- first camera

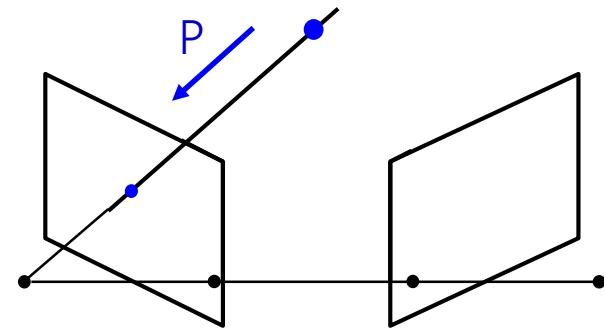
$$P = K [I \mid 0]$$

world coordinate frame aligned with first camera

- second camera

$$P' = K' [R \mid t]$$

Step 1: for a point x in the first image back project a ray with camera $P = K [I | 0]$



A point x back projects to a ray

$$\begin{pmatrix} x \\ y \\ z \end{pmatrix} = zK^{-1} \begin{pmatrix} x \\ y \\ 1 \end{pmatrix} = zK^{-1}x$$

where Z is the point's depth, since

$$x(z) = \begin{pmatrix} zK^{-1}x \\ 1 \end{pmatrix}$$

satisfies

$$Px(z) = K[I | 0]x(z) = x$$

Step 2: choose two points on the ray and project into the second image with camera P'

Consider two points on the ray $\mathbf{X}(z) = \begin{pmatrix} z\mathbf{K}^{-1}\mathbf{x} \\ 1 \end{pmatrix}$

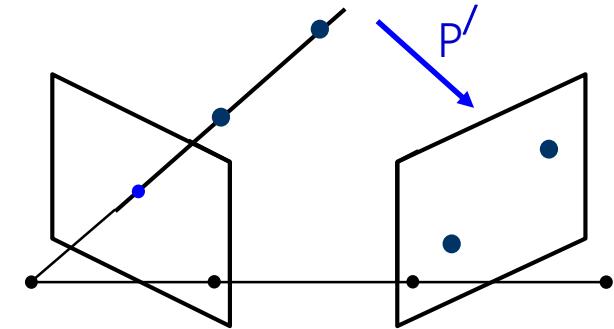
- $Z = 0$ is the camera centre $\begin{pmatrix} \mathbf{0} \\ 1 \end{pmatrix}$

- $Z = \infty$ is the point at infinity $\begin{pmatrix} \mathbf{K}^{-1}\mathbf{x} \\ 0 \end{pmatrix}$

Project these two points into the second view

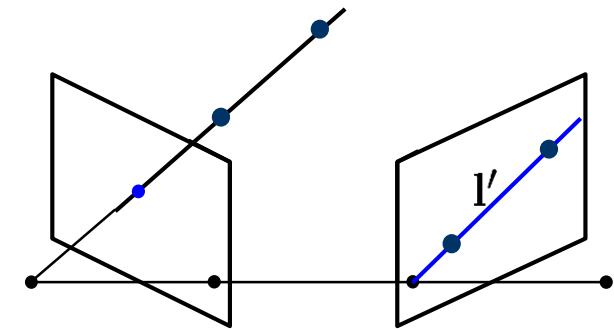
$$P' \begin{pmatrix} \mathbf{0} \\ 1 \end{pmatrix} = K'[R \mid t] \begin{pmatrix} \mathbf{0} \\ 1 \end{pmatrix} = K'\mathbf{t}$$

$$P' \begin{pmatrix} K^{-1}\mathbf{x} \\ 0 \end{pmatrix} = K'[R \mid t] \begin{pmatrix} K^{-1}\mathbf{x} \\ 0 \end{pmatrix} = K'R\mathbf{K}^{-1}\mathbf{x}$$



Step 3: compute the line through the two image points using the relation $\mathbf{l}' = \mathbf{p} \times \mathbf{q}$

Compute the line through the points $\mathbf{l}' = (\mathbf{K}'\mathbf{t}) \times (\mathbf{K}'\mathbf{R}\mathbf{K}^{-1}\mathbf{x})$



Using the identity $(\mathbf{M}\mathbf{a}) \times (\mathbf{M}\mathbf{b}) = \mathbf{M}^{-\top}(\mathbf{a} \times \mathbf{b})$ where $\mathbf{M}^{-\top} = (\mathbf{M}^{-1})^\top = (\mathbf{M}^\top)^{-1}$

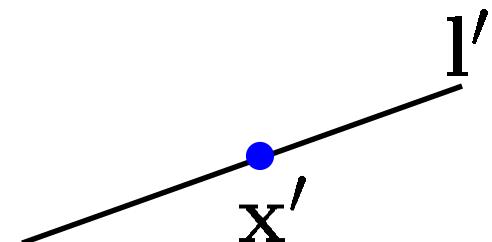
$$\mathbf{l}' = \mathbf{K}'^{-\top} \left(\mathbf{t} \times (\mathbf{R}\mathbf{K}^{-1}\mathbf{x}) \right) = \underbrace{\mathbf{K}'^{-\top} [\mathbf{t}]_\times \mathbf{R}\mathbf{K}^{-1}}_{\mathbf{F}} \mathbf{x}$$

\mathbf{F} is the fundamental matrix

$$\mathbf{l}' = \mathbf{Fx} \quad \mathbf{F} = \mathbf{K}'^{-\top} [\mathbf{t}]_\times \mathbf{R}\mathbf{K}^{-1}$$

Points \mathbf{x} and \mathbf{x}' correspond ($\mathbf{x} \leftrightarrow \mathbf{x}'$) then $\mathbf{x}'^\top \mathbf{l}' = 0$

$$\mathbf{x}'^\top \mathbf{Fx} = 0$$



Example 1: compute F for parallel cameras

$$P = K[I \mid \mathbf{0}] \quad P' = K'[R \mid t]$$

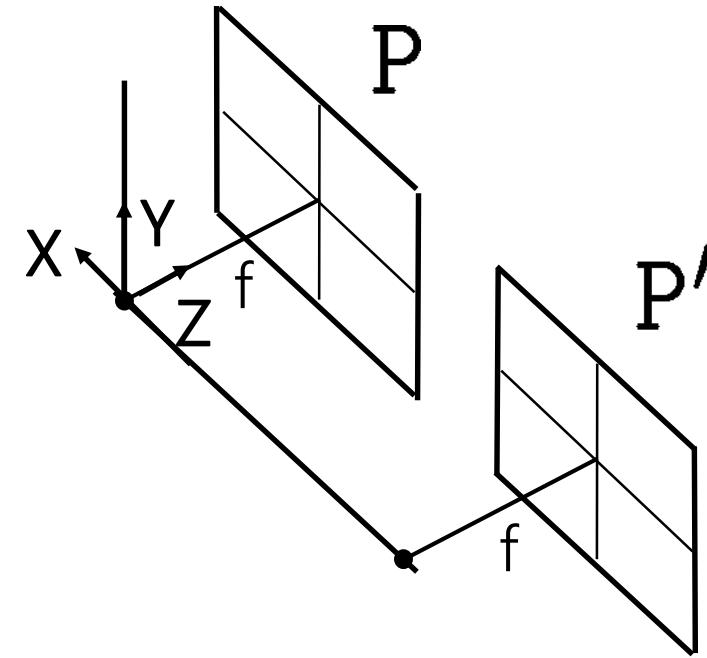
$$K = K' = \begin{bmatrix} f & 0 & 0 \\ 0 & f & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad R = I \quad t = \begin{pmatrix} t_x \\ 0 \\ 0 \end{pmatrix}$$

$$F = K'^{-\top}[t]_x R K^{-1}$$

$$= \begin{bmatrix} 1/f & 0 & 0 \\ 0 & 1/f & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & -t_x \\ 0 & t_x & 0 \end{bmatrix} \begin{bmatrix} 1/f & 0 & 0 \\ 0 & 1/f & 0 \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & -1 \\ 0 & 1 & 0 \end{bmatrix}$$

$$\mathbf{x}'^\top F \mathbf{x} = (x' \ y' \ 1) \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & -1 \\ 0 & 1 & 0 \end{bmatrix} \begin{pmatrix} x \\ y \\ 1 \end{pmatrix} = 0$$

reduces to $y = y'$, i.e. raster correspondence (horizontal scan-lines)



Example 1: compute F for parallel cameras

F is a rank 2 matrix

The epipole e is the null-space vector (kernel) of F (exercise), i.e. $Fe = 0$

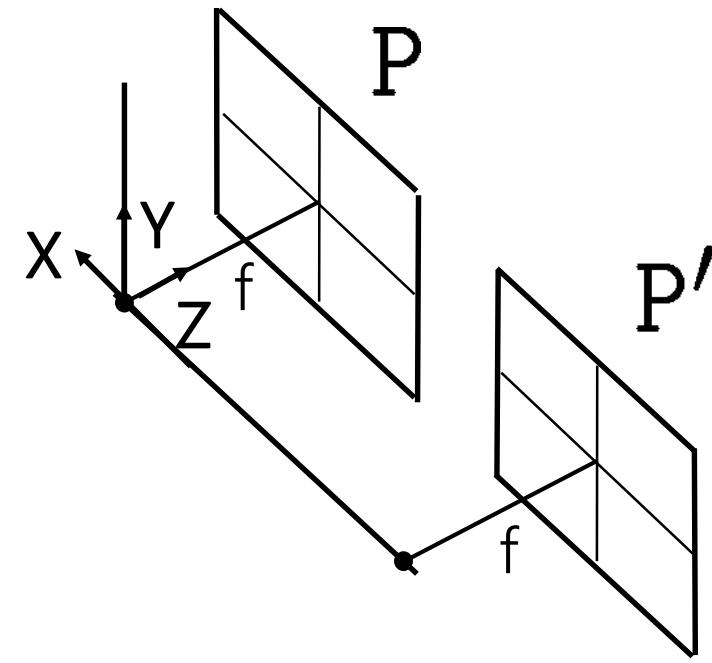
In this case

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & -1 \\ 0 & 1 & 0 \end{bmatrix} \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} = 0$$

so that

$$e = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix}$$

Geometric interpretation?



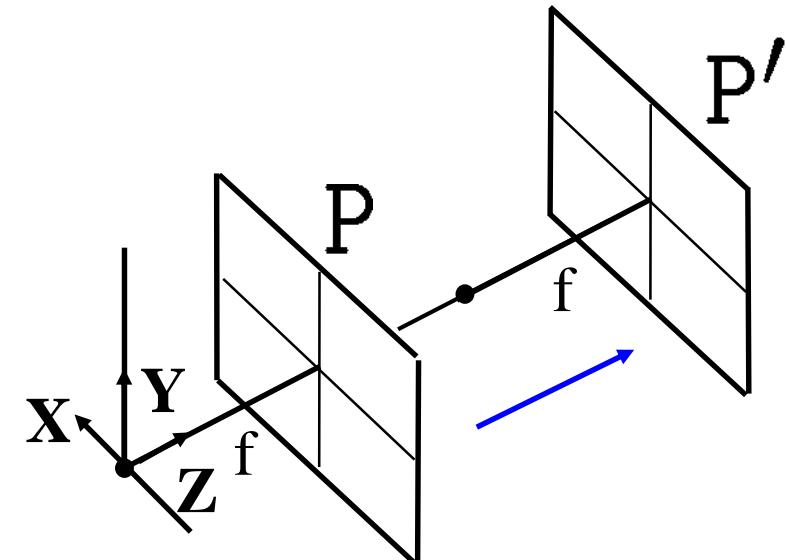
Example 2: compute F for forward-translating camera

$$P = K[I \mid \mathbf{0}] \quad P' = K'[R \mid t]$$

$$K = K' = \begin{bmatrix} f & 0 & 0 \\ 0 & f & 0 \\ 0 & 0 & 1 \end{bmatrix} \quad R = I \quad t = \begin{pmatrix} 0 \\ 0 \\ t_z \end{pmatrix}$$

$$F = K'^{-T}[t]_x R K^{-1}$$

$$\begin{aligned} &= \begin{bmatrix} 1/f & 0 & 0 \\ 0 & 1/f & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 0 & -t_z & 0 \\ t_z & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} 1/f & 0 & 0 \\ 0 & 1/f & 0 \\ 0 & 0 & 1 \end{bmatrix} \\ &= \begin{bmatrix} 0 & -1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} \end{aligned}$$

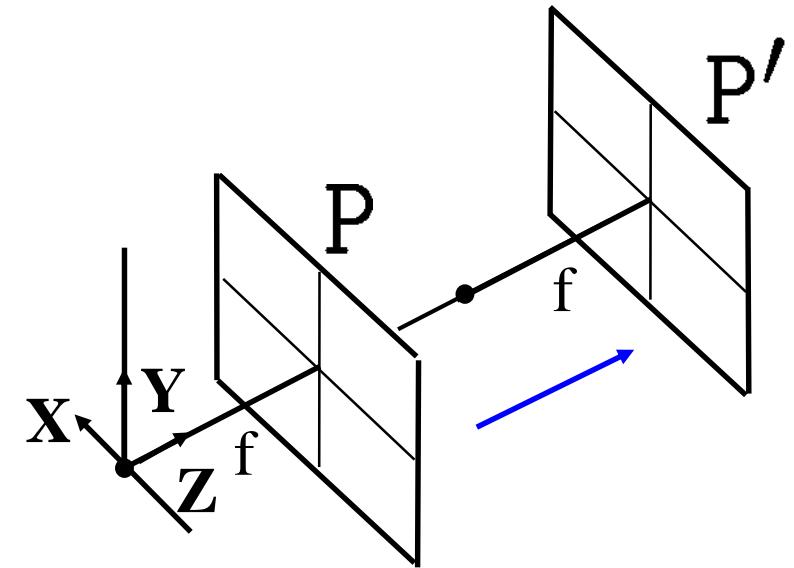


Example 2: compute F for forward-translating camera

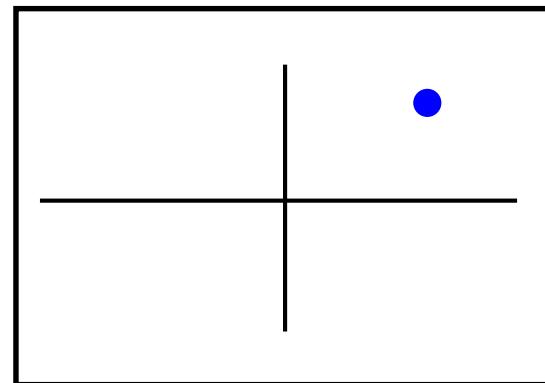
From $\mathbf{l}' = \mathbf{F}\mathbf{x}$ the epipolar line for the point $\mathbf{x} = (x, y, 1)^\top$ is

$$\mathbf{l}' = \begin{bmatrix} 0 & -1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} \begin{pmatrix} x \\ y \\ 1 \end{pmatrix} = \begin{pmatrix} -y \\ x \\ 0 \end{pmatrix}$$

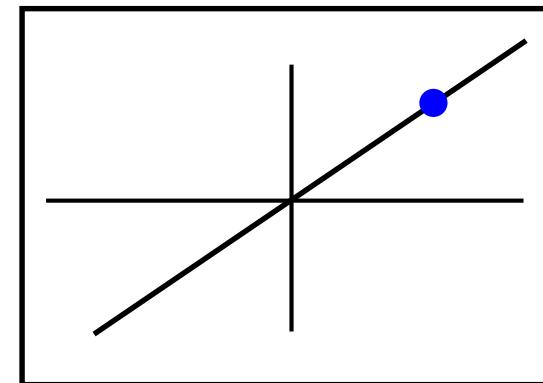
The points $(x, y, 1)^\top$ and $(0, 0, 1)^\top$ lie on this line



first image



second image





2017-10-16

CM50248 – Visual Understanding 1 – Lecture 5

69



2017-10-16

CM50248 – Visual Understanding 1 – Lecture 5

70

Summary: properties of the fundamental matrix

- F is a rank 2 homogeneous matrix with 7 degrees of freedom.
- Point correspondence:
if x and x' are corresponding image points, then
 $x'^\top Fx = 0$.
- Epipolar lines:
 - ◊ $l' = Fx$ is the epipolar line corresponding to x .
 - ◊ $l = F^\top x'$ is the epipolar line corresponding to x' .
- Epipoles:
 - ◊ $Fe = 0$.
 - ◊ $F^\top e' = 0$.
- Computation from camera matrices P, P' :
$$P = K[I \mid 0], \quad P' = K'[R \mid t], \quad F = K'^{-\top} [t]_\times R K^{-1}$$

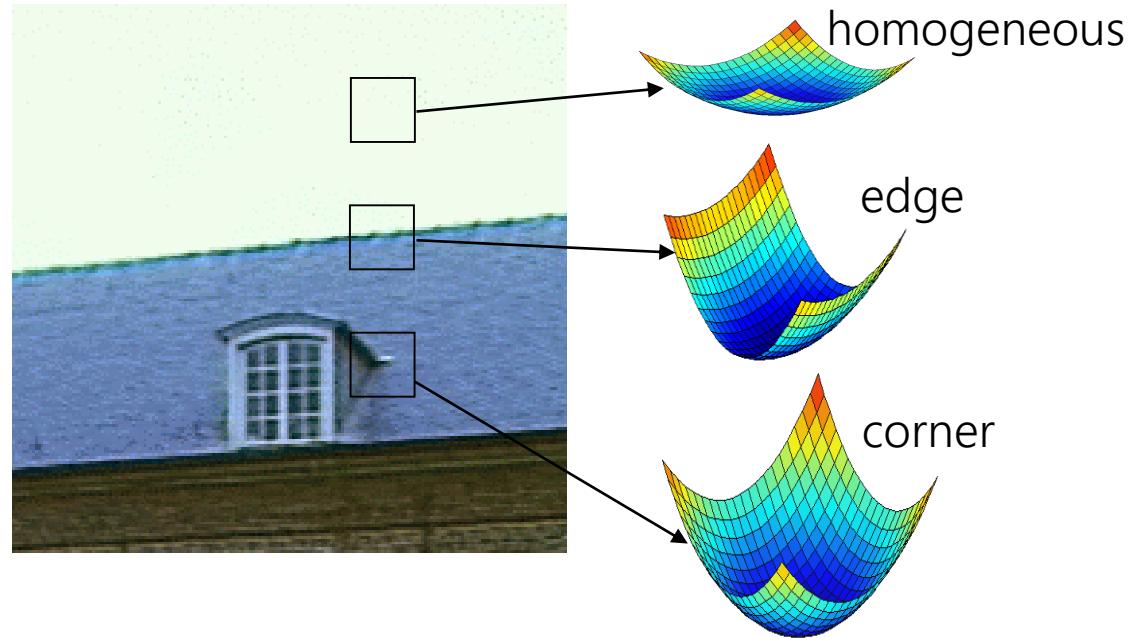
Fundamental matrix estimation

Automatic computation of F

- (i) Interest points
 - (ii) Putative correspondences
 - (iii) RANSAC
 - (iv) Non-linear re-estimation of F
 - (v) Guided matching
- (repeat (iv) and (v) until stable)

Feature points

- Find points that differ as much as possible from neighbouring points, for example:
 - Harris corners
 - SIFT
 - SURF
 - etc.



RANSAC

Step 1. Extract features

Step 2. Compute a set of potential matches

Step 3. do

Step 3.1 select minimal sample (i.e. 7 matches)

Step 3.2 compute solution(s) for F

Step 3.3 determine inliers (verify hypothesis)

until $\Gamma(\#inliers, \#samples) < 95\%$

} (generate hypothesis)

$$\Gamma = 1 - \left(1 - \left(\frac{\#inliers}{\#matches}\right)^7\right)^{\#samples}$$

Step 4. Compute F based on all inliers

Step 5. Look for additional matches

Step 6. Refine F based on all correct matches

#inliers	90%	80%	70%	60%	50%
#samples	5	13	35	106	382

Finding more matches

- Restrict search range to neighbourhood of epipolar line
 - ± 1.5 pixels
- relax disparity restriction (along epipolar line)

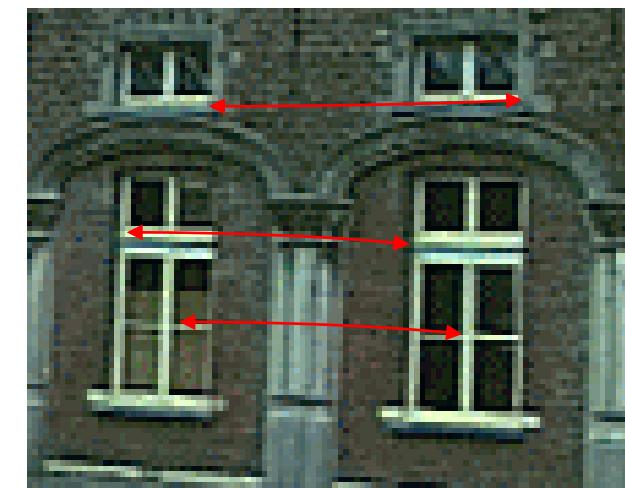
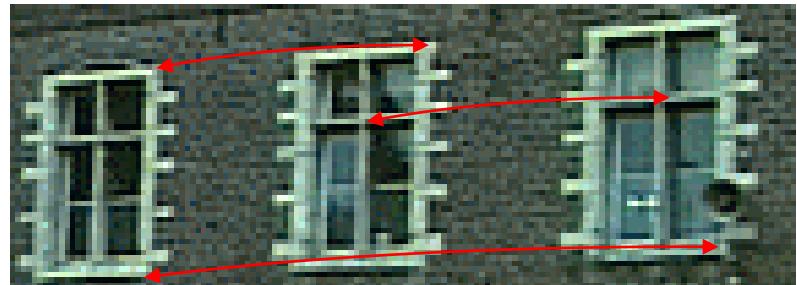


Degenerate cases

- Planar scene
- Pure rotation
- No unique solution
 - Remaining DOF filled by noise
 - Use simpler model (e.g. homography)
- Model selection (Torr et al., ICCV'98, Kanatani, Akaike)
 - Compare H and F according to expected residual error
(compensate for model complexity)

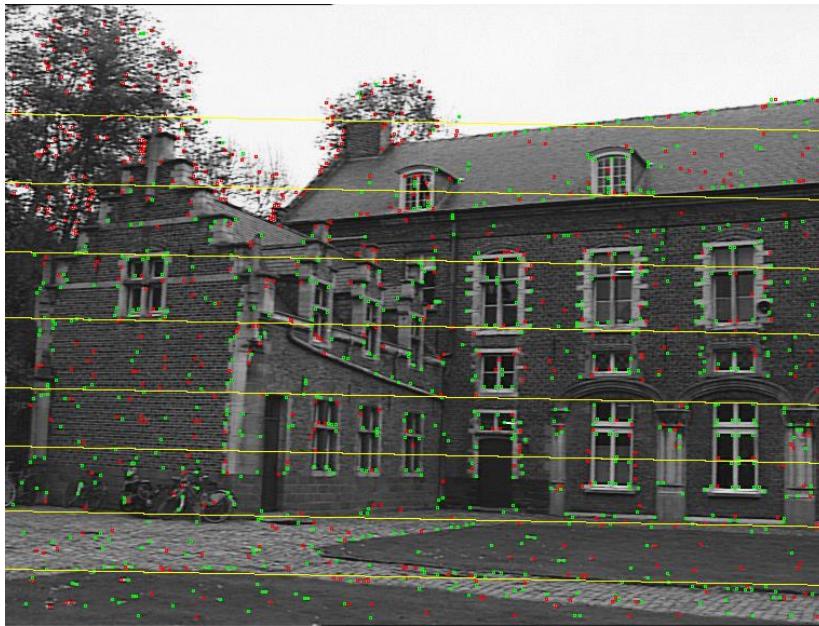
More problems

- Absence of sufficient features
(i.e. no texture)
- Repeated structure ambiguity
- Robust matcher also finds support for wrong hypothesis
 - solution: detect repetition
(Schaffalitzky and Zisserman,
BMVC'98)



Two-view geometry

geometric relations between two views is fully described by recovered 3×3 matrix \mathbf{F}



Triangulation

Triangulation

- **Problem:**

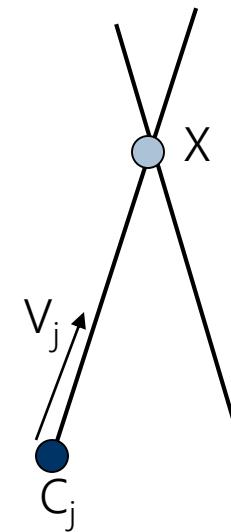
Given some points $\{(u_j, v_j)\}$ in correspondence across two or more images (taken from calibrated cameras), compute the 3D location X

Method I: intersecting viewing rays in 3D

- Minimise $\arg \min_{\mathbf{X}} \sum_j \|\mathbf{C}_j + s \mathbf{V}_j - \mathbf{X}\|$

- \mathbf{X} is the unknown 3D point
- \mathbf{C}_j is the optical center of camera j
- \mathbf{V}_j is the *viewing ray* for pixel (u_j, v_j)
- s_j is unknown distance along \mathbf{V}_j

- Advantage: geometrically intuitive



Method II: linear triangulation (algebraic solution)

Use the equations $\mathbf{x} = \mathbf{P}\mathbf{X}$ and $\mathbf{x}' = \mathbf{P}'\mathbf{X}$ to solve for \mathbf{X}

For the first camera:

$$\mathbf{P} = \begin{bmatrix} p_{11} & p_{12} & p_{13} & p_{14} \\ p_{21} & p_{22} & p_{23} & p_{24} \\ p_{31} & p_{32} & p_{33} & p_{34} \end{bmatrix} = \begin{bmatrix} \mathbf{p}^{1T} \\ \mathbf{p}^{2T} \\ \mathbf{p}^{3T} \end{bmatrix}$$

where \mathbf{p}^{iT} are the rows of \mathbf{P}

- eliminate unknown scale in $\lambda\mathbf{x} = \mathbf{P}\mathbf{X}$ by forming a cross product $\mathbf{x} \times (\mathbf{P}\mathbf{X}) = \mathbf{0}$

$$x(\mathbf{p}^{3T}\mathbf{X}) - (\mathbf{p}^{1T}\mathbf{X}) = 0$$

$$y(\mathbf{p}^{3T}\mathbf{X}) - (\mathbf{p}^{2T}\mathbf{X}) = 0$$

$$x(\mathbf{p}^{2T}\mathbf{X}) - y(\mathbf{p}^{1T}\mathbf{X}) = 0$$

- rearrange as (first two equations only)

$$\begin{bmatrix} xp^{3T} - p^{1T} \\ yp^{3T} - p^{2T} \end{bmatrix} \mathbf{X} = \mathbf{0}$$

Method II: linear triangulation (algebraic solution)

Similarly for the second camera:

$$\begin{bmatrix} x' \mathbf{p}'^3{}^\top - \mathbf{p}'^1{}^\top \\ y' \mathbf{p}'^3{}^\top - \mathbf{p}'^2{}^\top \end{bmatrix} \mathbf{X} = \mathbf{0}$$

Collecting together gives

$$\mathbf{A}\mathbf{X} = \mathbf{0}$$

where \mathbf{A} is the 4×4 matrix

$$\mathbf{A} = \begin{bmatrix} x\mathbf{p}^3{}^\top - \mathbf{p}^1{}^\top \\ y\mathbf{p}^3{}^\top - \mathbf{p}^2{}^\top \\ x' \mathbf{p}'^3{}^\top - \mathbf{p}'^1{}^\top \\ y' \mathbf{p}'^3{}^\top - \mathbf{p}'^2{}^\top \end{bmatrix}$$

from which \mathbf{X} can be solved up to scale.

Problem: does not minimize anything meaningful

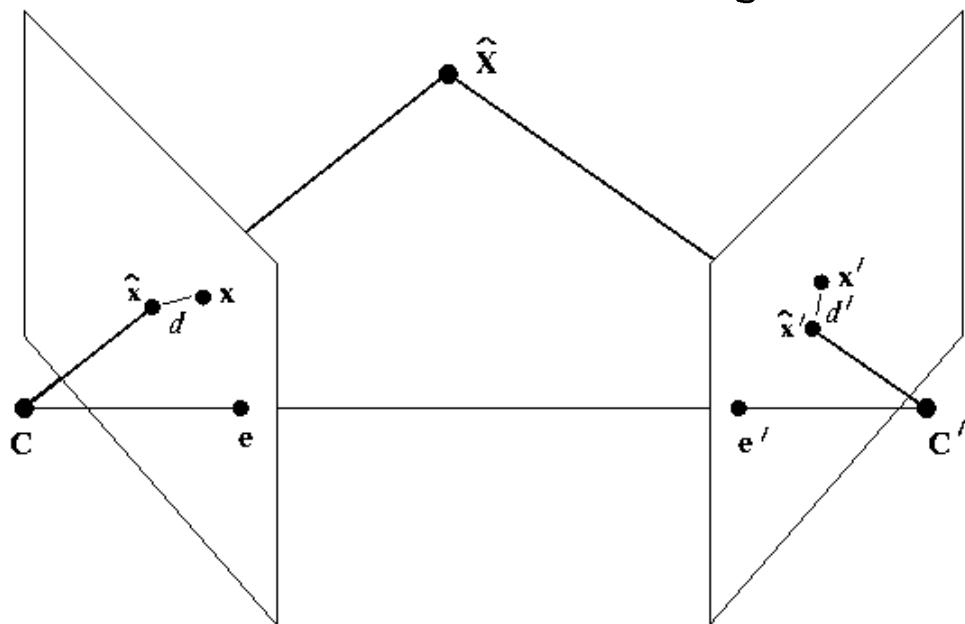
Advantage: extends to more than two views

Method III: minimising a geometric/statistical error

The idea is to estimate a 3D point $\hat{\mathbf{x}}$ which exactly satisfies the supplied camera geometry, so it projects as

$$\hat{\mathbf{x}} = \mathbf{P}\hat{\mathbf{x}} \quad \hat{\mathbf{x}}' = \mathbf{P}'\hat{\mathbf{x}}$$

and the aim is to estimate $\hat{\mathbf{x}}$ from the image measurements \mathbf{x} and \mathbf{x}' .



$$\min_{\hat{\mathbf{x}}} \quad \mathcal{C}(\mathbf{x}, \mathbf{x}') = d(\mathbf{x}, \hat{\mathbf{x}})^2 + d(\mathbf{x}', \hat{\mathbf{x}}')^2$$

where $d(*, *)$ is the Euclidean distance between the points.

Conclusion + Questions